Querying Data in Ignite

Edward Curren
ENTERPRISE ARCHITECT

@EdwardCurren http://www.edwardcurren.com

Overview

Architecture

Affinity function & affinity collocation

“Fat” keys

SQL & other types of queries

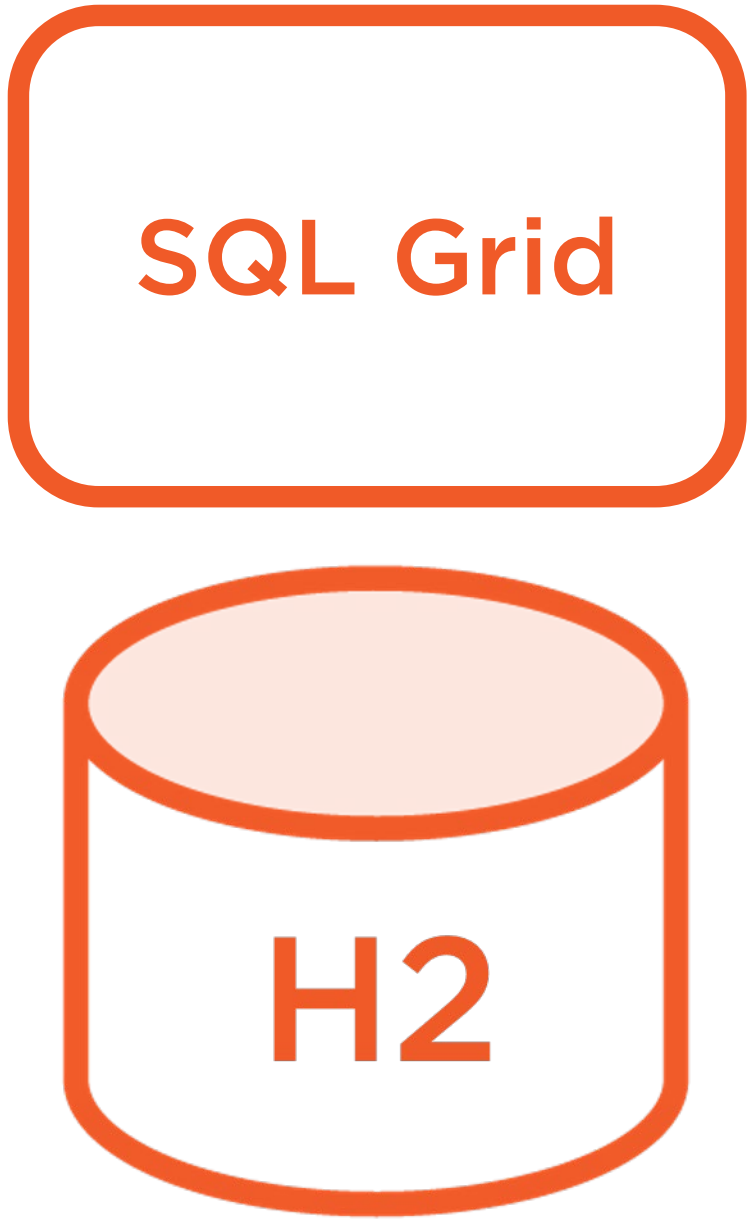
SQL Grid
H2

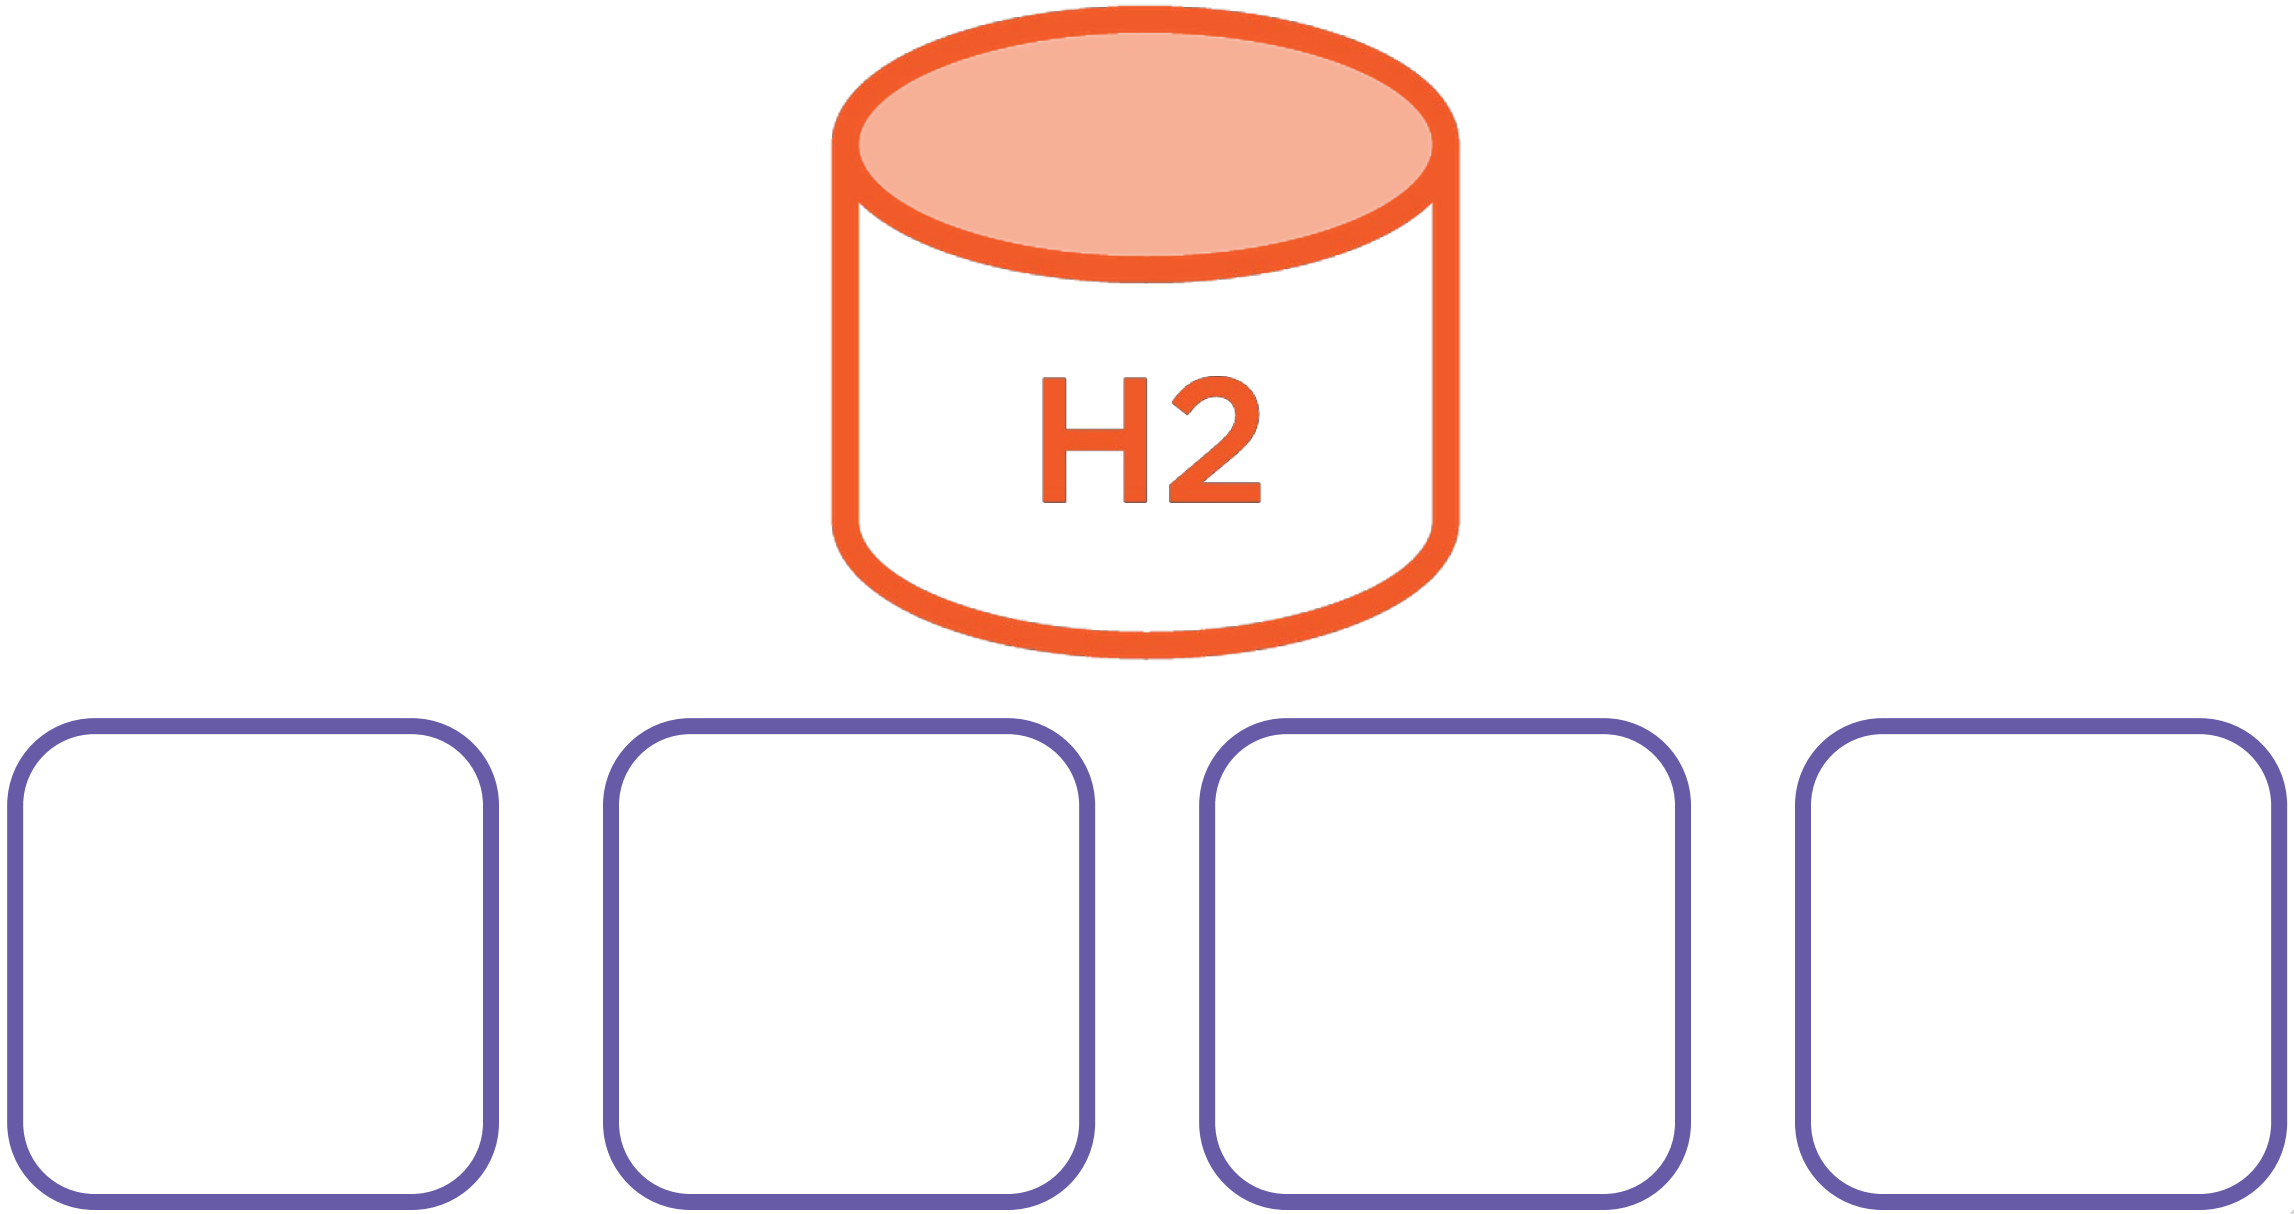
H2

- SQL Parsing
- SQL Optimization
- Execution Planning

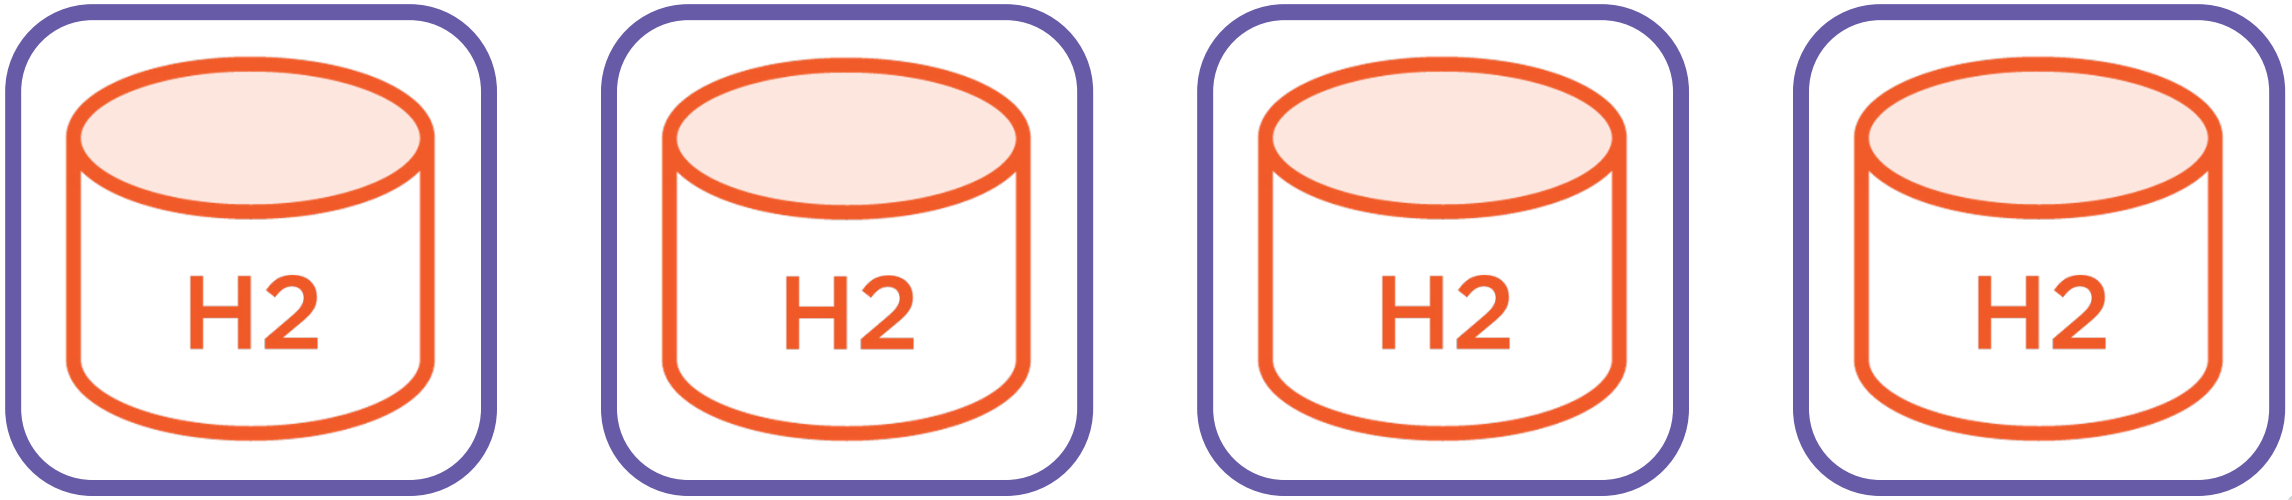

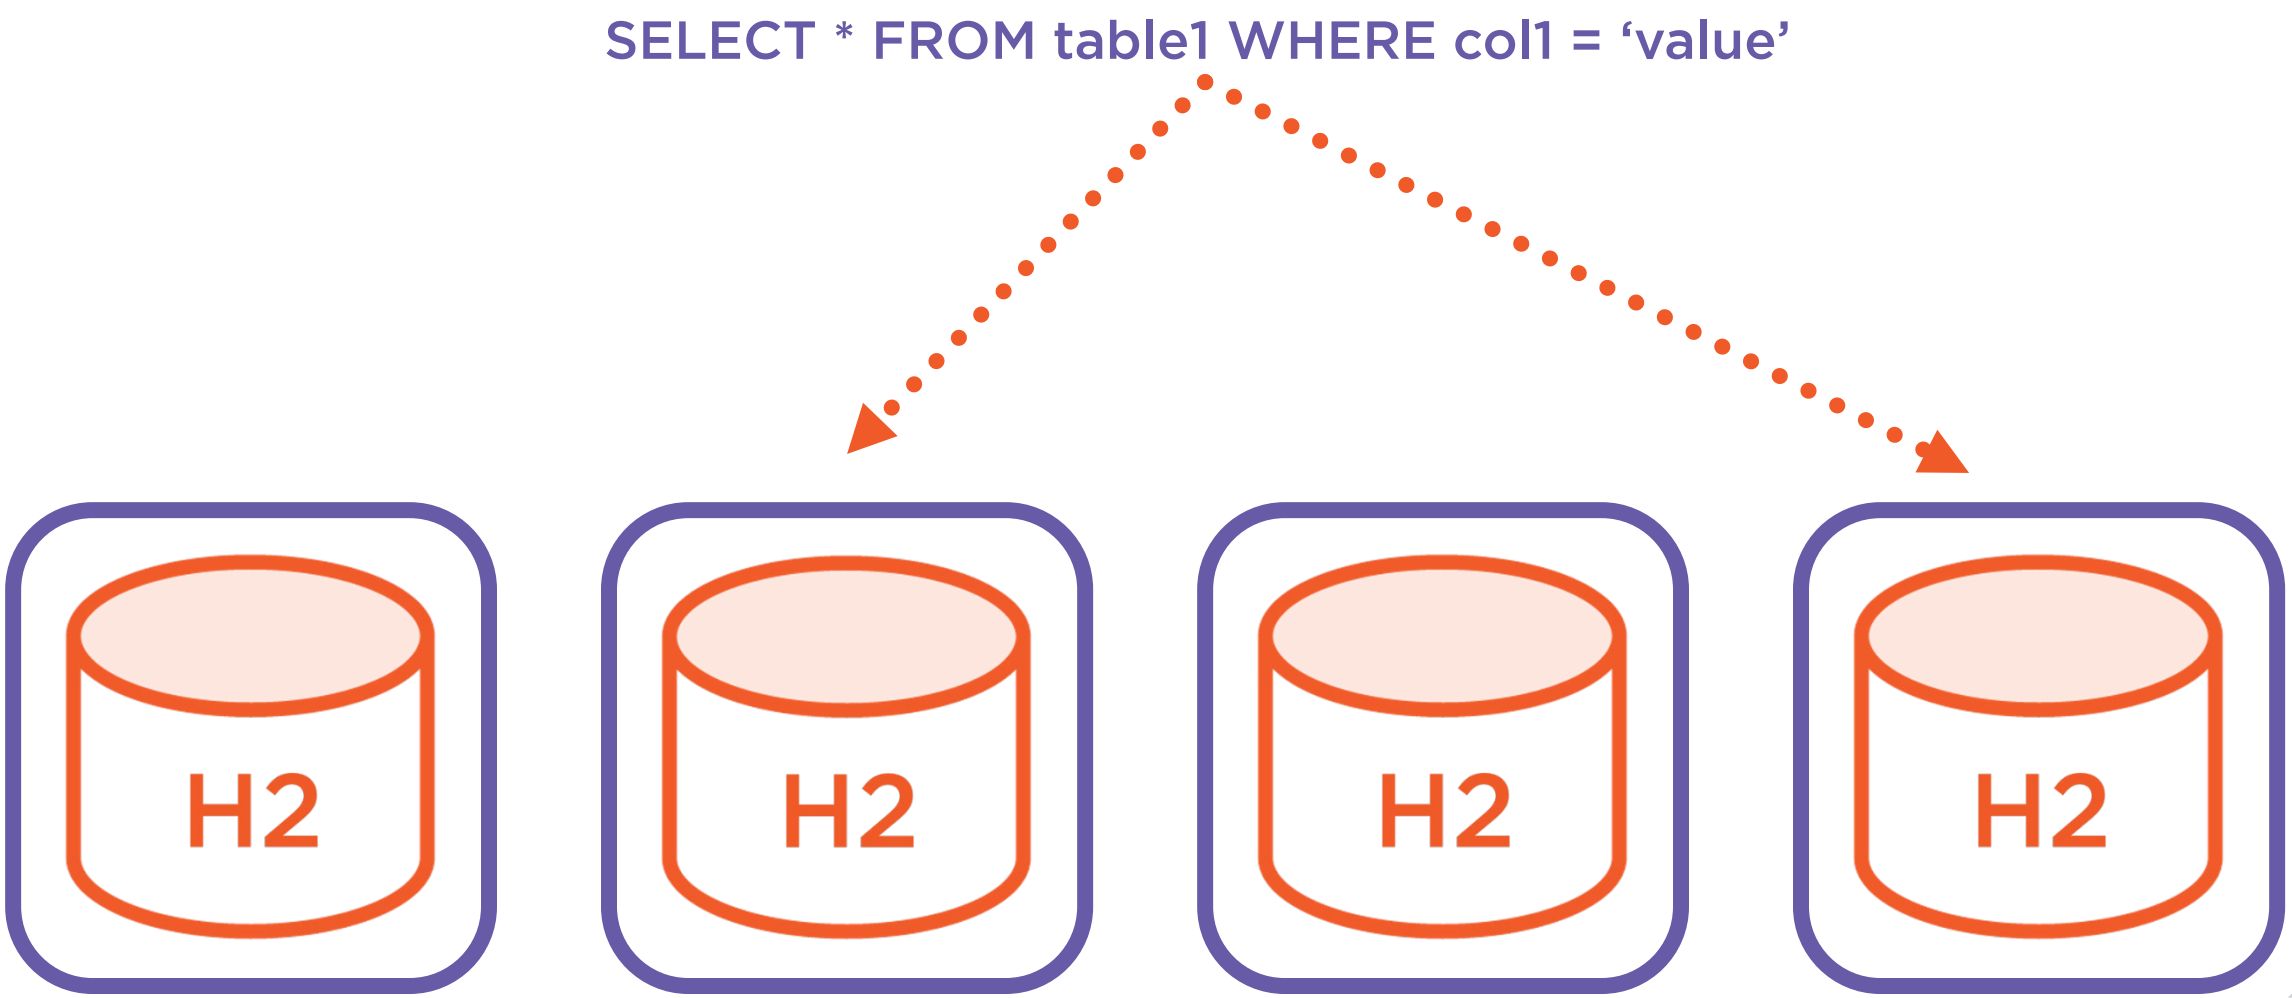
SELECT * FROM table1 WHERE col1 = ‘value’
H2
H2
H2
H2

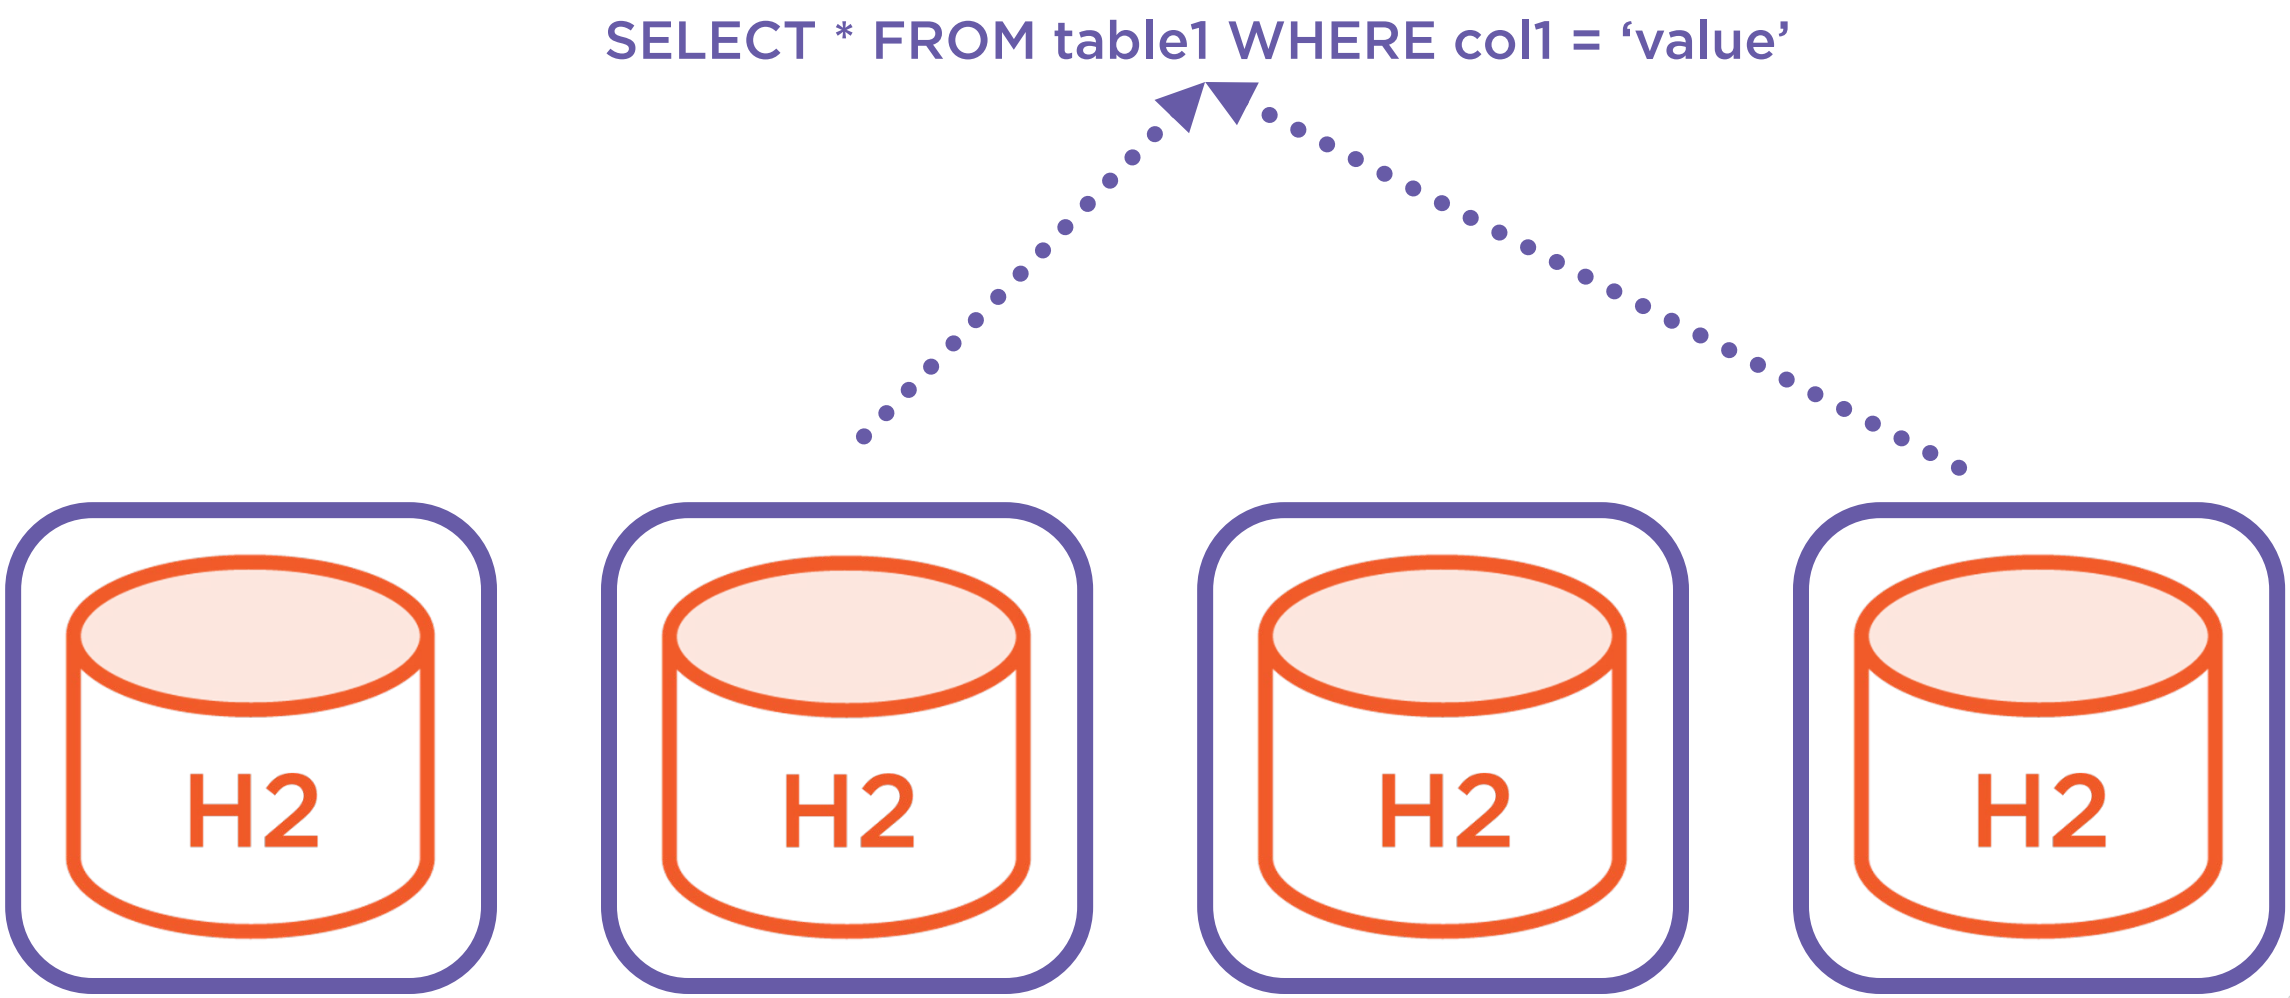
SELECT * FROM table1 WHERE col1 = ‘value’
H2
H2
H2
H2

SELECT * FROM table1 WHERE col1 = ‘value’

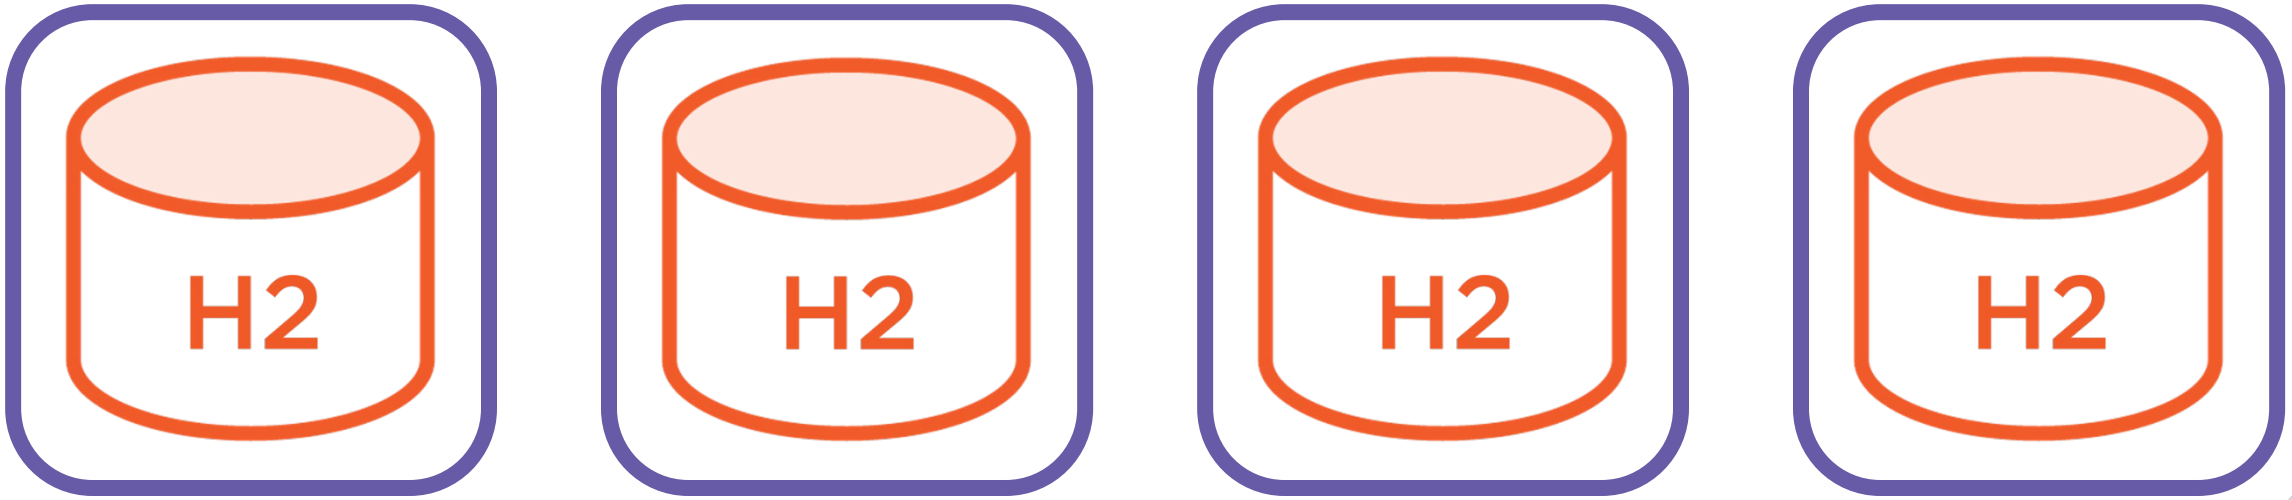

SELECT * FROM table1 WHERE col1 = ‘value’

INSERT INTO table2 (col1, col2) VALUES (val1, val2)

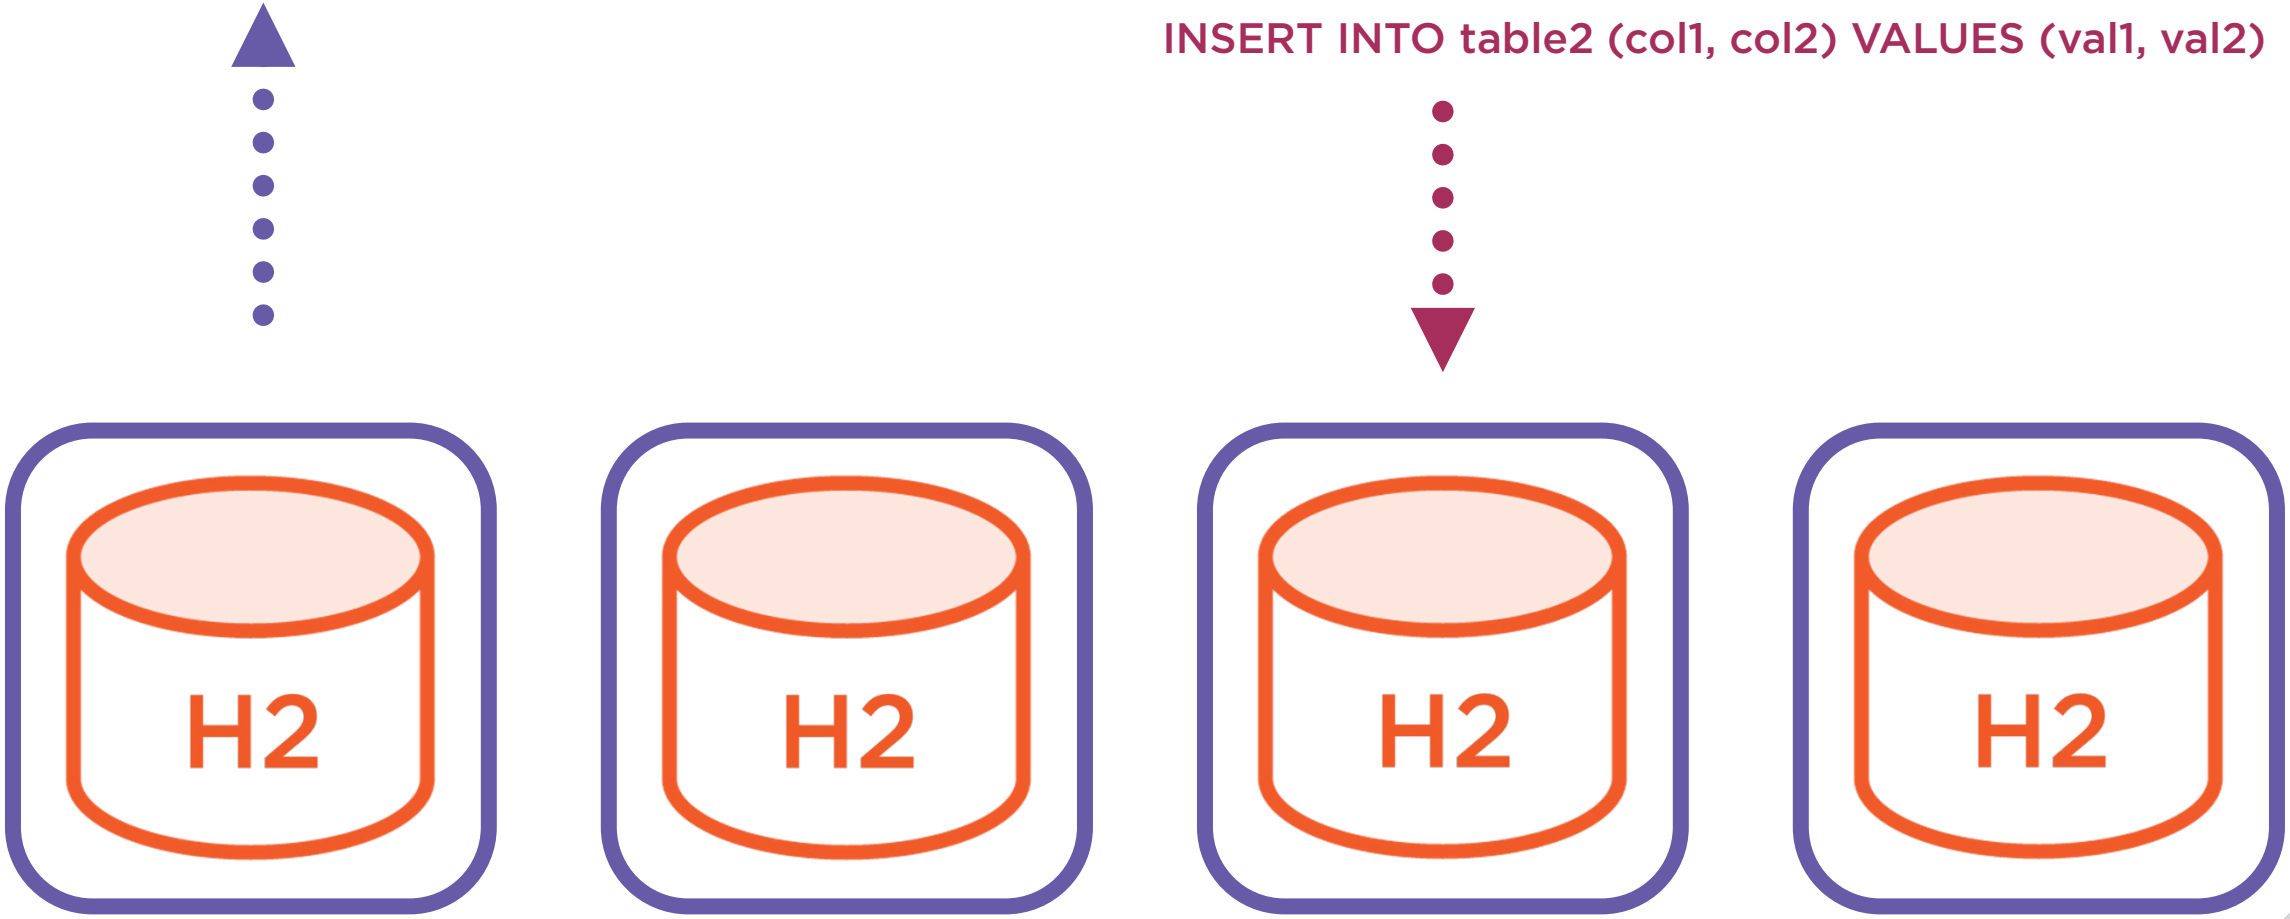

DELETE FROM table1 WHERE col1 = ‘value’

UPDATE table2 SET col1 = ‘val3’ WERE col1 = ‘val1’

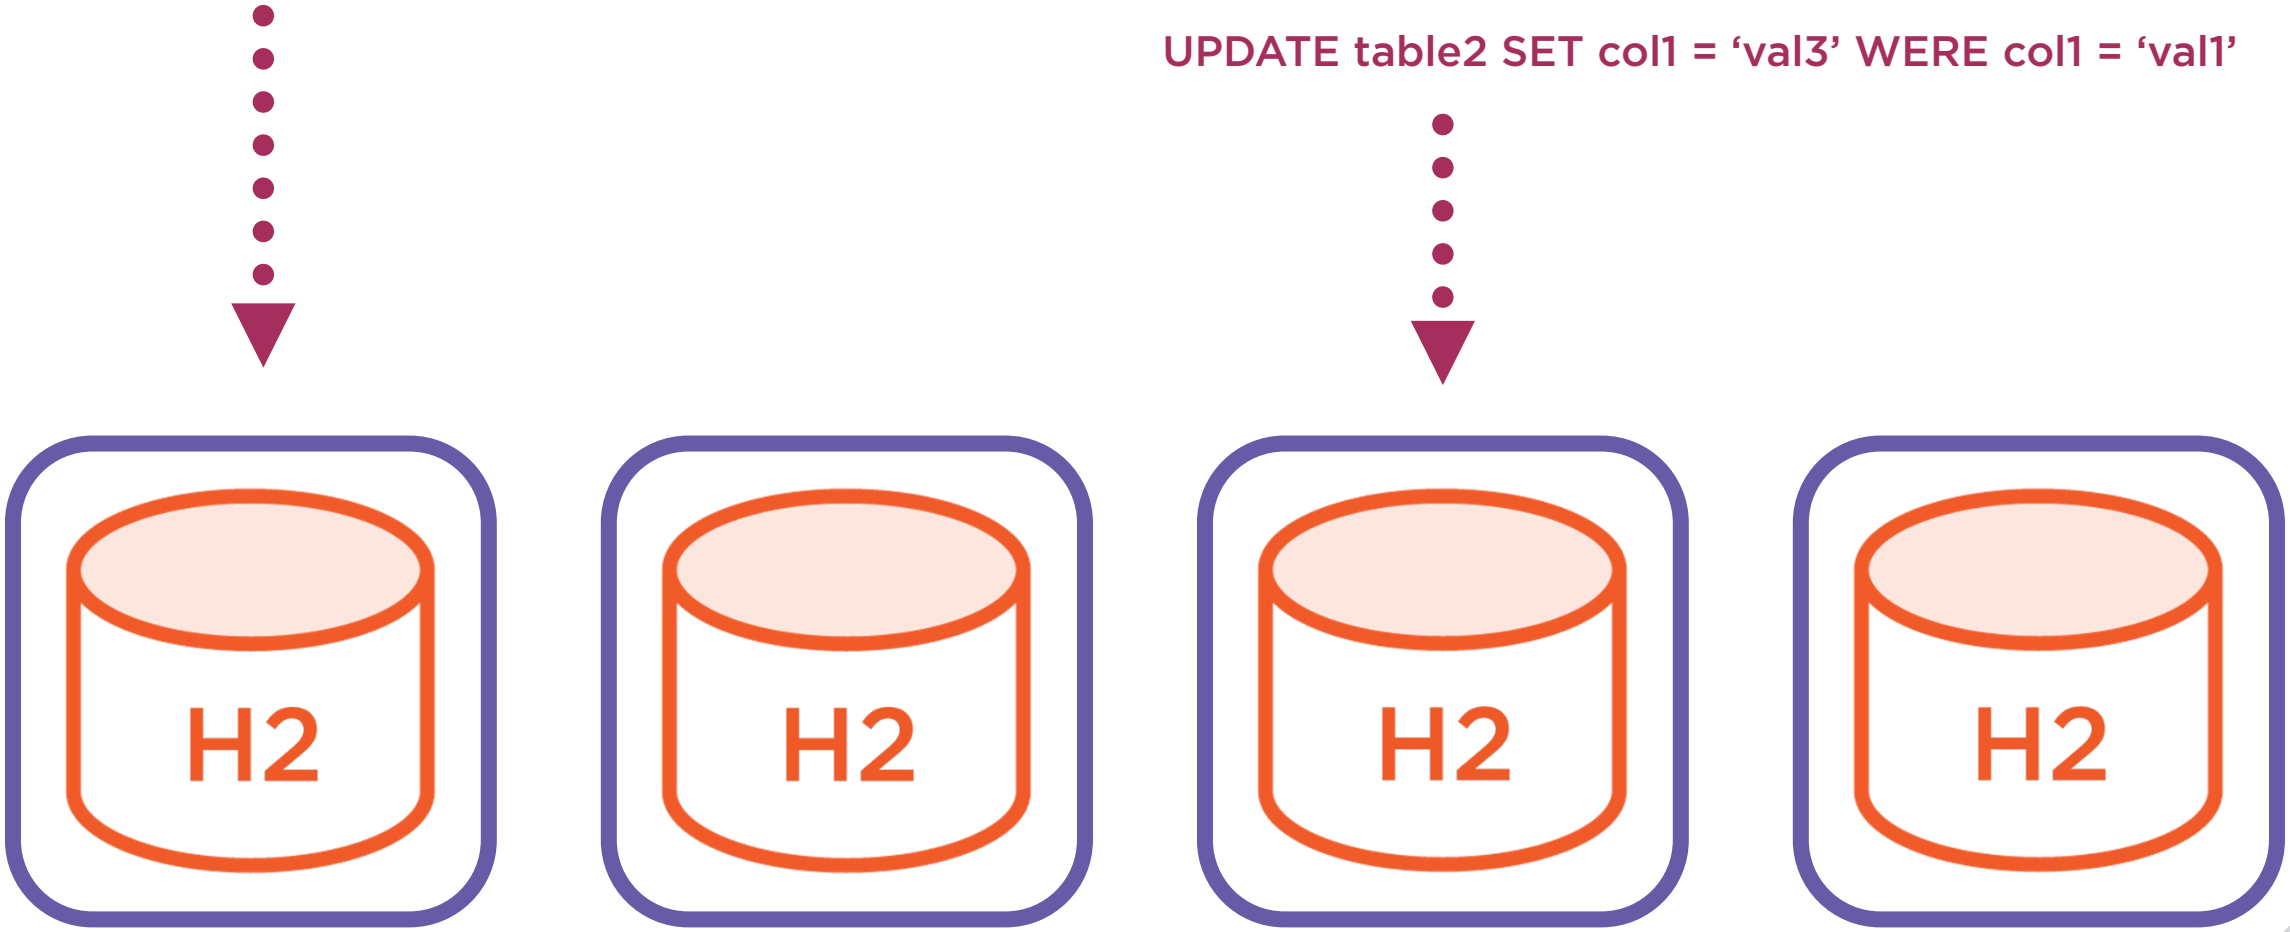

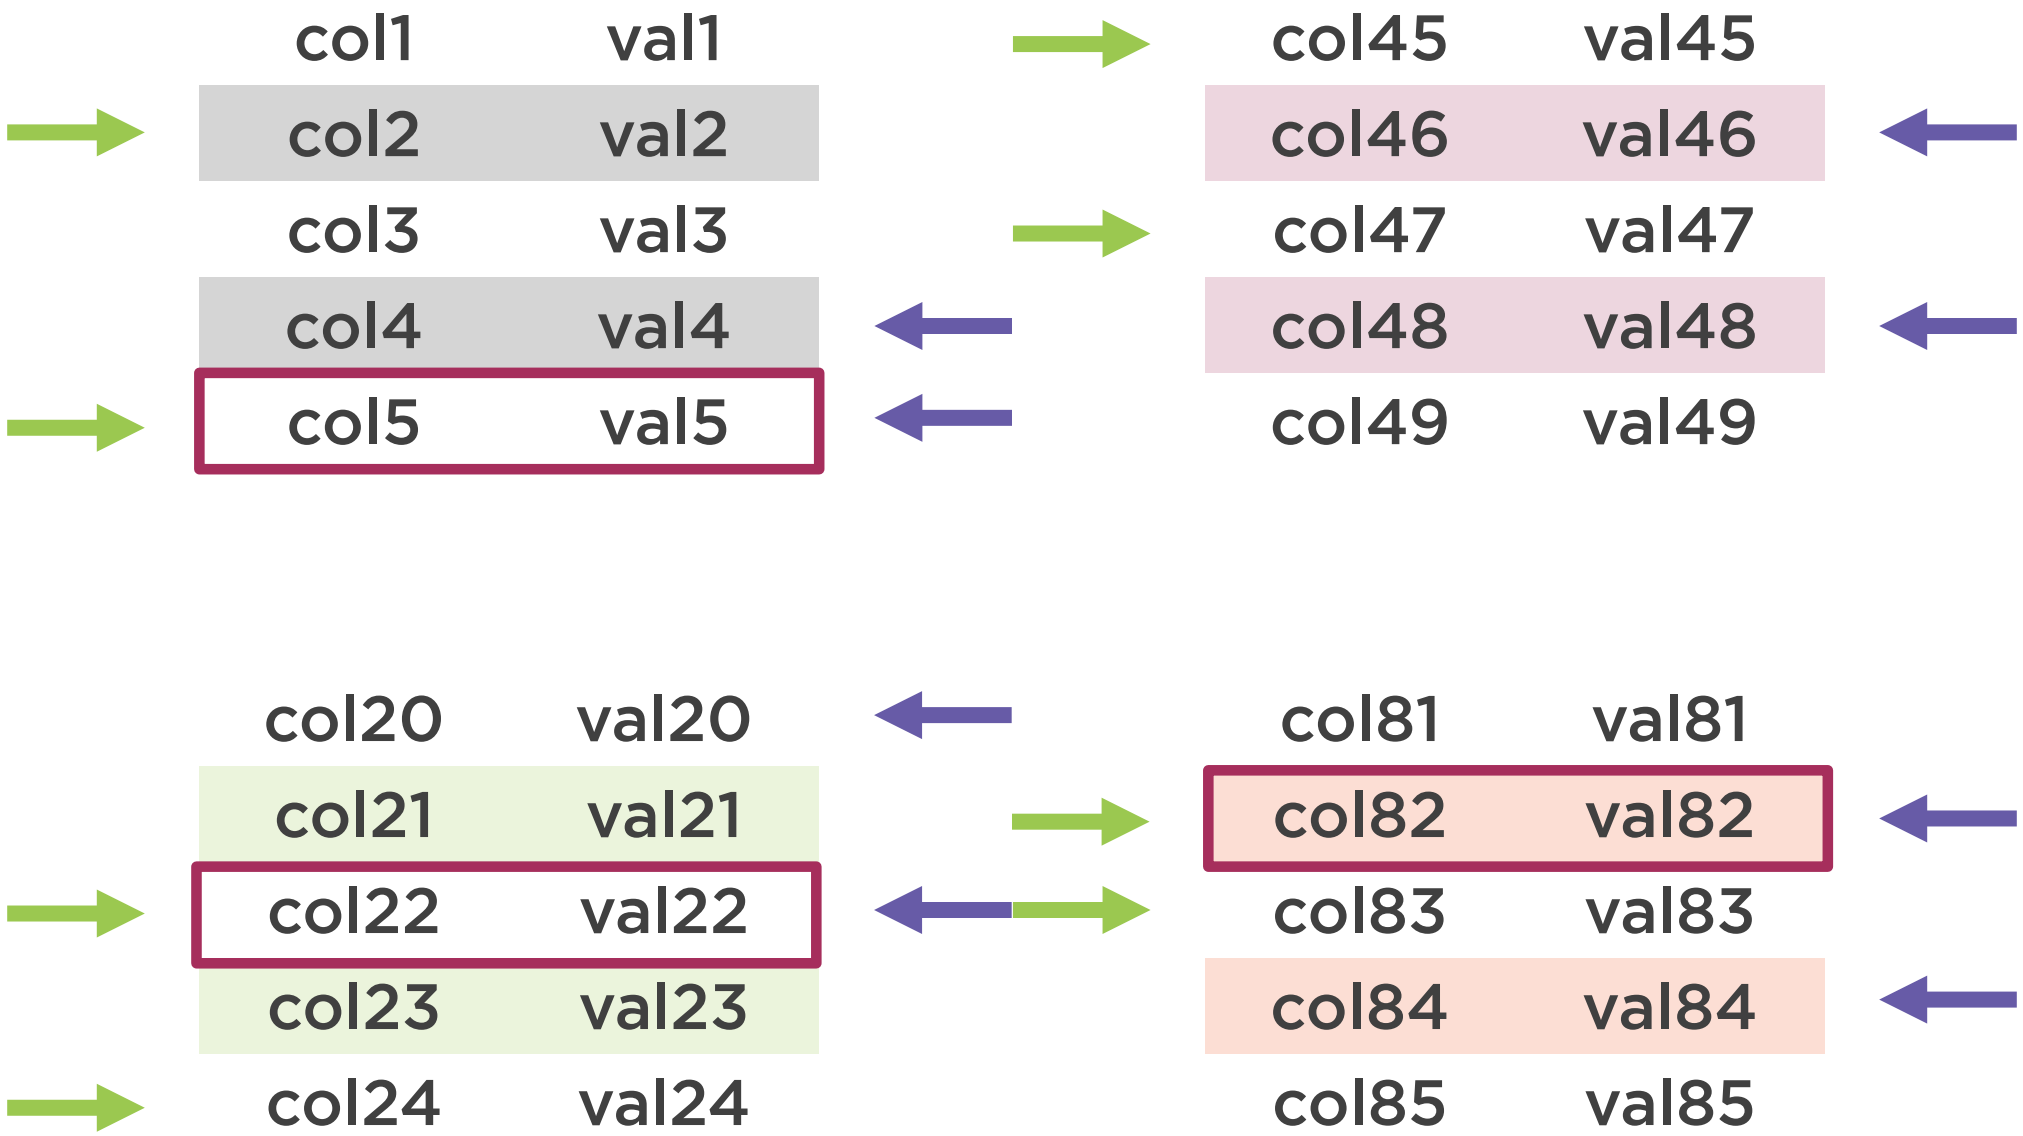
col1 val1
col2 val2
col3 val3
col4 val4
col5 val5
col45 val45
col46 val46
col47 val47
col48 val48
col49 val49
col20 val20
col21 val21
col22 val22
col23 val23
col24 val24
col81 val81
col82 val82
col83 val83
col84 val84
col85 val85

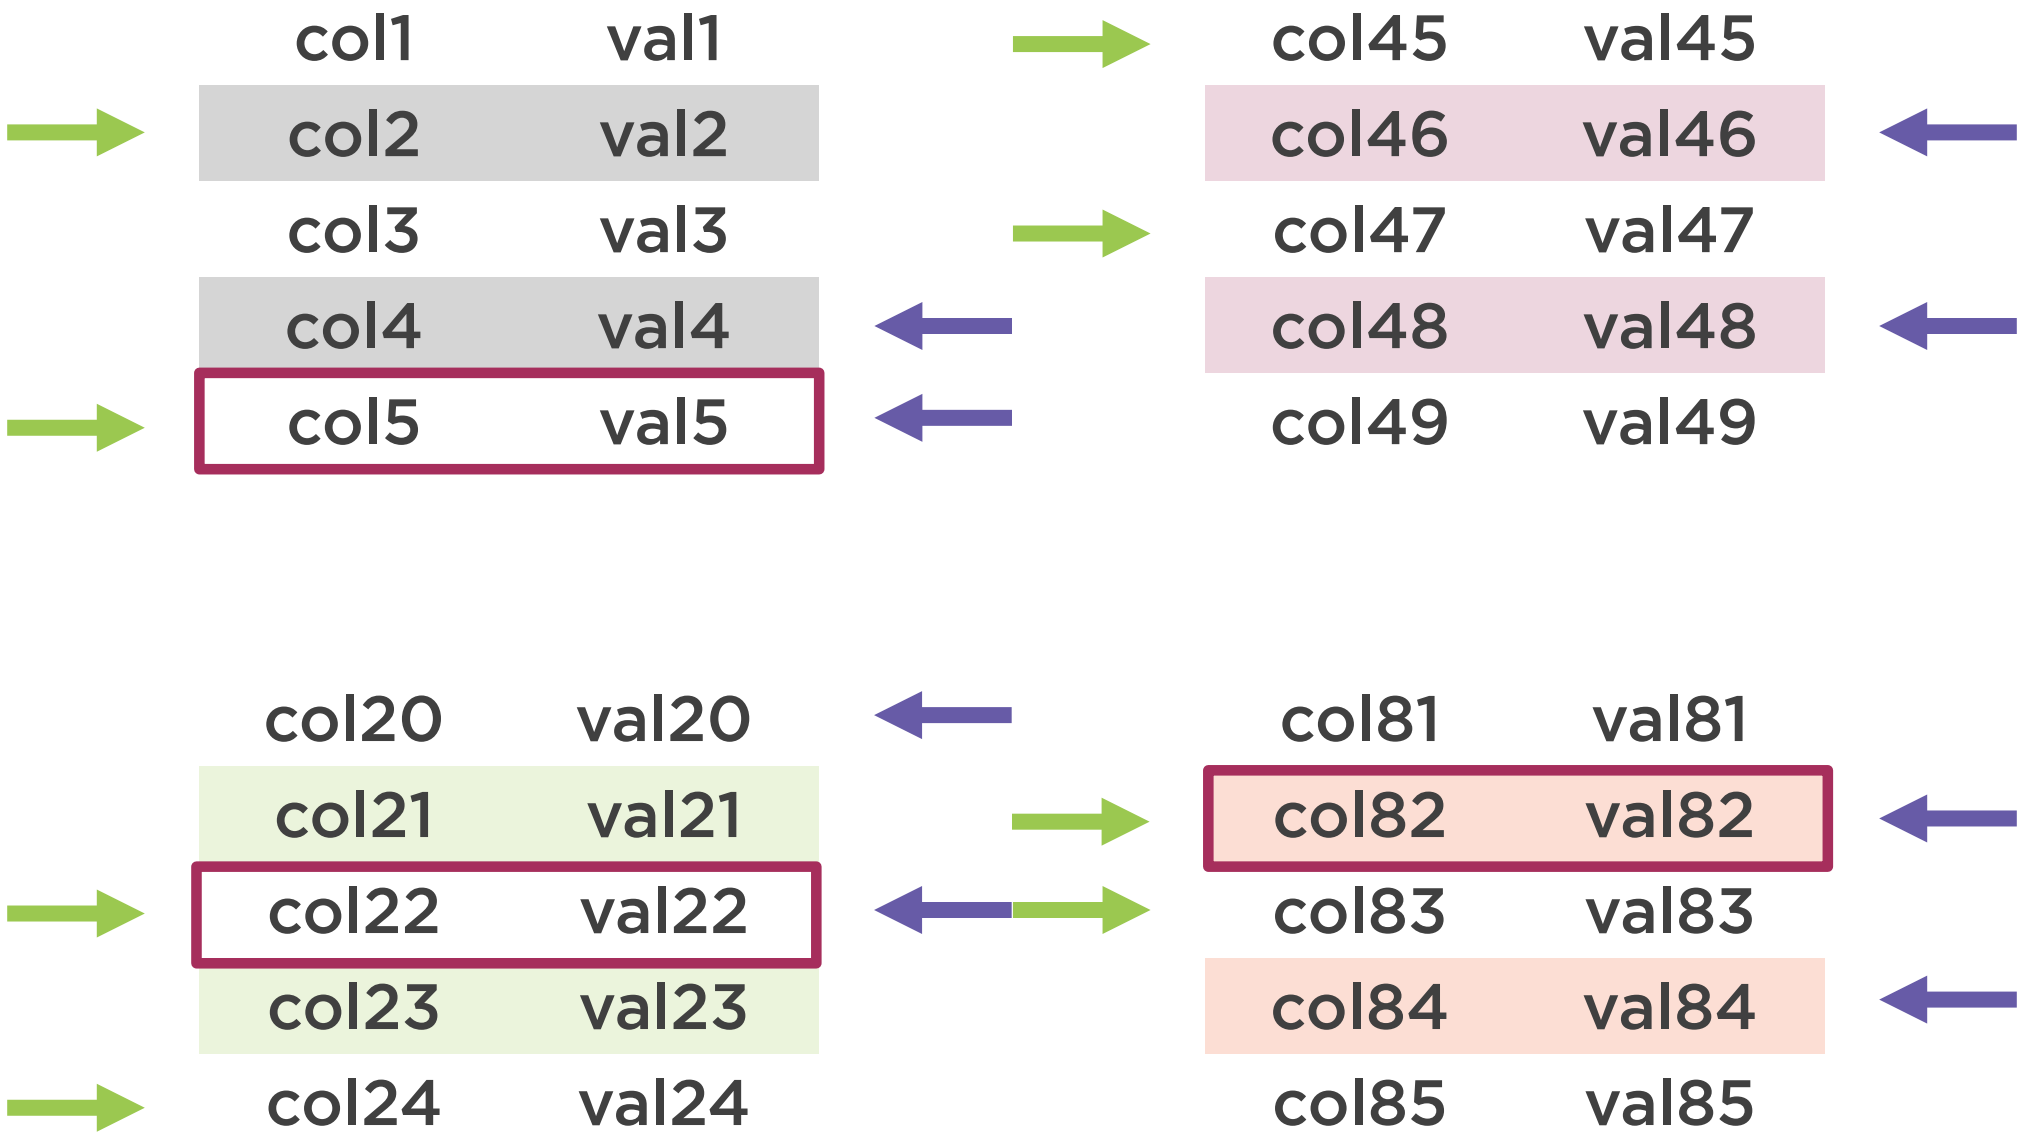
col1 val1
col2 val2
col3 val3
col4 val4
col5 val5
col45 val45
col46 val46
col47 val47
col48 val48
col49 val49
col20 val20
col21 val21
col22 val22
col23 val23
col24 val24
col81 val81
col82 val82
col83 val83
col84 val84
col85 val85

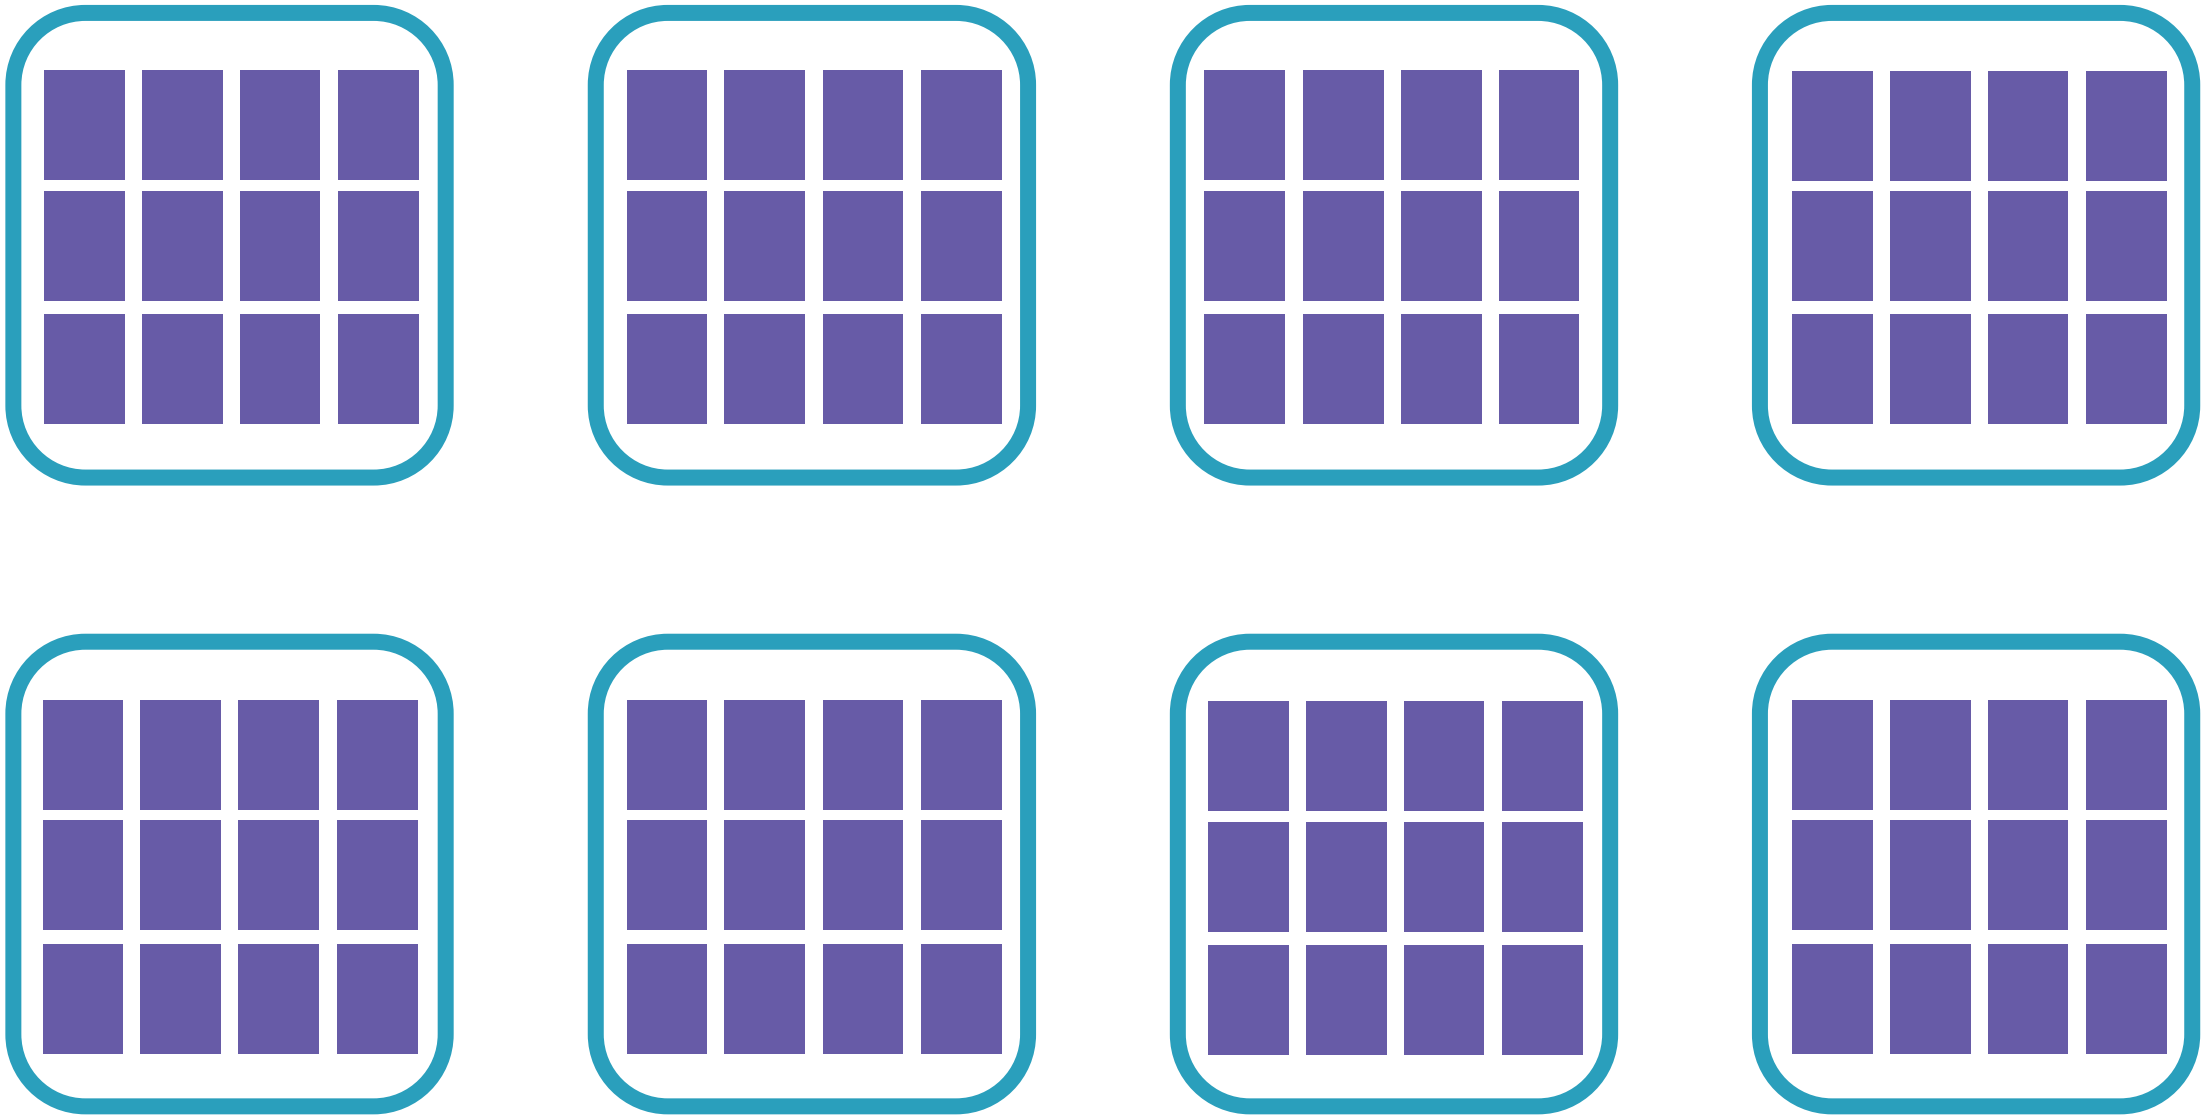

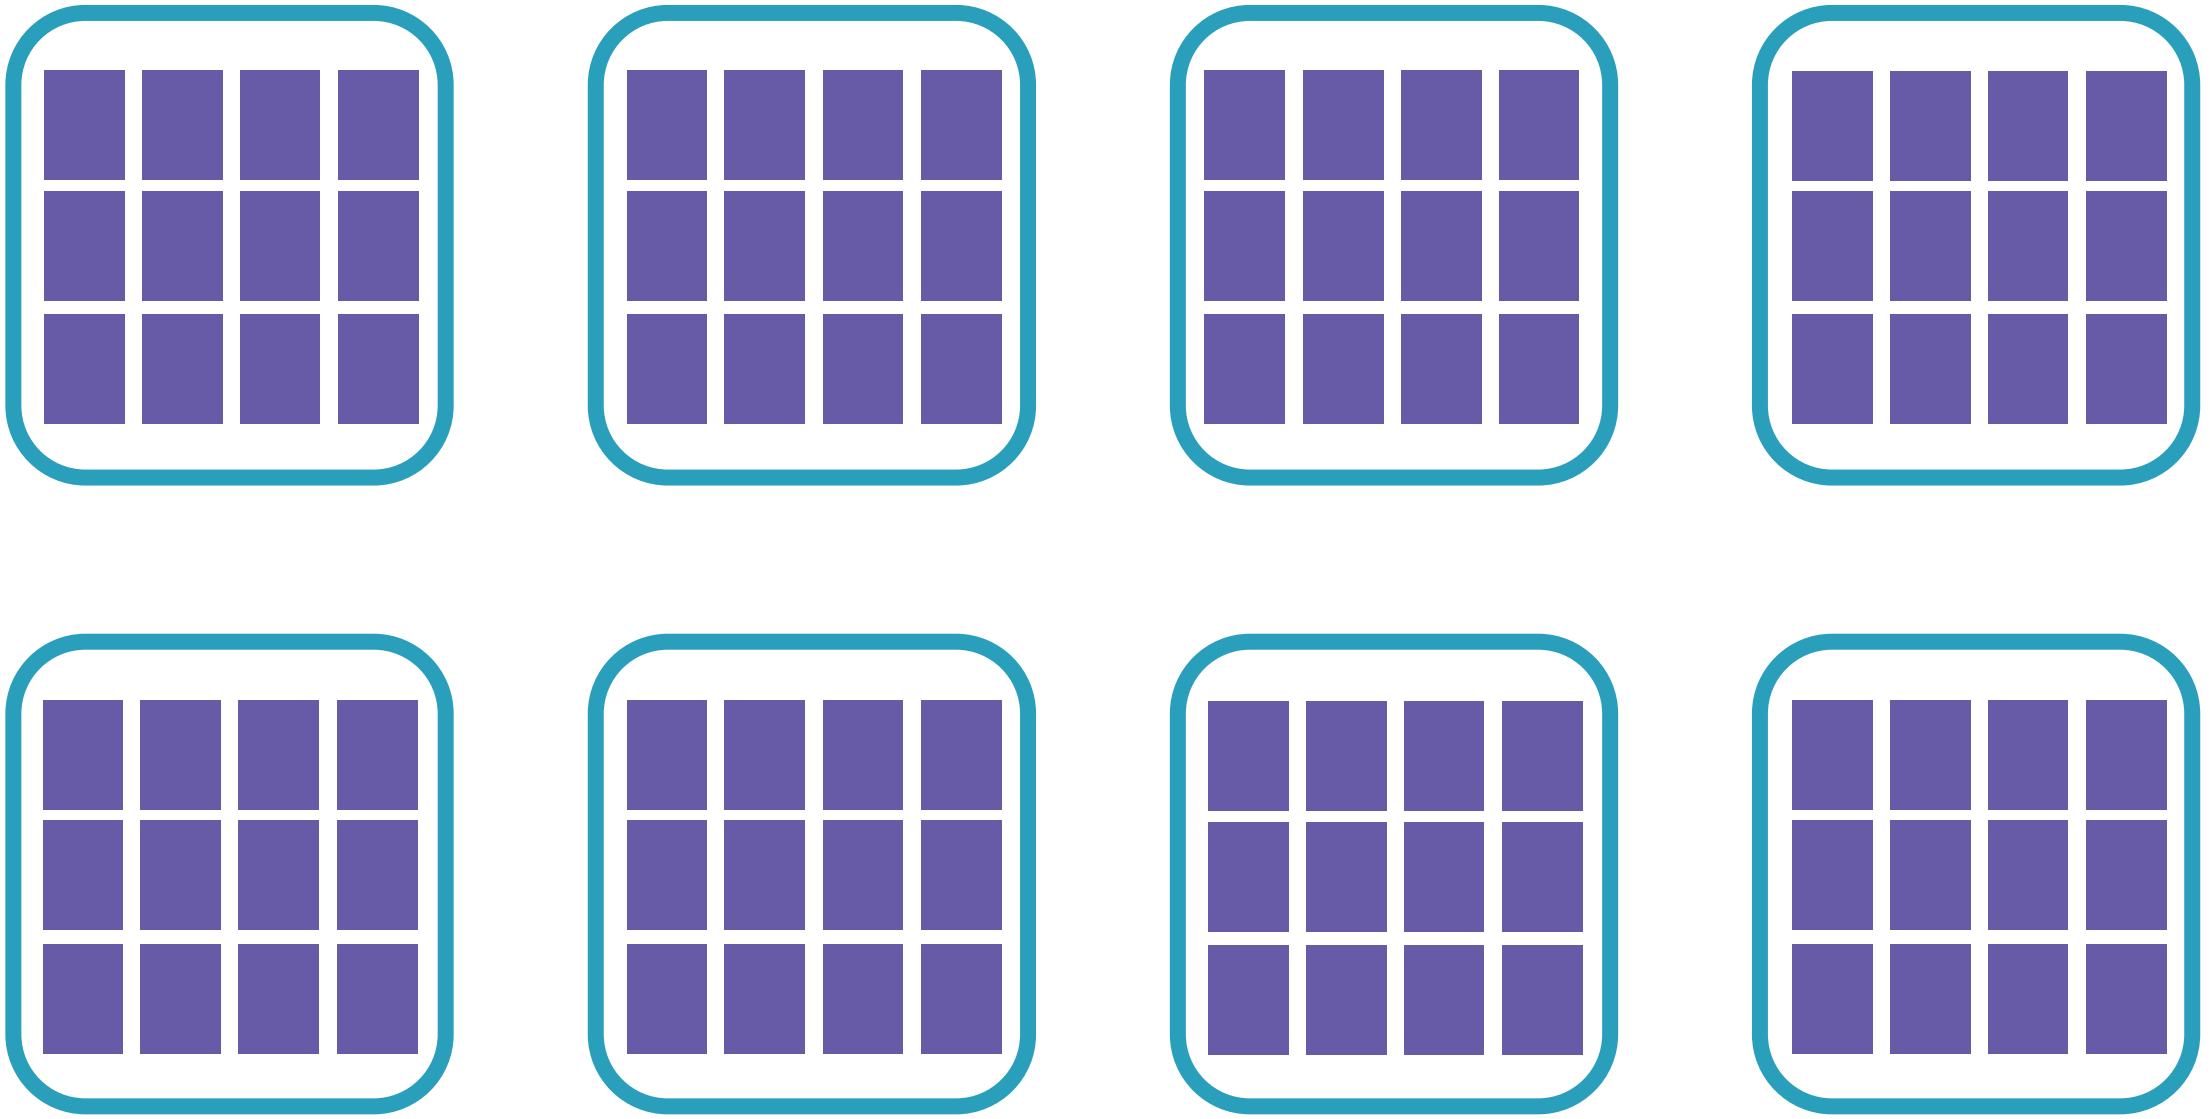

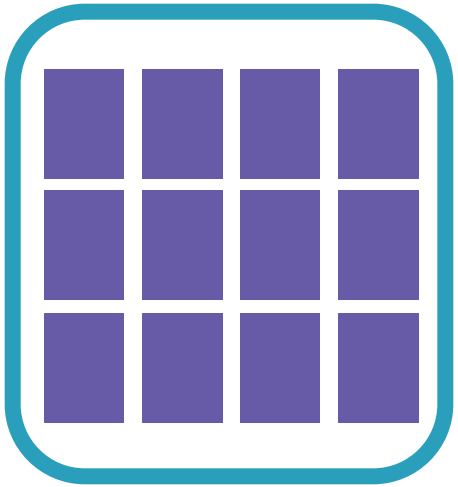

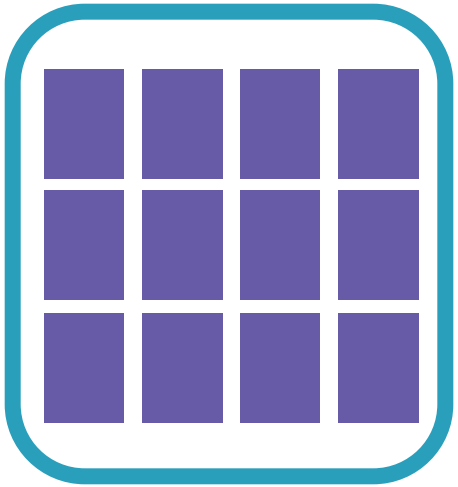

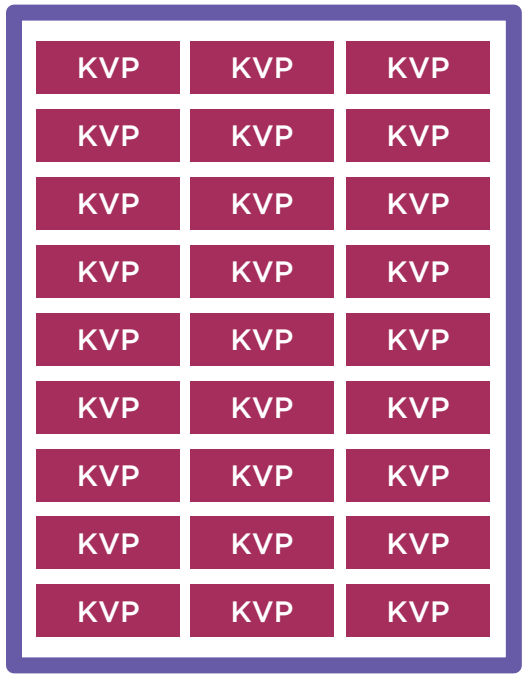
KVP KVP KVP
KVP KVP KVP
KVP KVP KVP
KVP KVP KVP
KVP KVP KVP
KVP KVP KVP
KVP KVP KVP
KVP KVP KVP
KVP KVP KVP

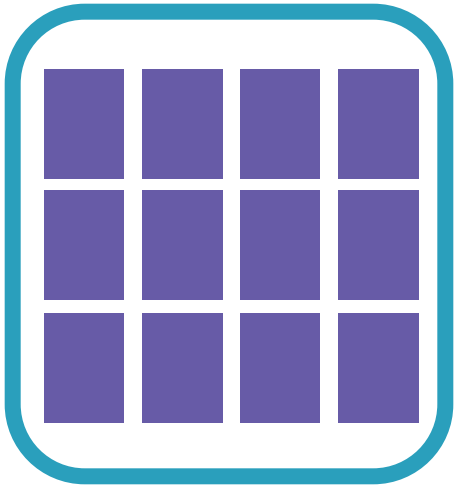

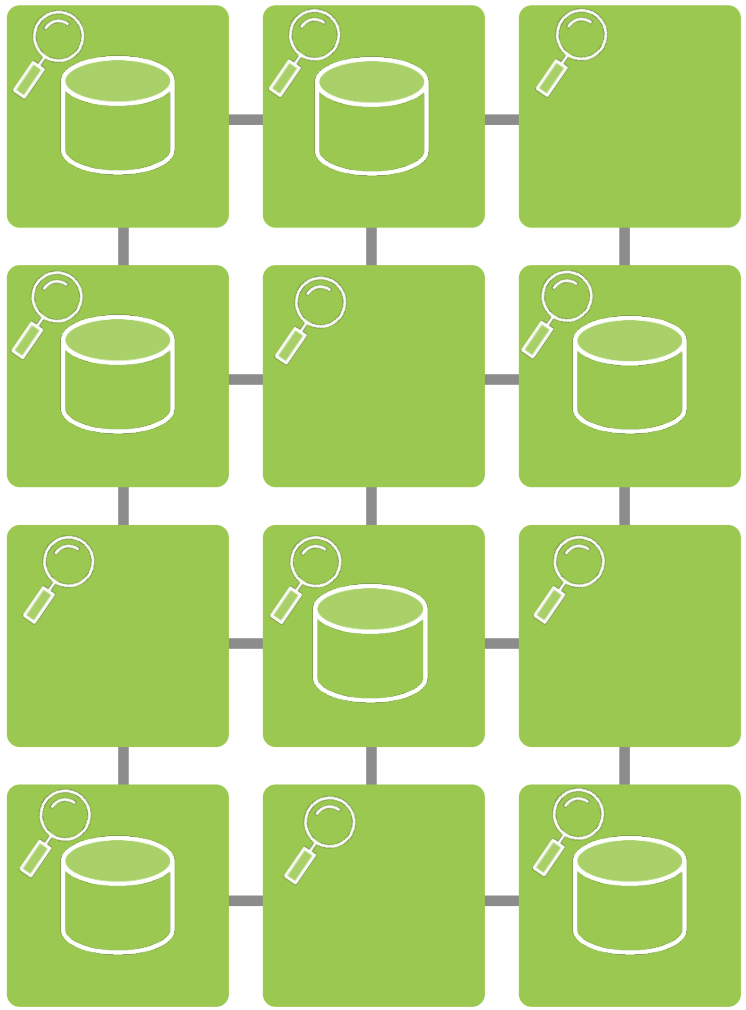

Resources

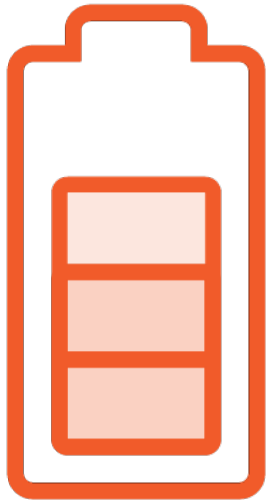

Network

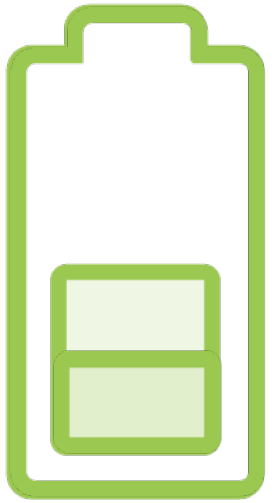

CPU

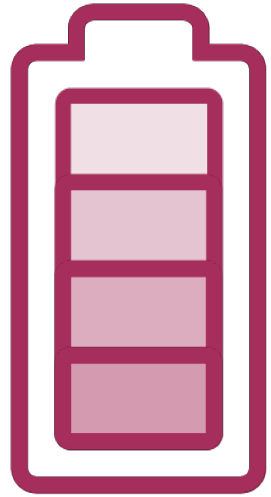

Memory

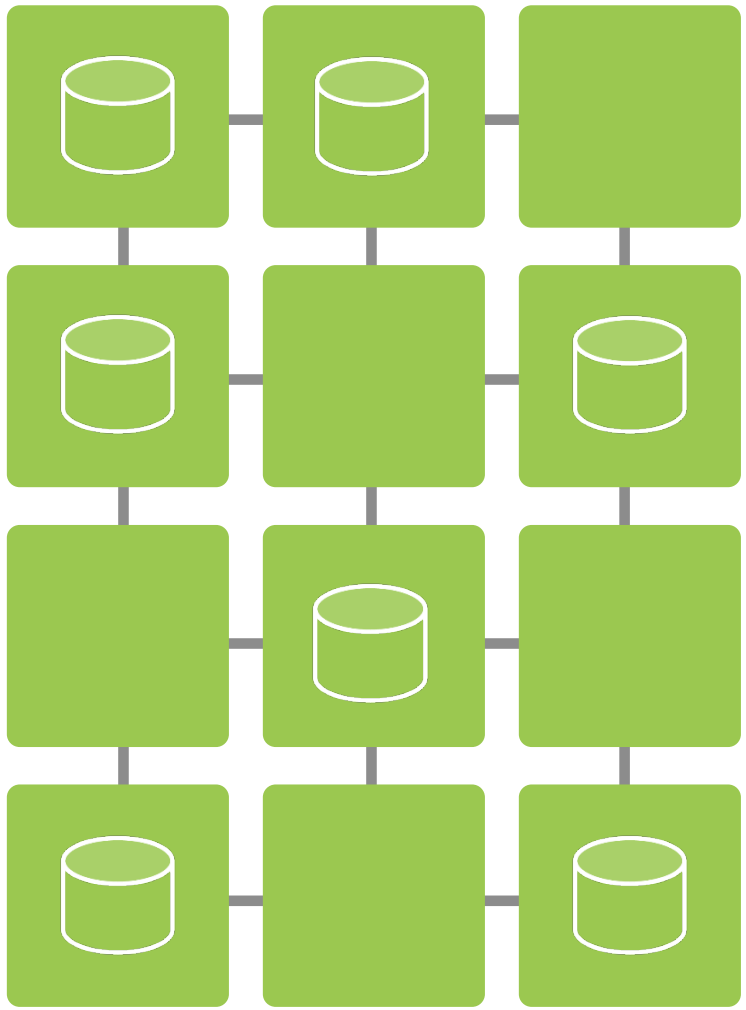

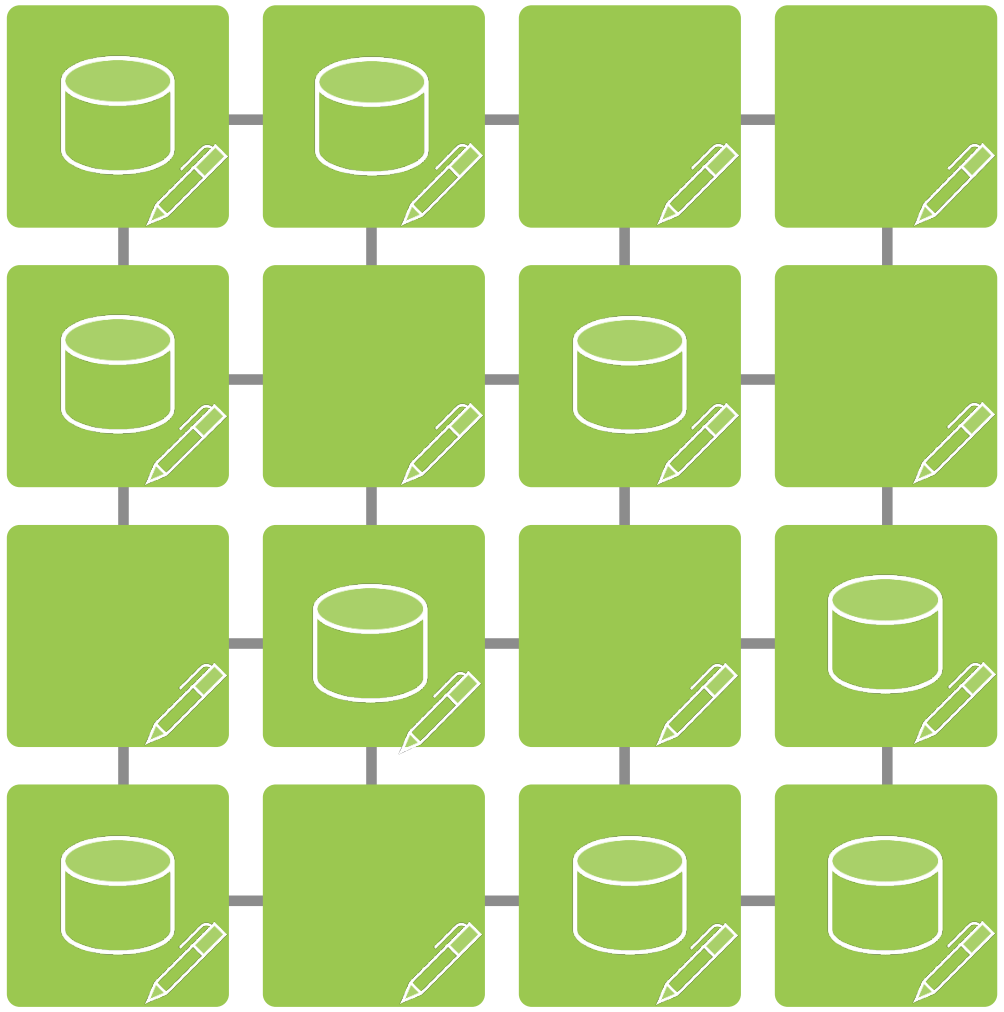

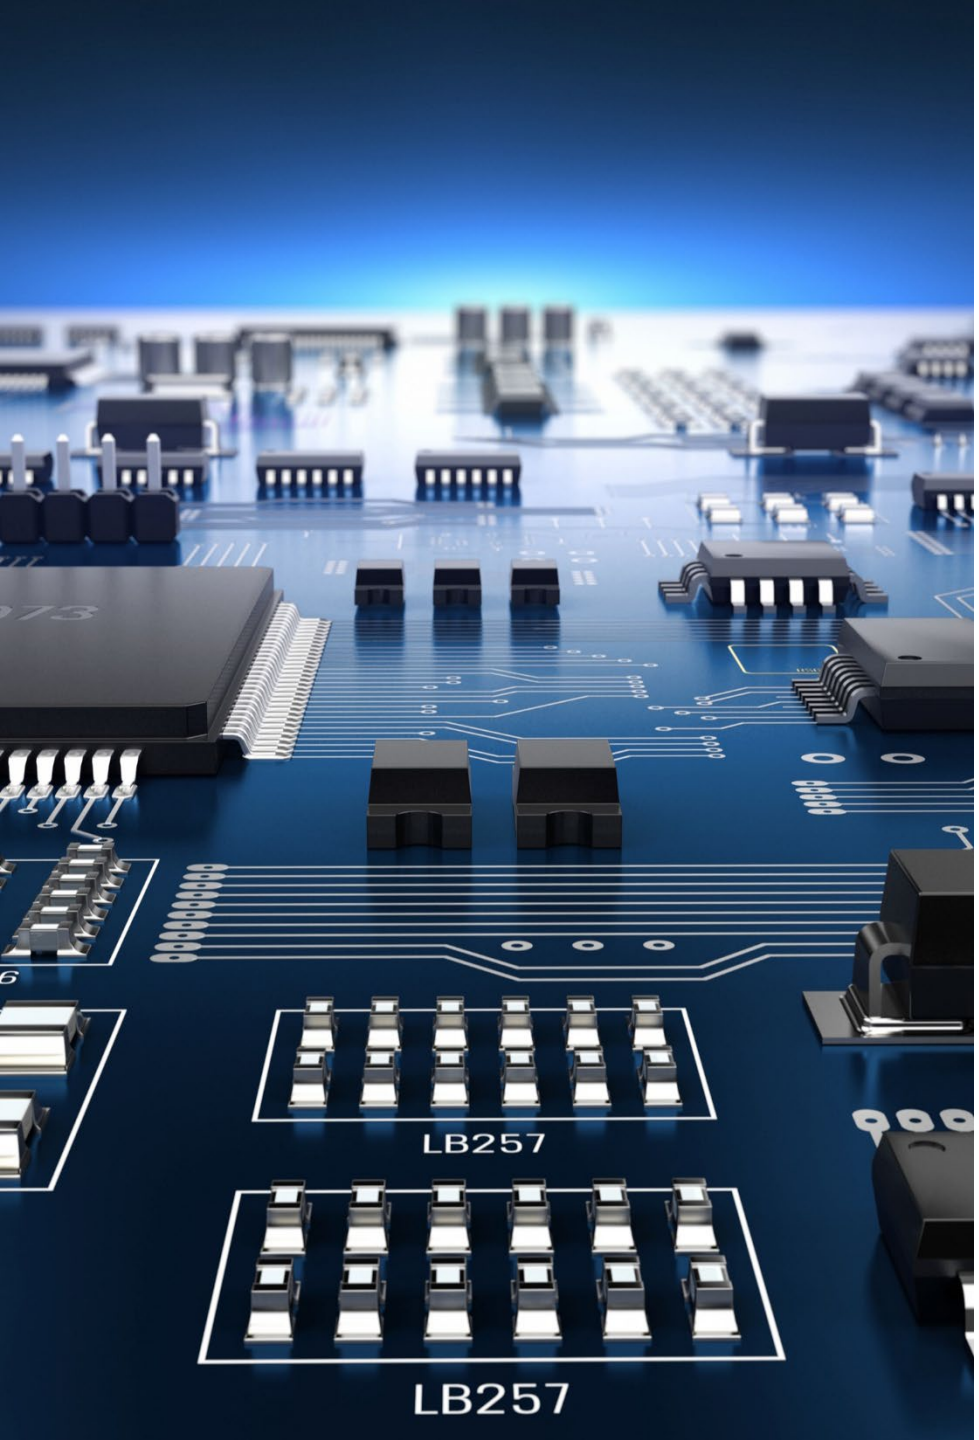

Low overhead

Load balanced

Scalable

Autonomous

Fault tolerant

Key123ABC

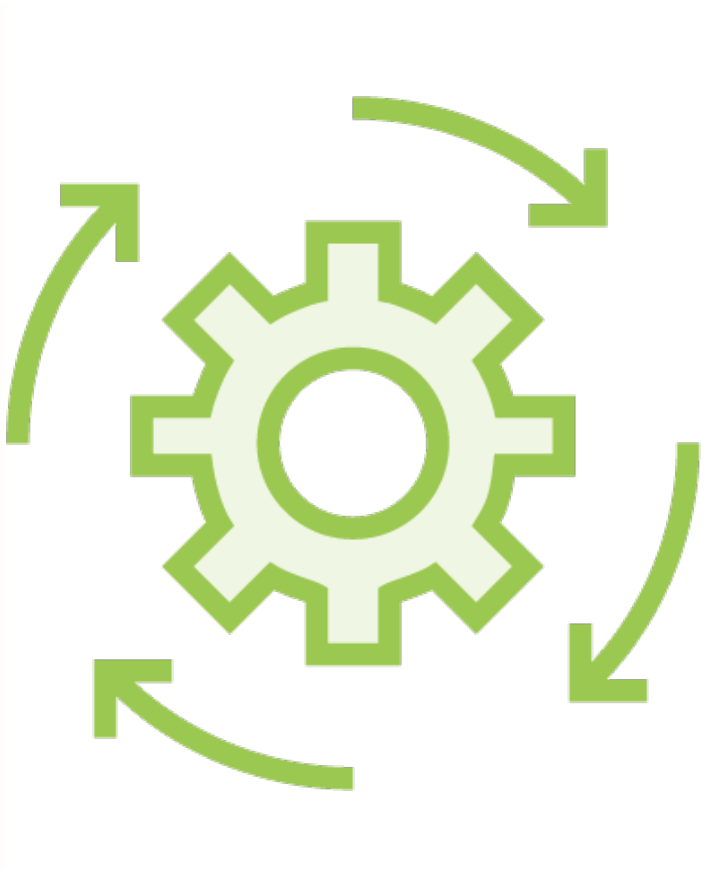

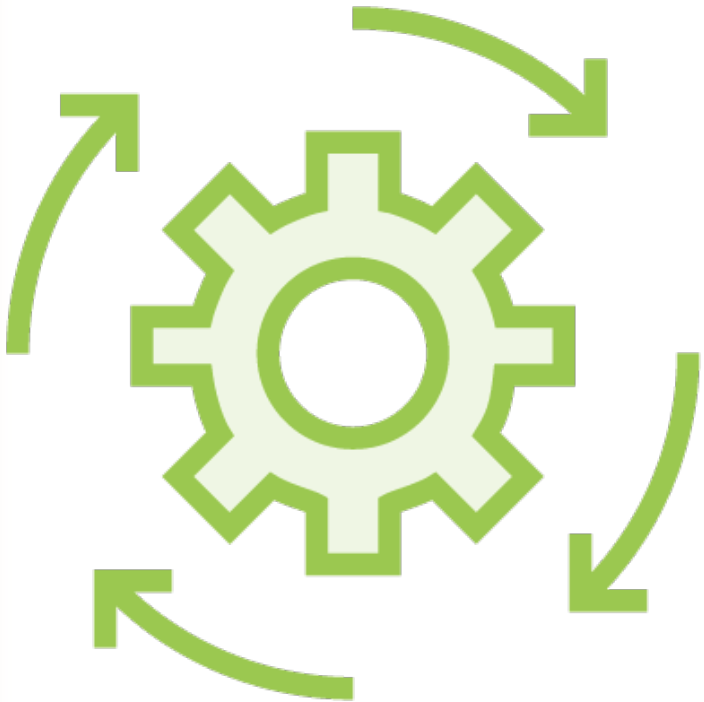

Rendezvous Hashing

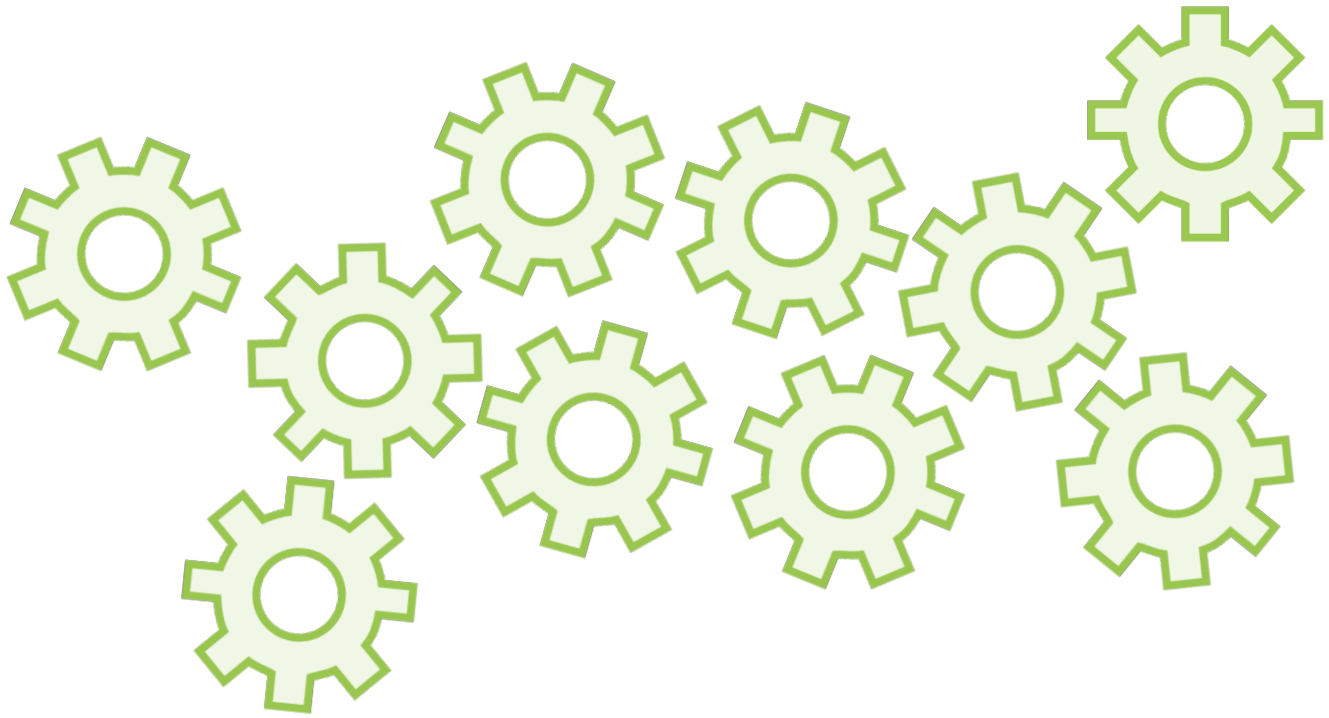

Attributes of the Affinity Function

Low Overhead: We get the node and partition with one call

Fault Tolerant: As the topology changes – the load is rebalanced

Smart Load Balancing: Load is distributed in proportion with capacity

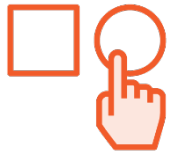

Deterministic: A given key will always result in the same partition / node

Affinity and Collocation

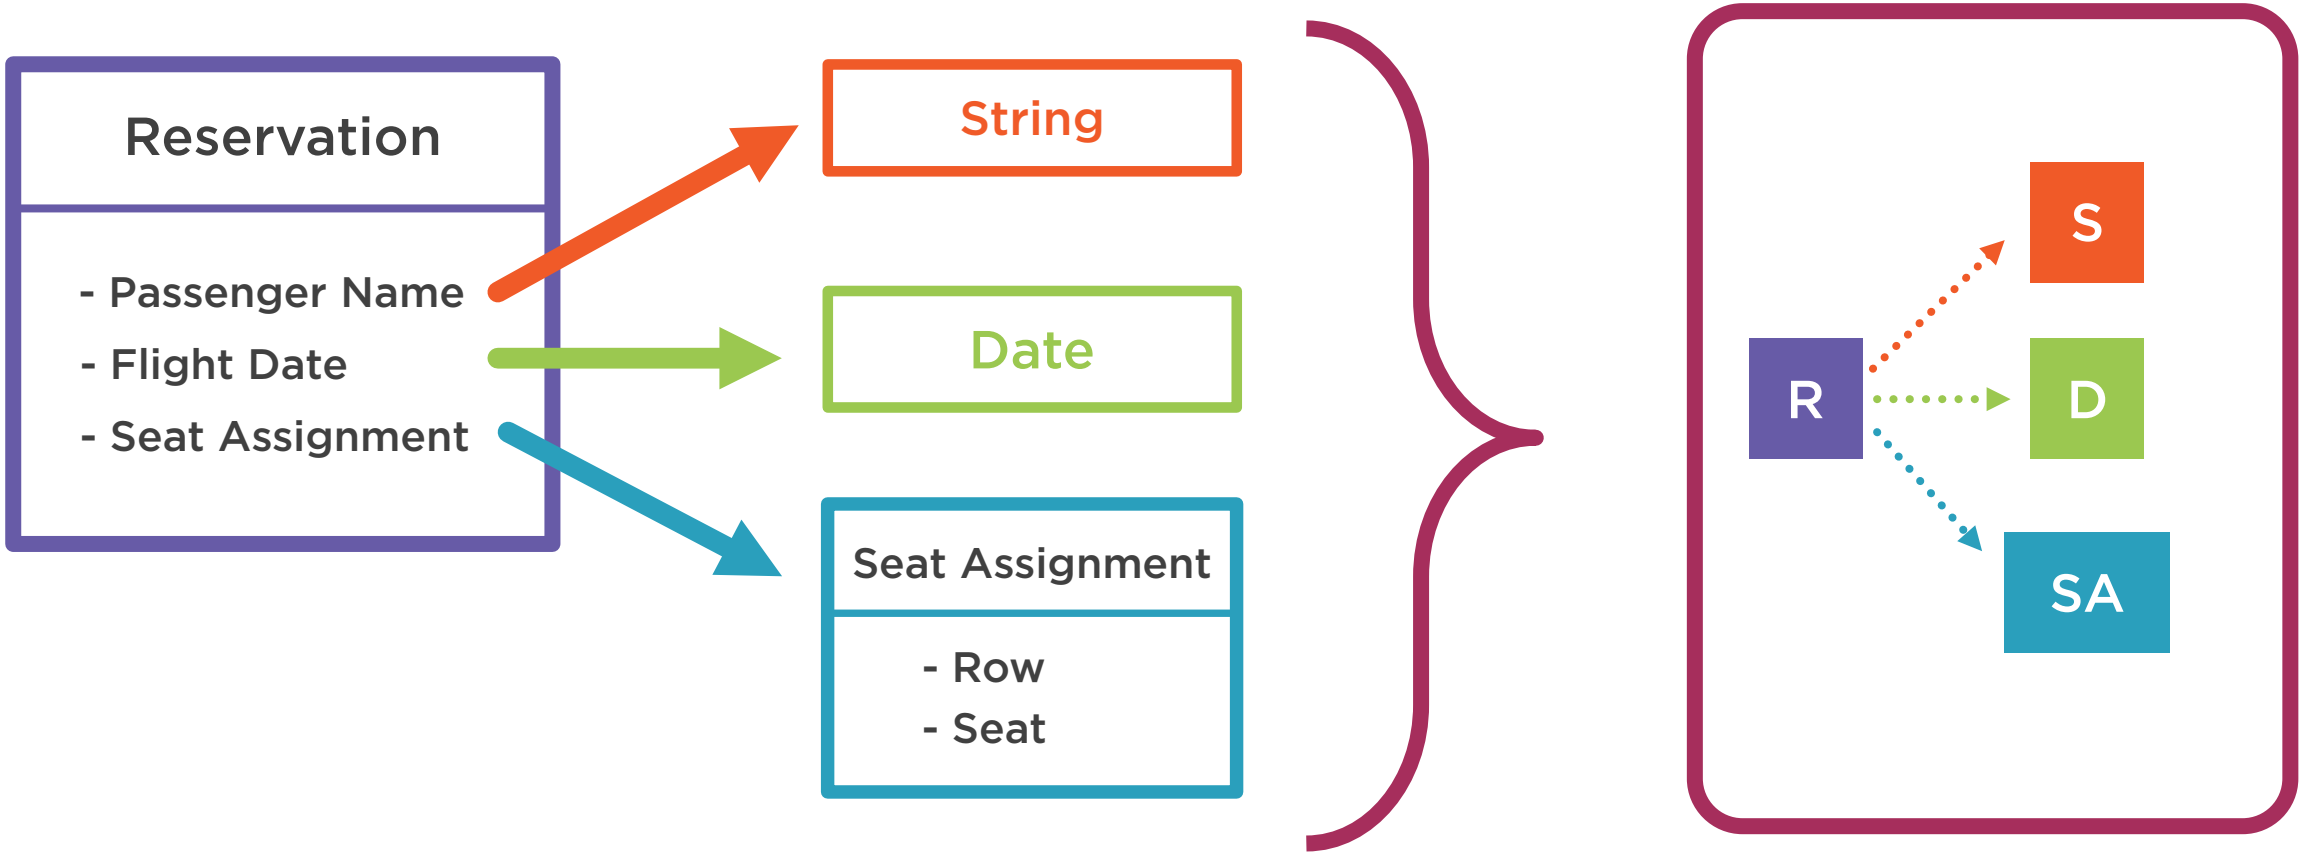
Reservation
- Passenger Name
- Flight Date
- Seat Assignment
String
Date
Seat Assignment
- Row
- Seat
R
S
D
SA

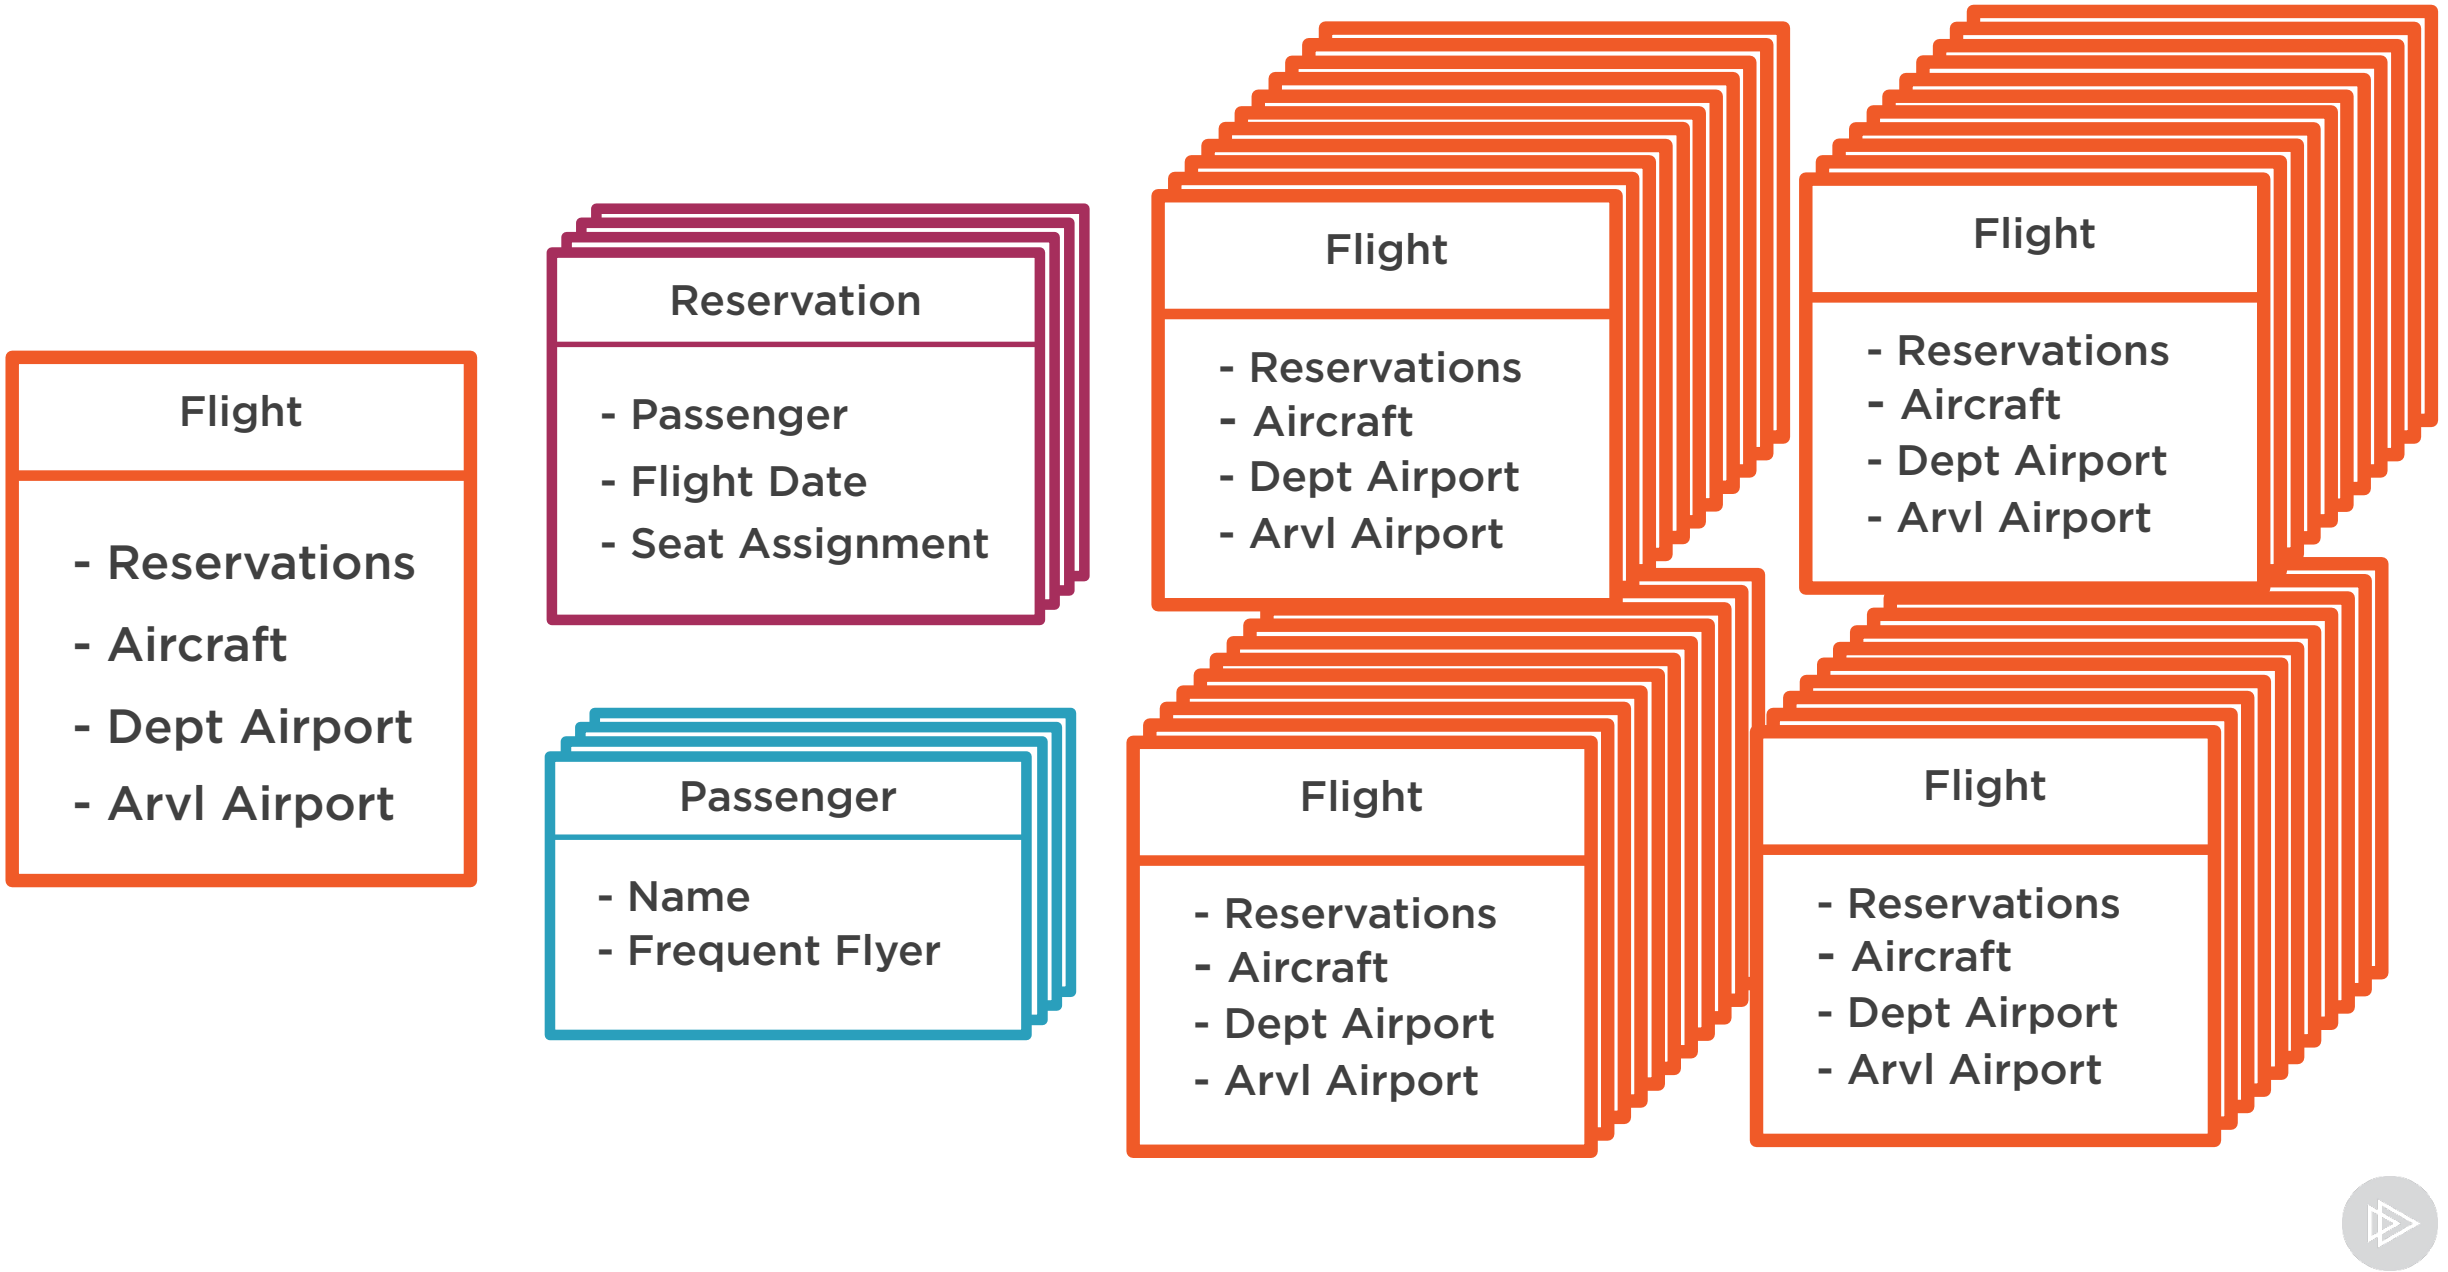
Flight
- Reservations
- Aircraft
- Dept Airport
- Arvl Airport
Reservation
- Passenger
- Flight Date
- Seat Assignment
Passenger
- Name
- Frequent Flyer
Flight
- Reservations
- Aircraft
- Dept Airport
- Arvl Airport
Flight
- Reservations
- Aircraft
- Dept Airport
- Arvl Airport
Flight
- Reservations
- Aircraft
- Dept Airport
- Arvl Airport
Flight
- Reservations
- Aircraft
- Dept Airport
- Arvl Airport

Embedded Objects vs. Referenced Objects

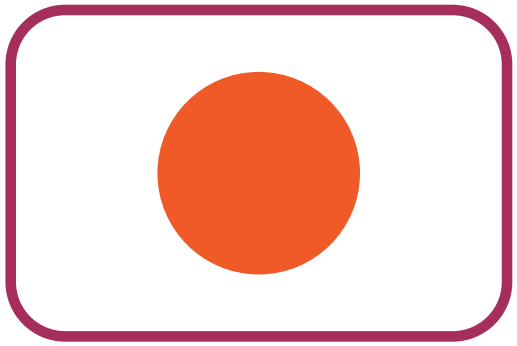

Embedded

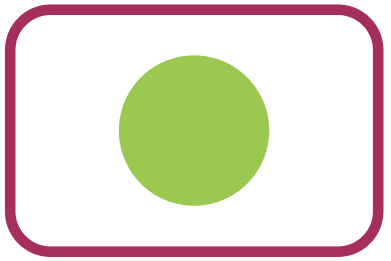

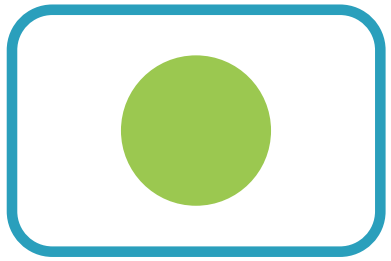

By Reference

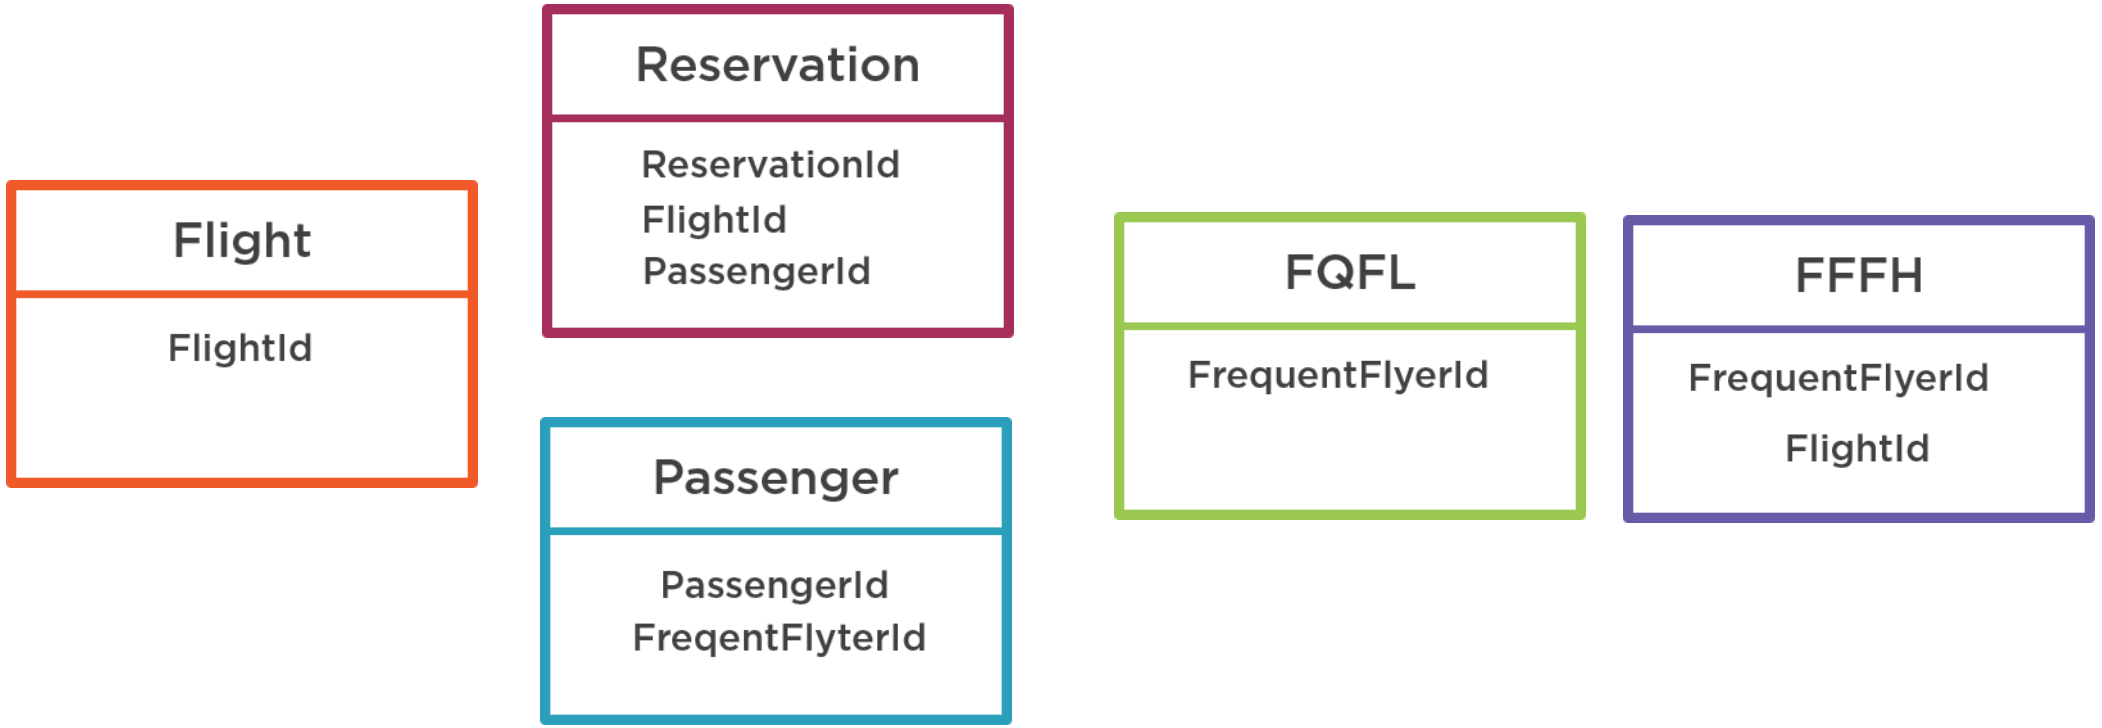
Flight
FlightId
Reservation
ReservationId
FlightId
PassengerId
Passenger
PassengerId
FreqentFlyterId
FQFL
FrequentFlyerId
FFFH
FrequentFlyerId
FlightId

Flight

FlightId

Reservation

ReservationId
FlightId
PassengerId

Passenger

PassengerI
FreqentFlyte

FQFL

FrequentFlyerId

FFFH

FrequentFlyerId
FlightId

863

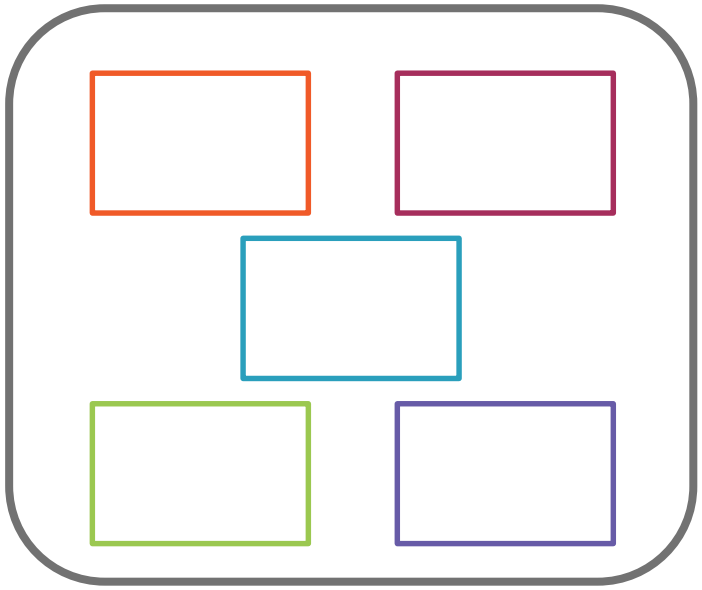

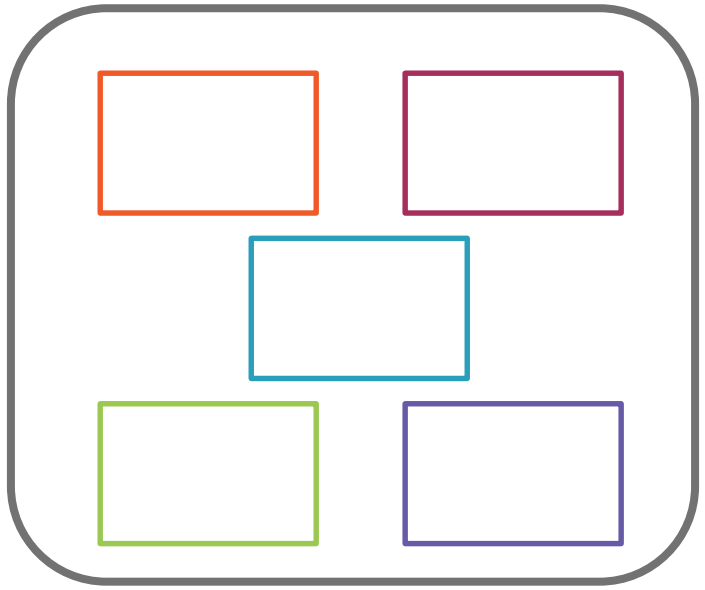

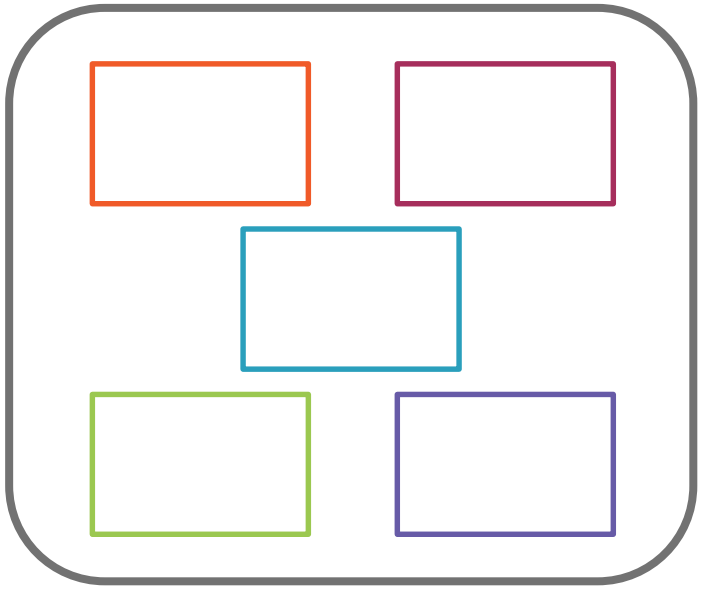

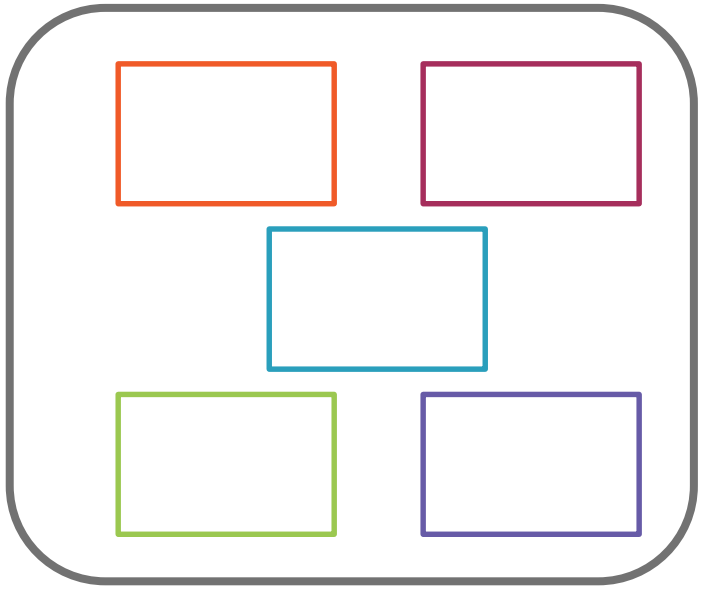

Flight

FlightId

Reservation

Reservation
FlightId
PassengerId

Passenger

PassengerId
FreqentFlyterId

FQFL

FrequentFlyerId

FFFH

FrequentFlyerId
FlightId

1057

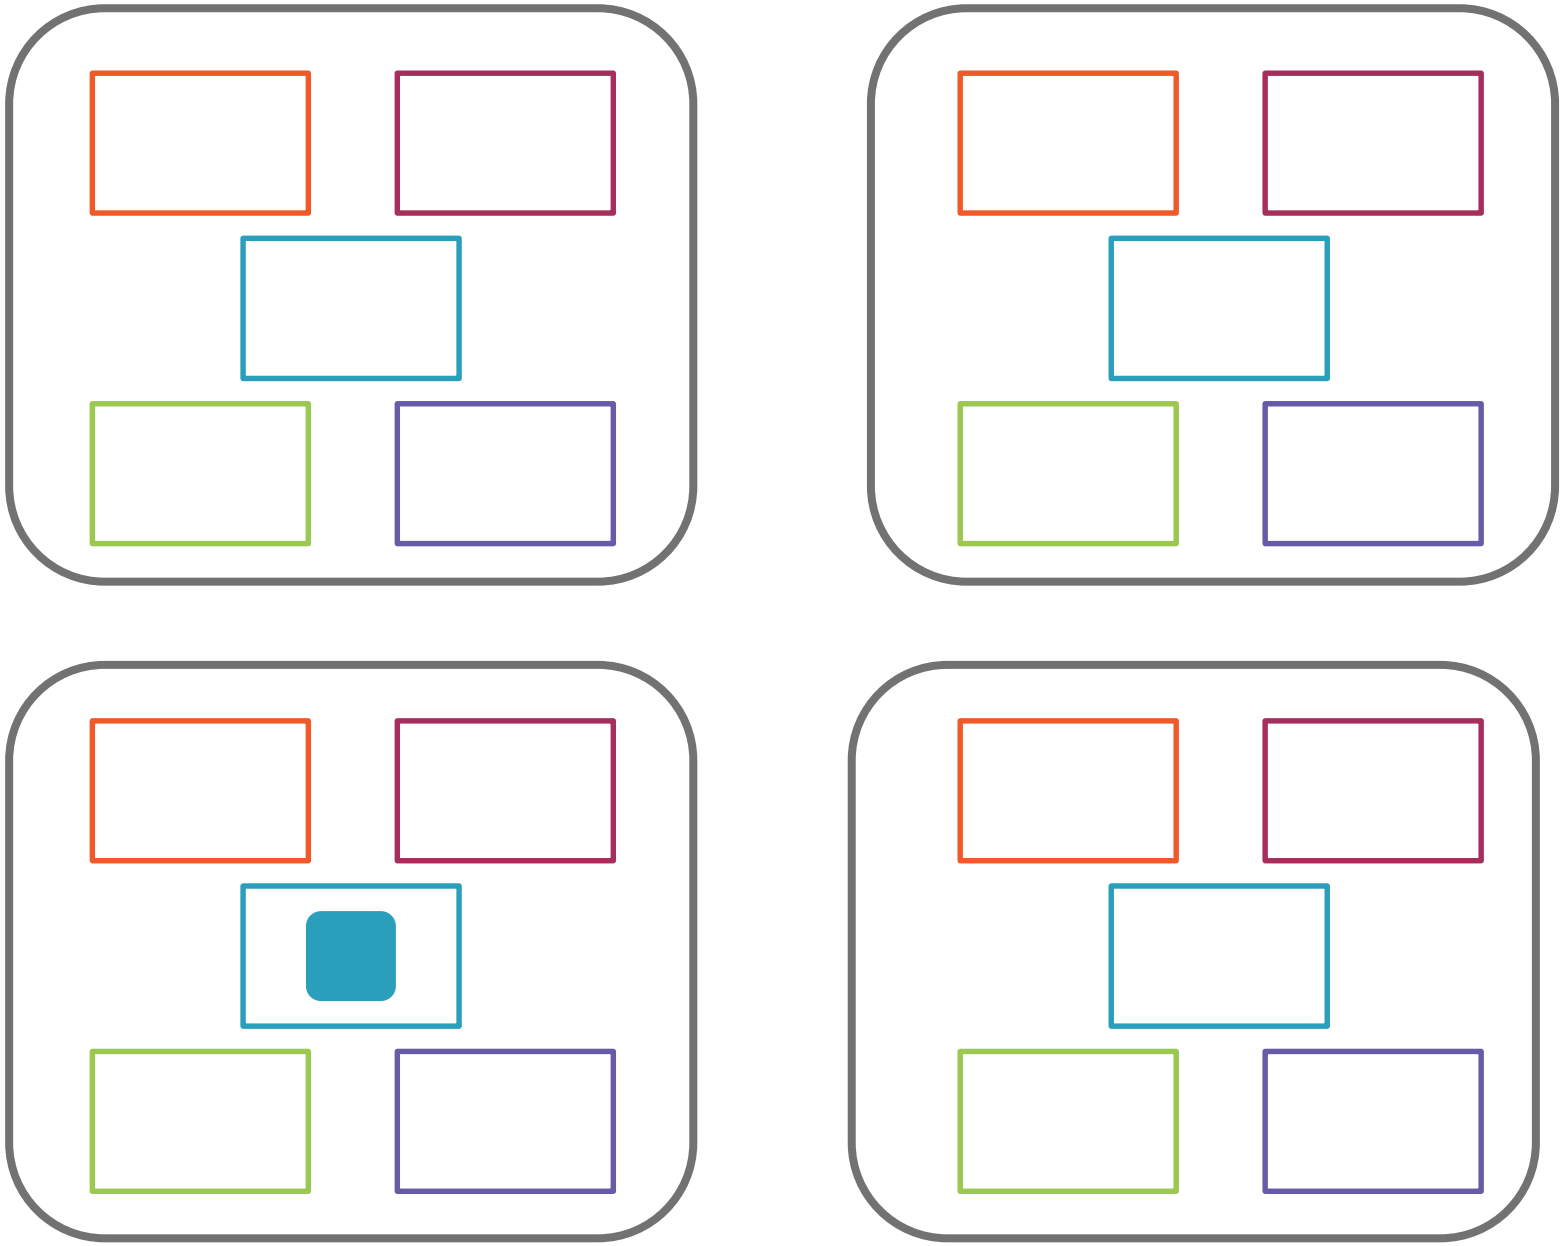

Flight

FlightId

Reservation

ReservationId
FlightId
PassengerId

Passenger

PassengerId
FreqentFlyterId

FQFL

FrequentFlyerId

FFFH

FrequentFlyerId
FlightId

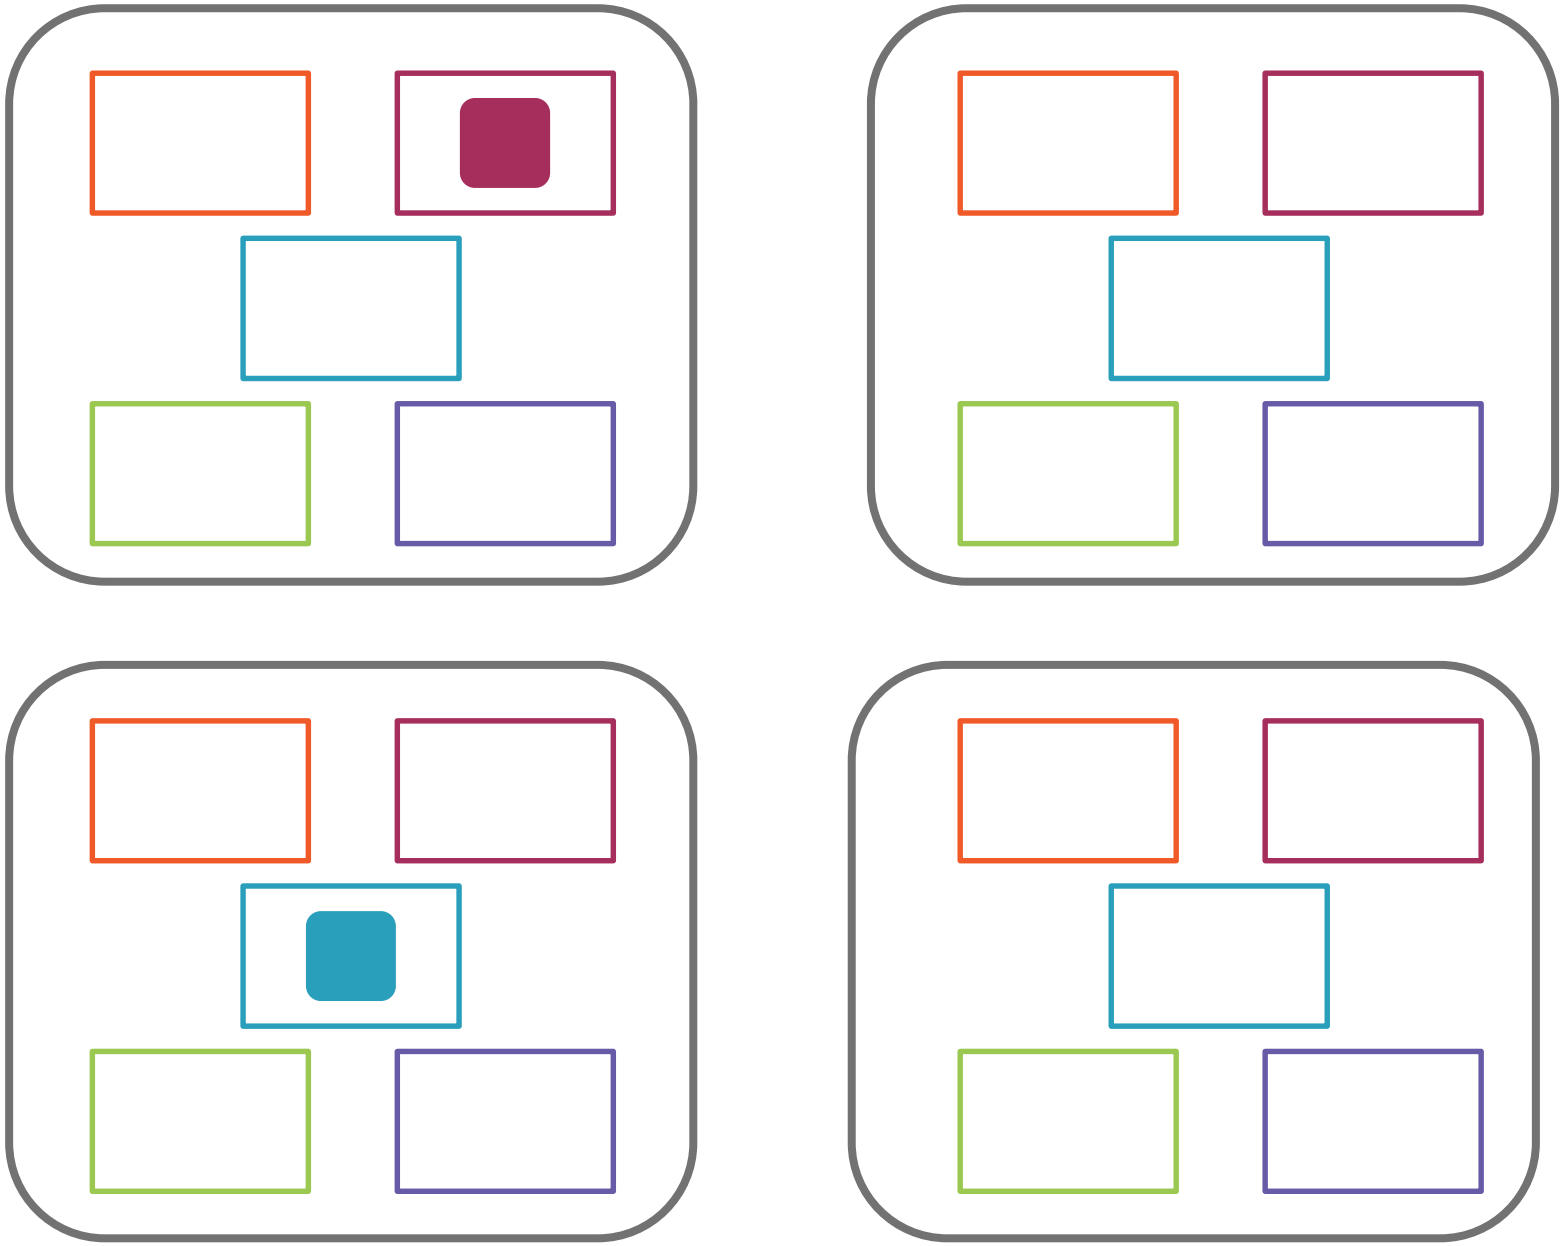

Flight

FlightId

Reservation

ReservationId
FlightId
PassengerId

Passenger

PassengerId
FreqentFlyterId

FQFL

FrequentFlyerId

FFFH

FrequentFlyerId
FlightId

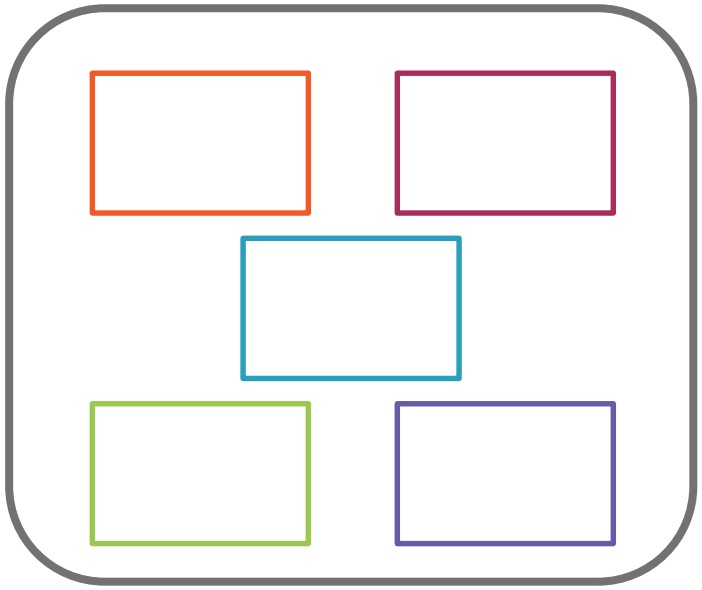

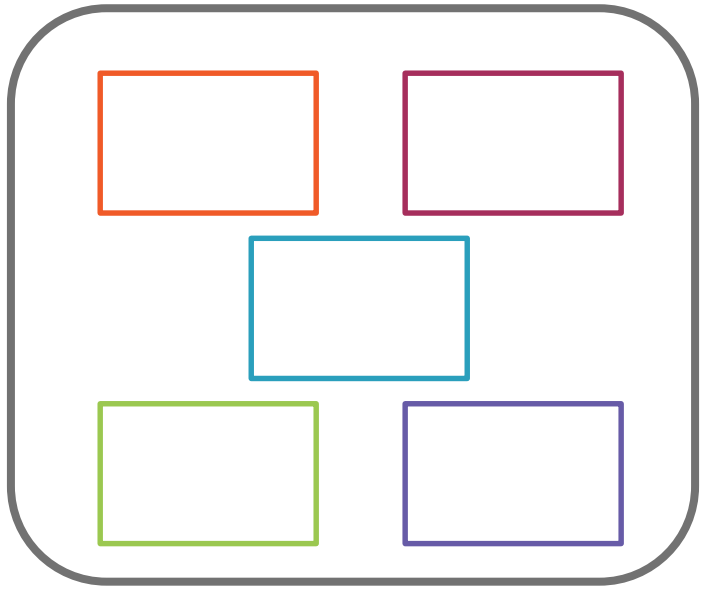

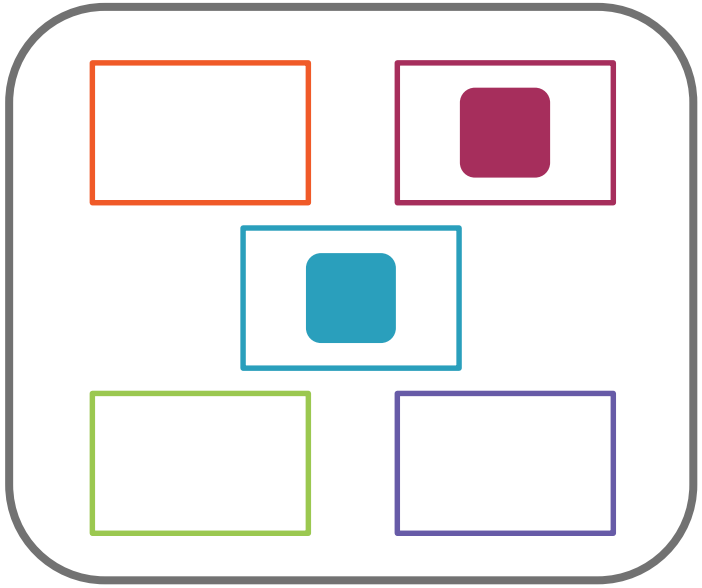

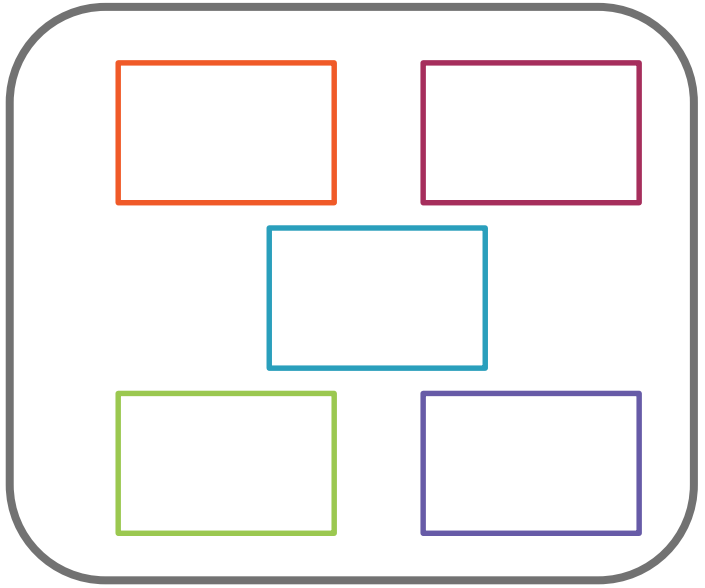

Affinity Key Construction

```
public class Reservation
{
     private int reservationId;
     private int passengerId;

     public Reservation(int passengerId, int reservationId) {

          …
     }

     public AffinityKey getKey() {
          return new AffinityKey(passengerId, reservationId);
     }

     …
}
```

Flight

FlightId

Reservation

Reservation
FlightId
PassengerId

Passenger

PassengerId
FreqentFlyterId

FQFL

FrequentFlyerId

FFFH

FrequentFlyerId
FlightId

1057

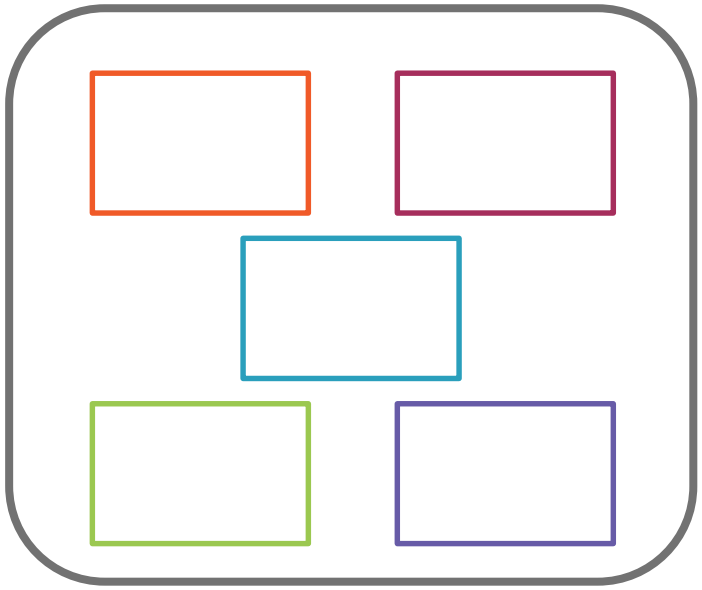

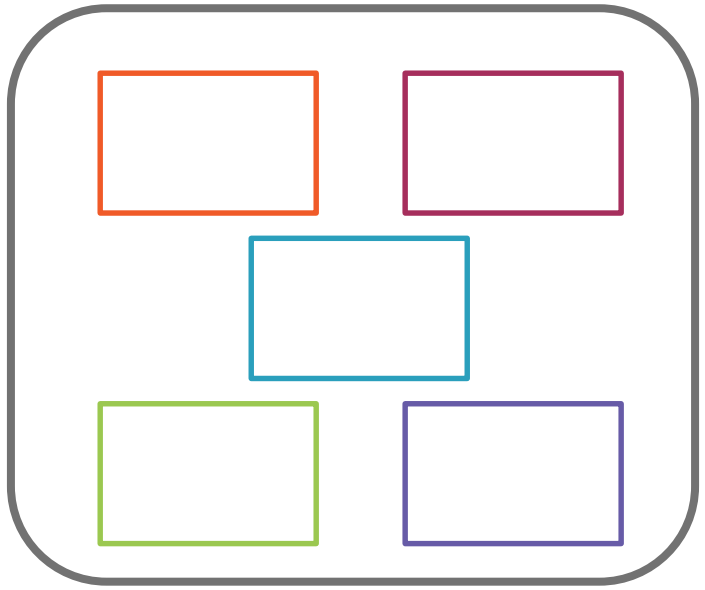

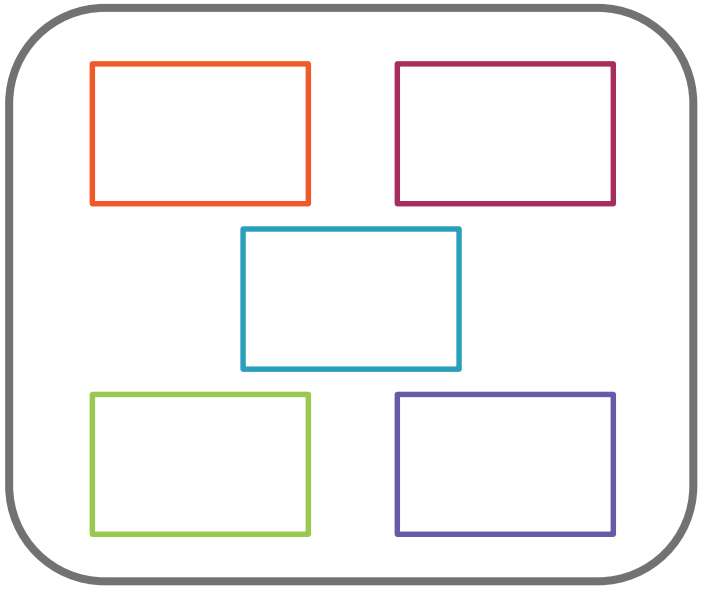

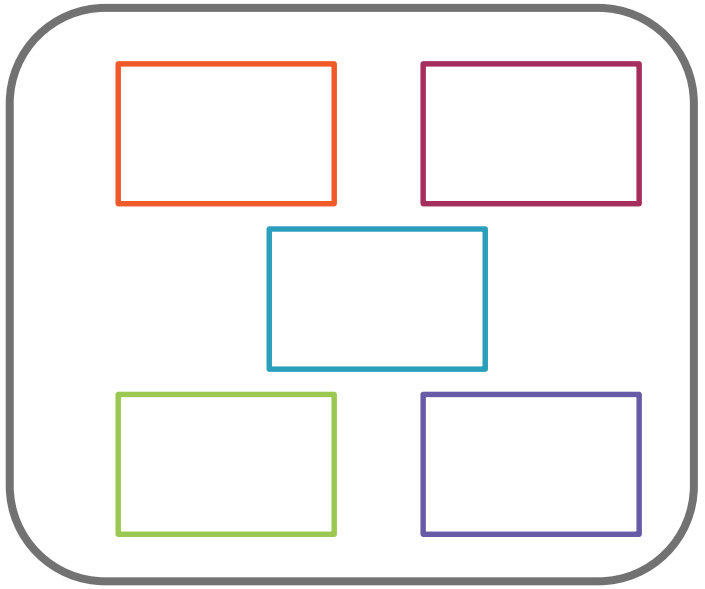

Flight

FlightId

Reservation

ReservationId
FlightId
PassengerId

Passenger

PassengerId
FreqentFlyterId

FQFL

FrequentFlyerId

FFFH

FrequentFlyerId
FlightId

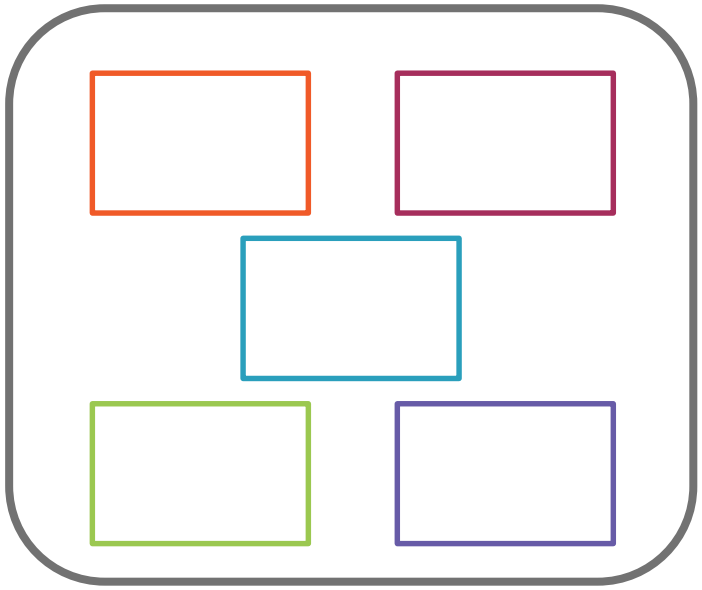

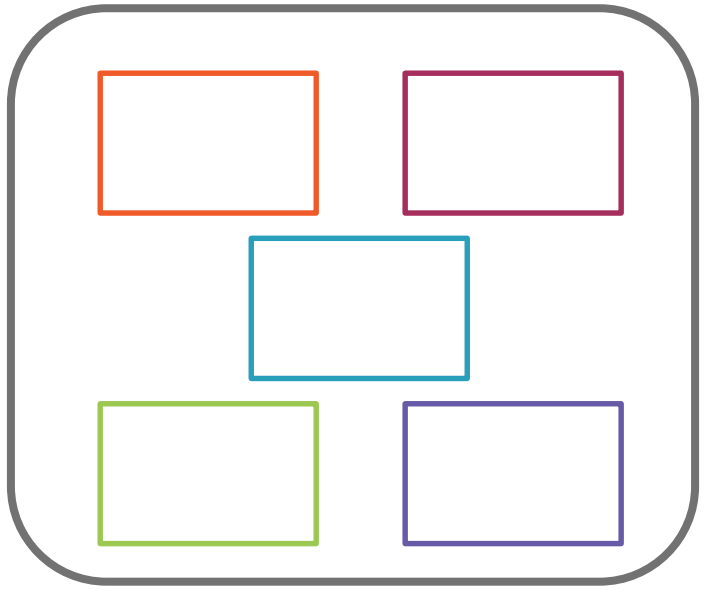

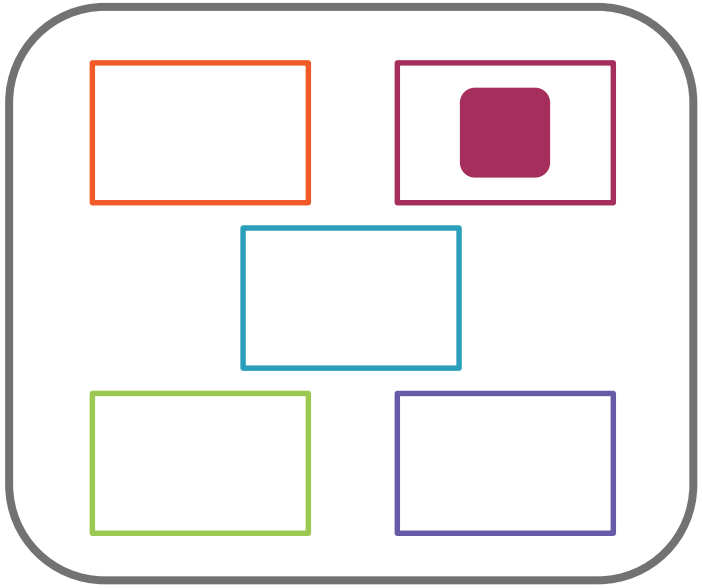

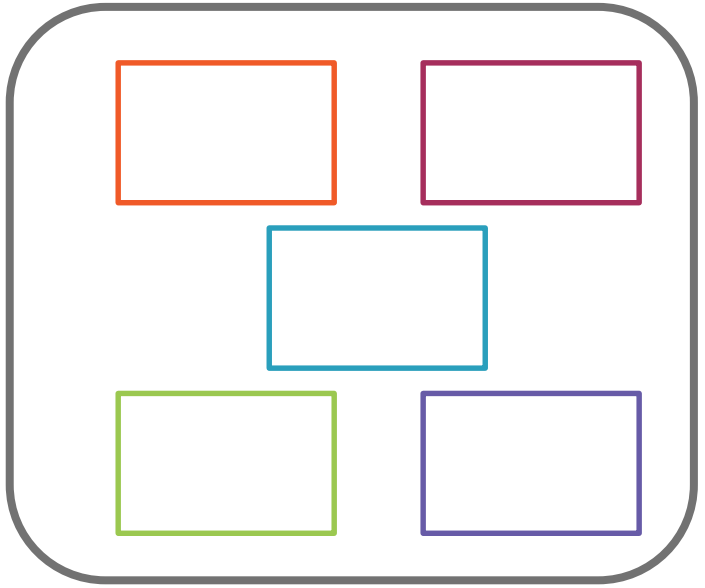

Flight

FlightId

Reservation

ReservationId
FlightId
PassengerId

Passenger

PassengerI
FreqentFlyte

FQFL

FrequentFlyerId

FFFH

FrequentFlyerId
FlightId

863, 1057

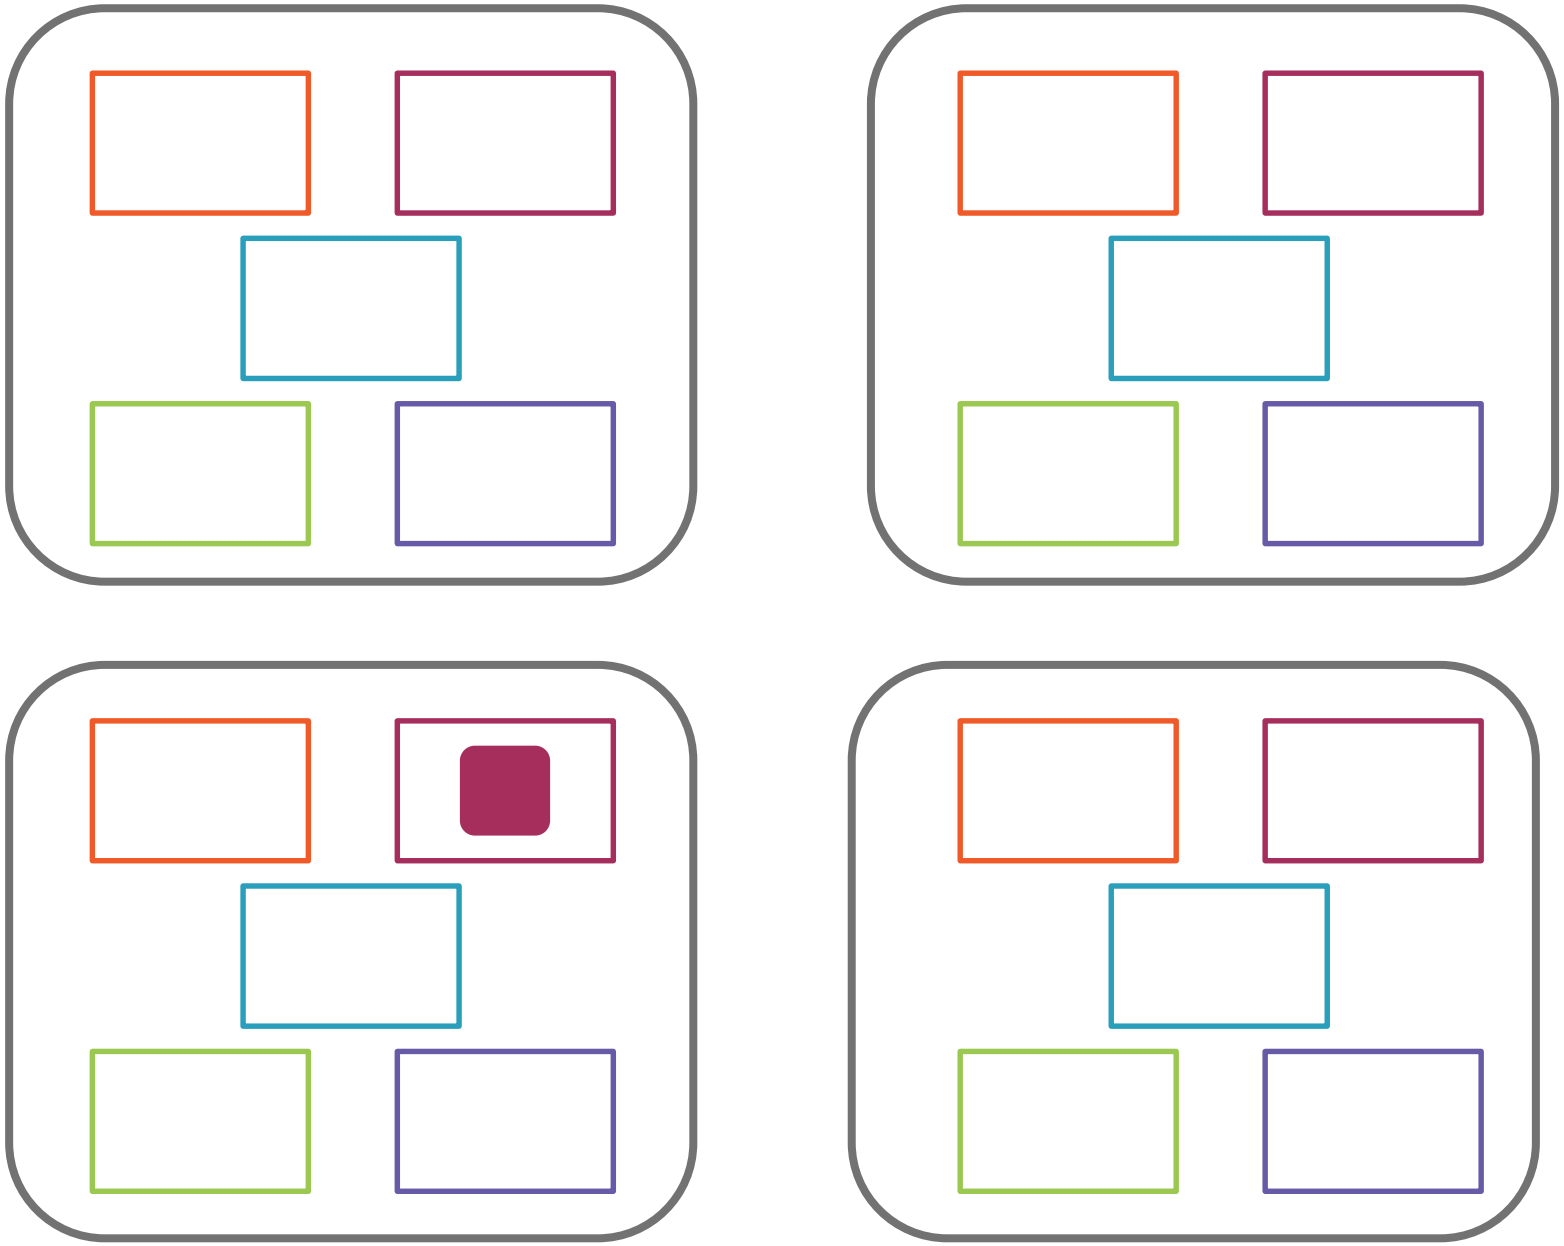

Flight

FlightId

Reservation

ReservationId
FlightId
PassengerId

Passenger

PassengerId
FreqentFlyterId

FQFL

FrequentFlyerId

FFFH

FrequentFlyerId
FlightId

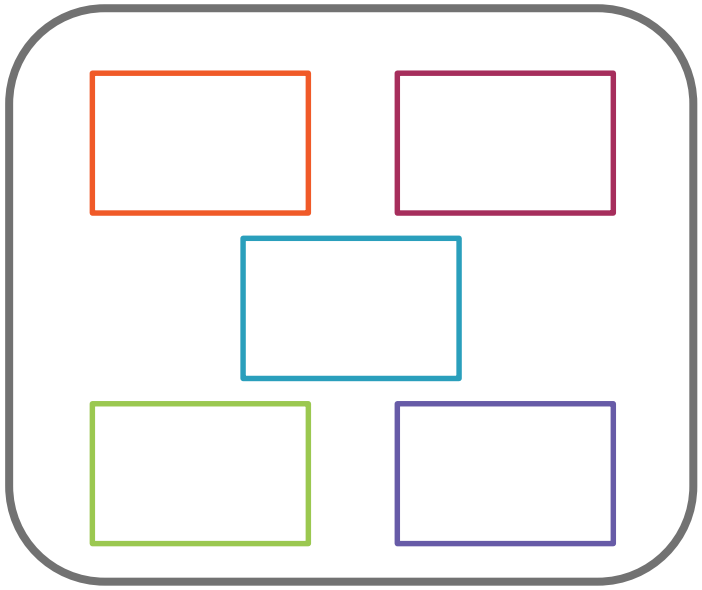

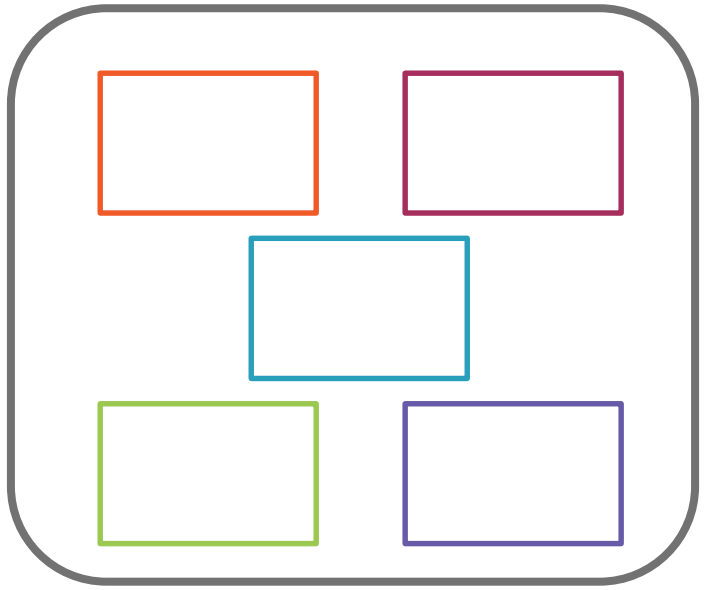

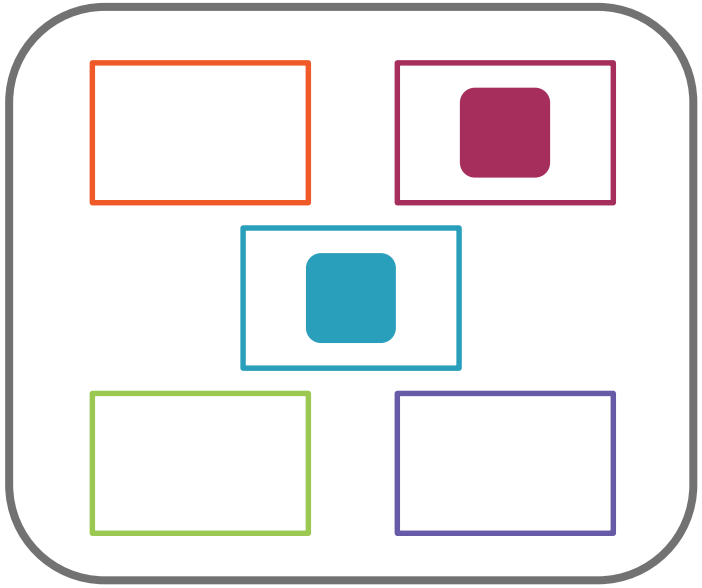

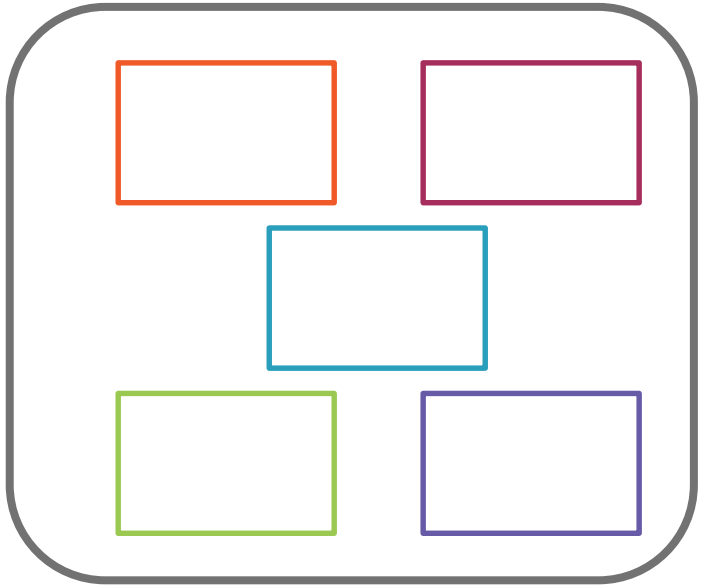

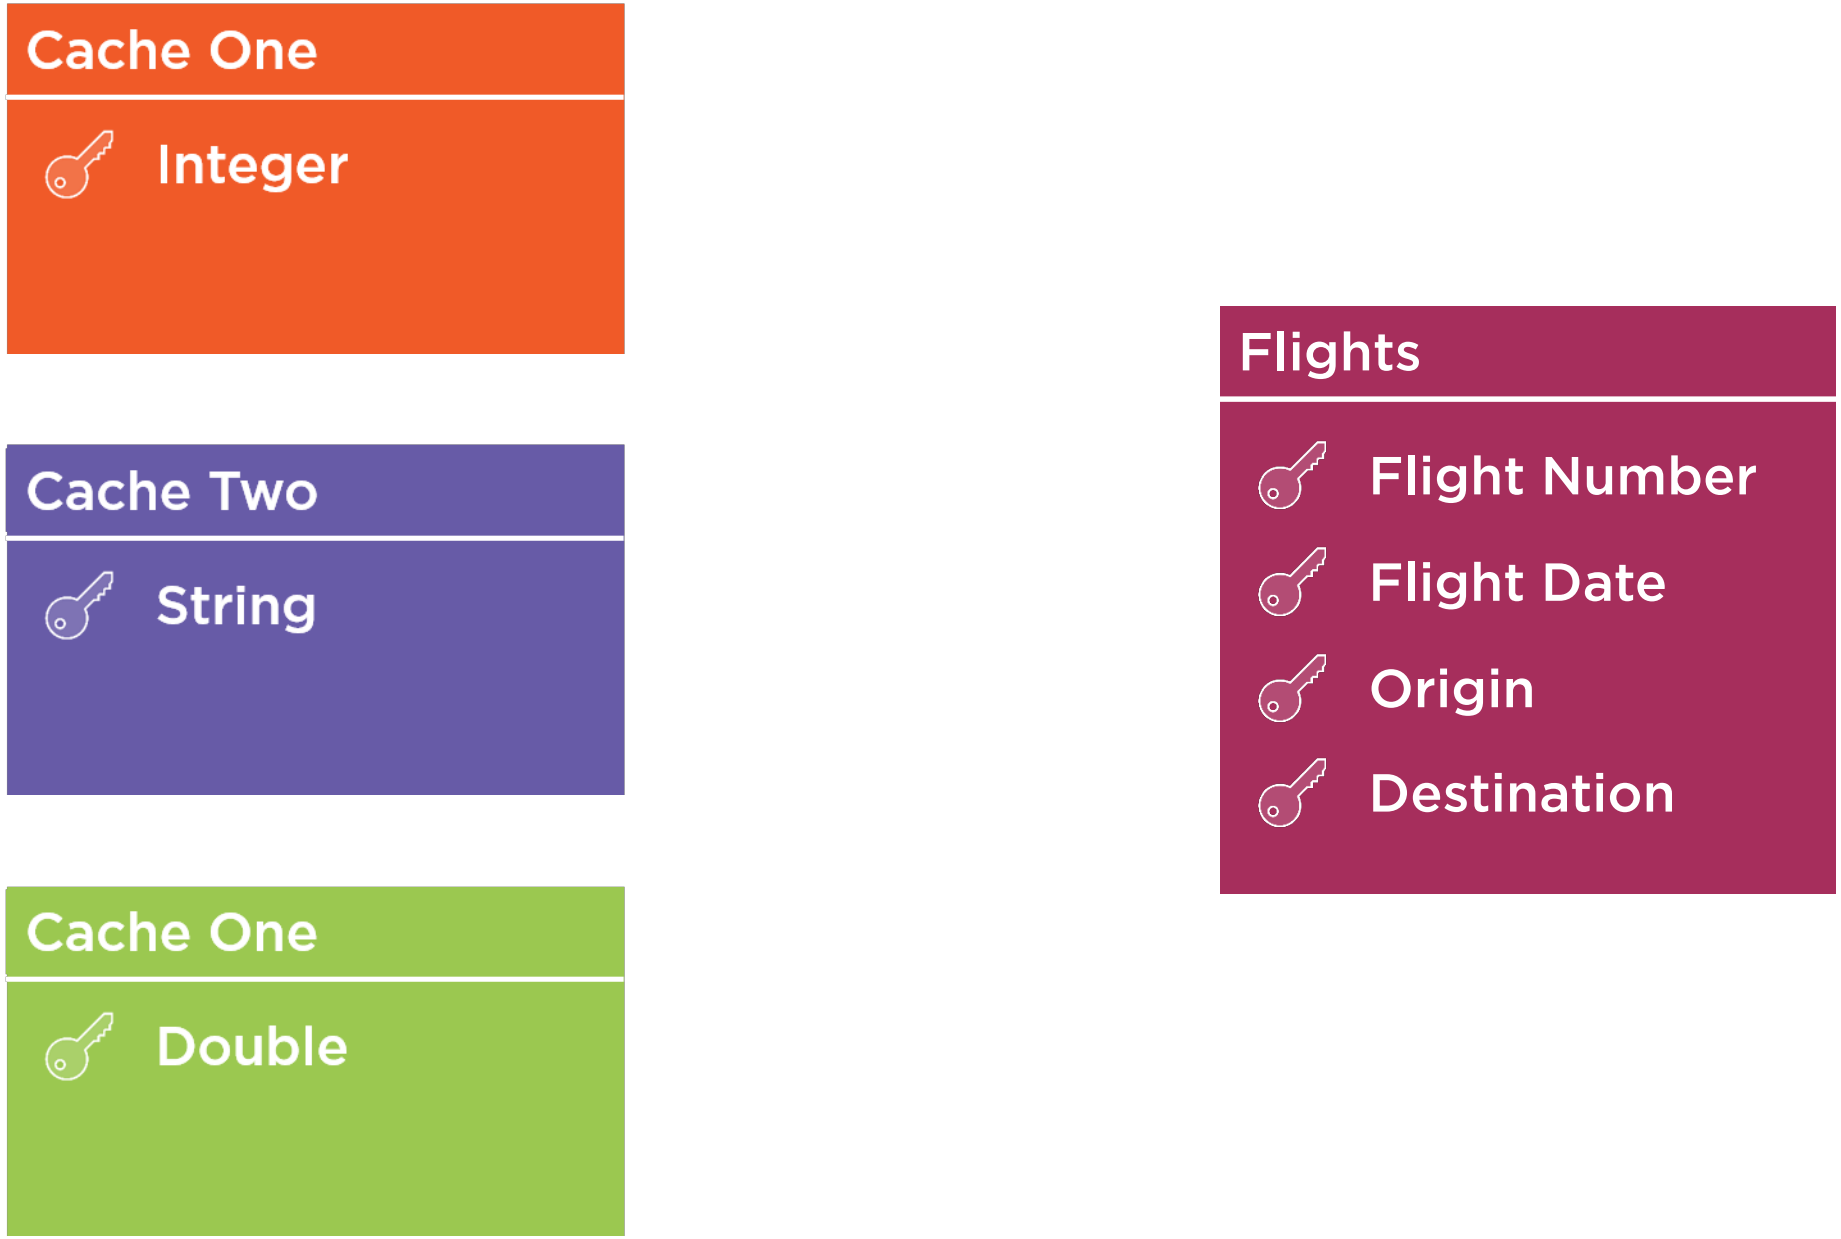
Cache One
Integer
Cache Two
String
Cache One
Double
Flights
Flight Number
Flight Date
Origin
Destination


Flights
Flight Number
Flight Date
Origin
Destination


```
de
boolean equals(Object o)

de
int hashCode()
```

0x817FA2

0x817FA2

Data Queries

QueryCursor = IgniteCache.query(Query)

Extensions of the Abstract Query Class

SqlQuery

SqlFieldsQuery

Extensions of the Abstract Query Class

Wrapper for a SQL statement

QueryCursor<List<?>>

SqlFieldsQuery

Extensions of the Abstract Query Class

SqlQuery

SqlFieldsQuery

Extensions of the Abstract Query Class

SqlQuery

Pass type and WHERE clause

QueryCursor<Entry<Key, Value>>

Flight

FlightId

Reservation

ReservationId
FlightId
PassengerId

Passenger

PassengerId
FreqentFlyterId

FQFL

FrequentFlyerId

FFFH

FrequentFlyerId
FlightId

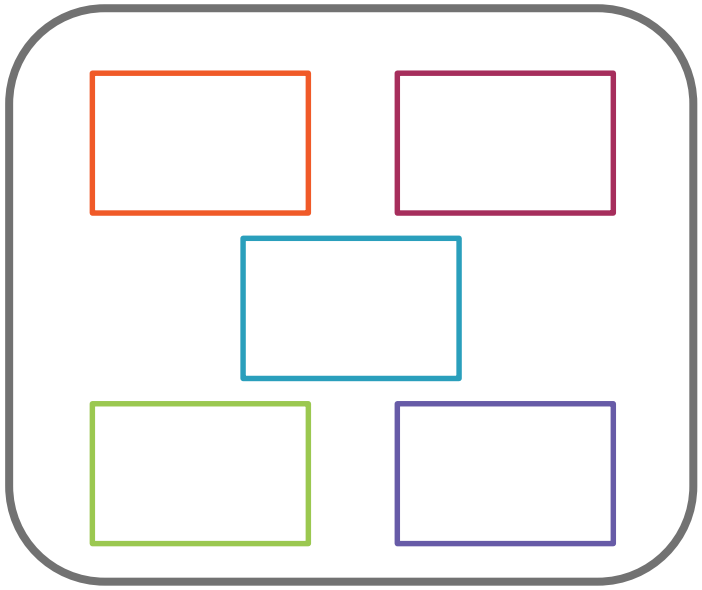

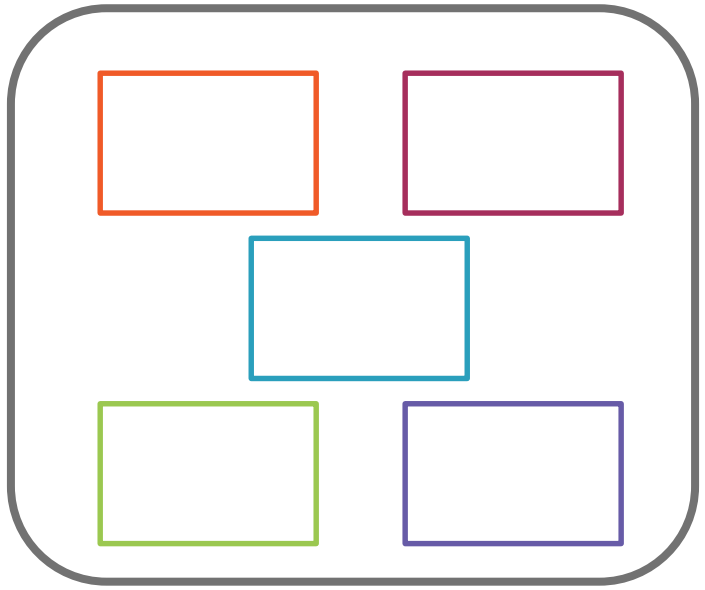

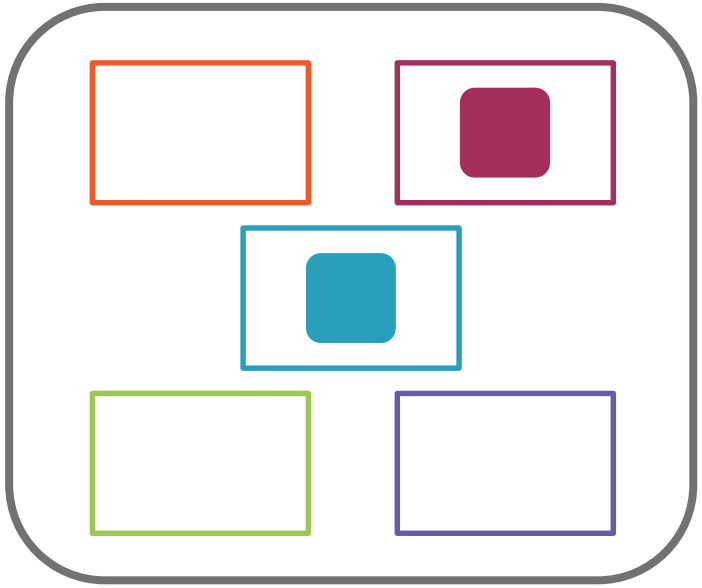

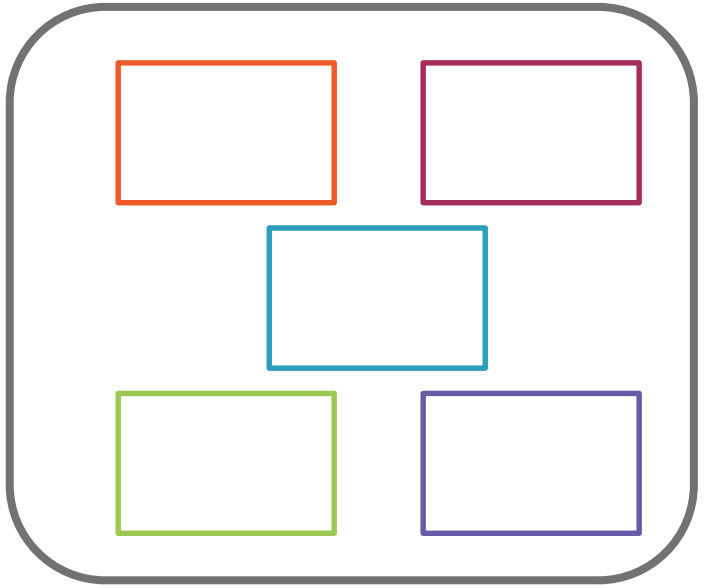

Continuous Query

Initial Query

Remote Filter

Local Listener

The initial query is run immediately.

The remote filter is the ongoing cache listener.

The local listener receives the data that is found on the remote nodes.

Receive 1 notification

All nodes will flush queues

Cache counters sent with data

Client sends confirmation

Bit of code that runs on the server

Sends the data that's changed

Locks the entry while it's updating

The SQL API & DML

Use SqlFieldsQuery

```
SqlQuery(Class<?> type, String sql)
```

Apache Ignite SqlQuery class

Only returns data from a cache. Cannot do any DML.

Use SqlFieldsQuery

Subquery must have all data collocated

Prefer JCache API methods

- Bulk methods
  - putAll() / getAll()
- Atomic methods
  - putIfAbsent()
  - getAndPutIfAbsent()

DML methods return rows affected

SQL API transaction support coming soon

Writing Data to a Cache

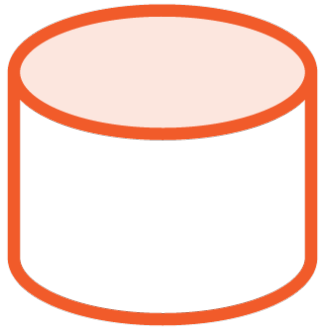

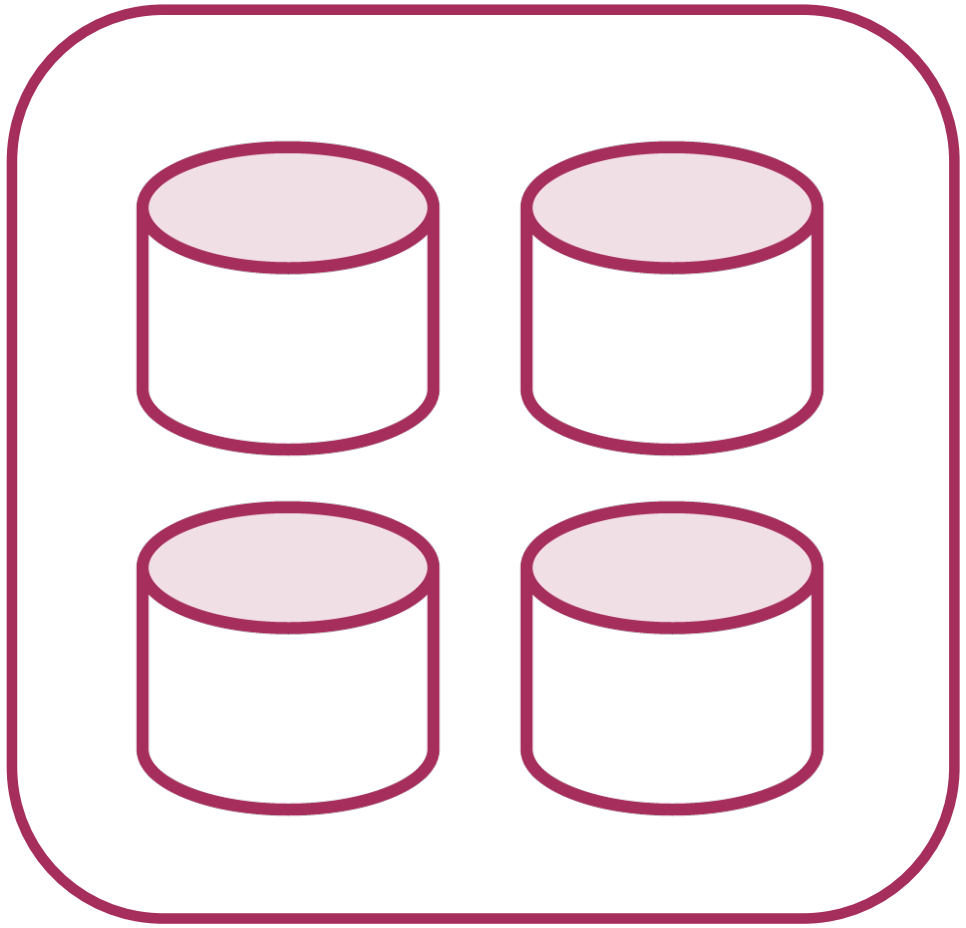

Writing Data to a Cache

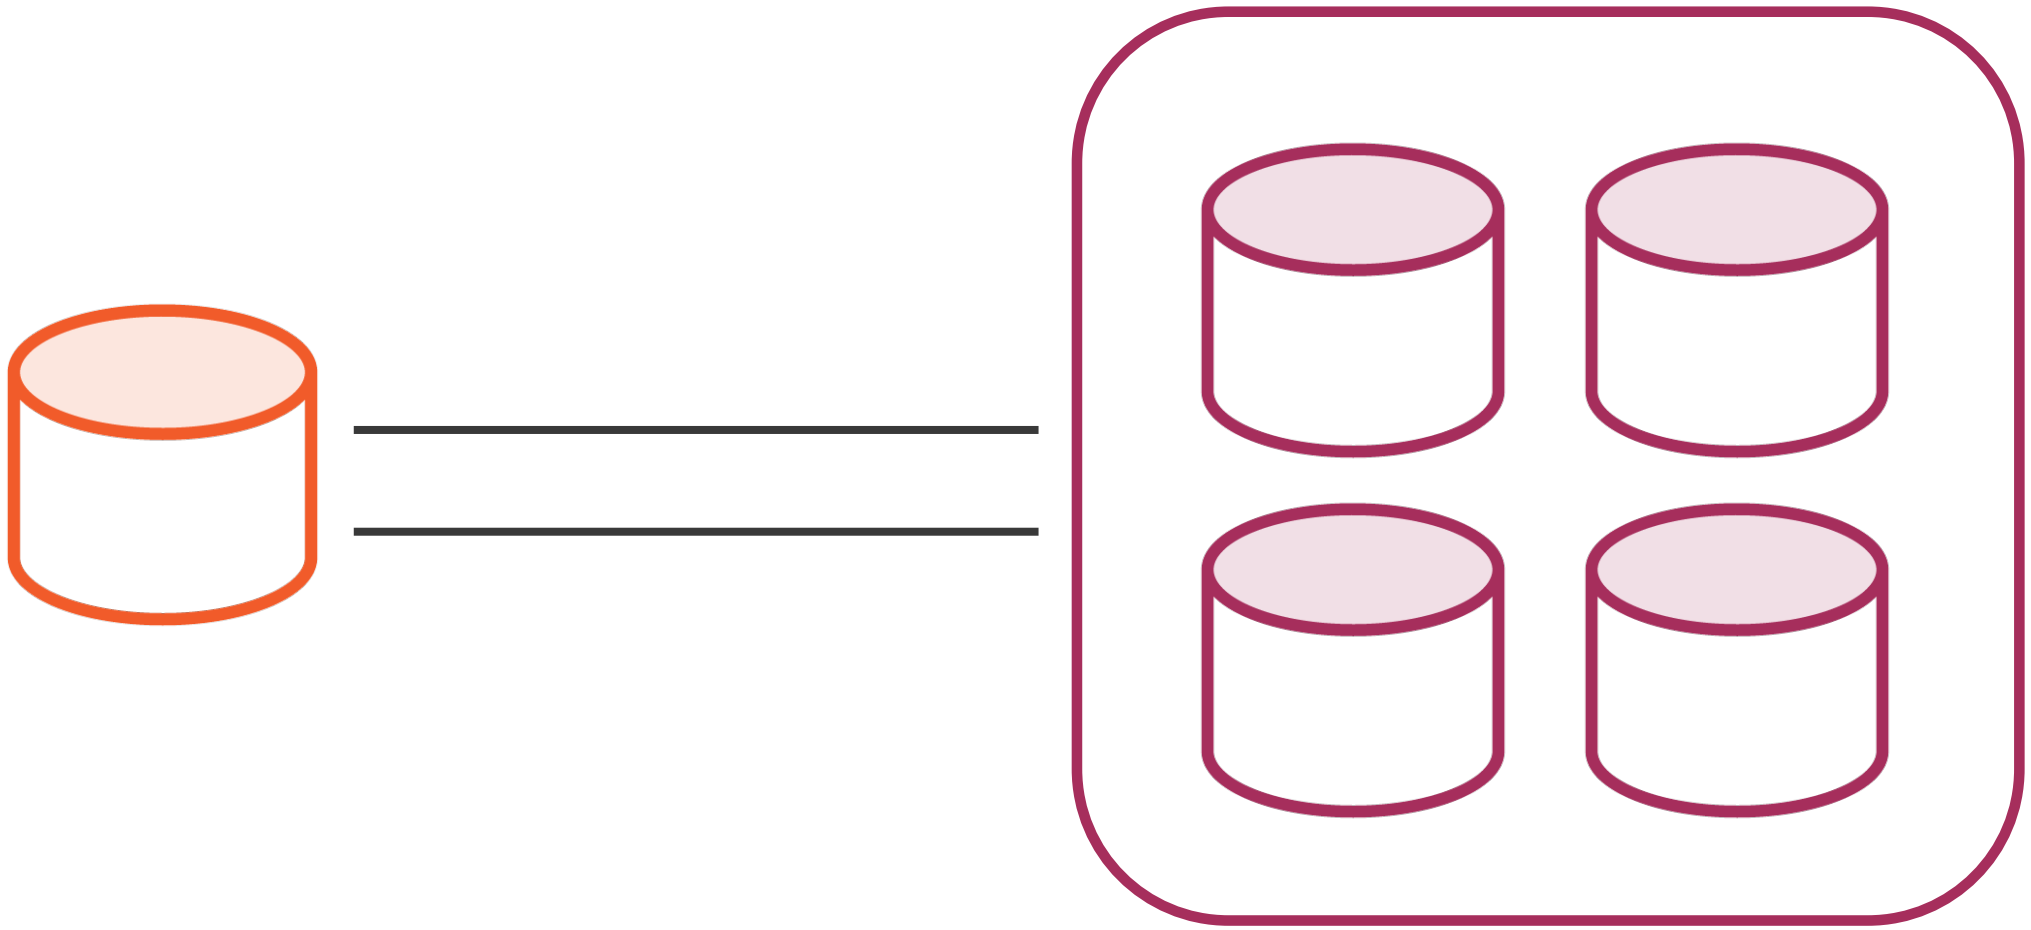

Writing Data to a Cache

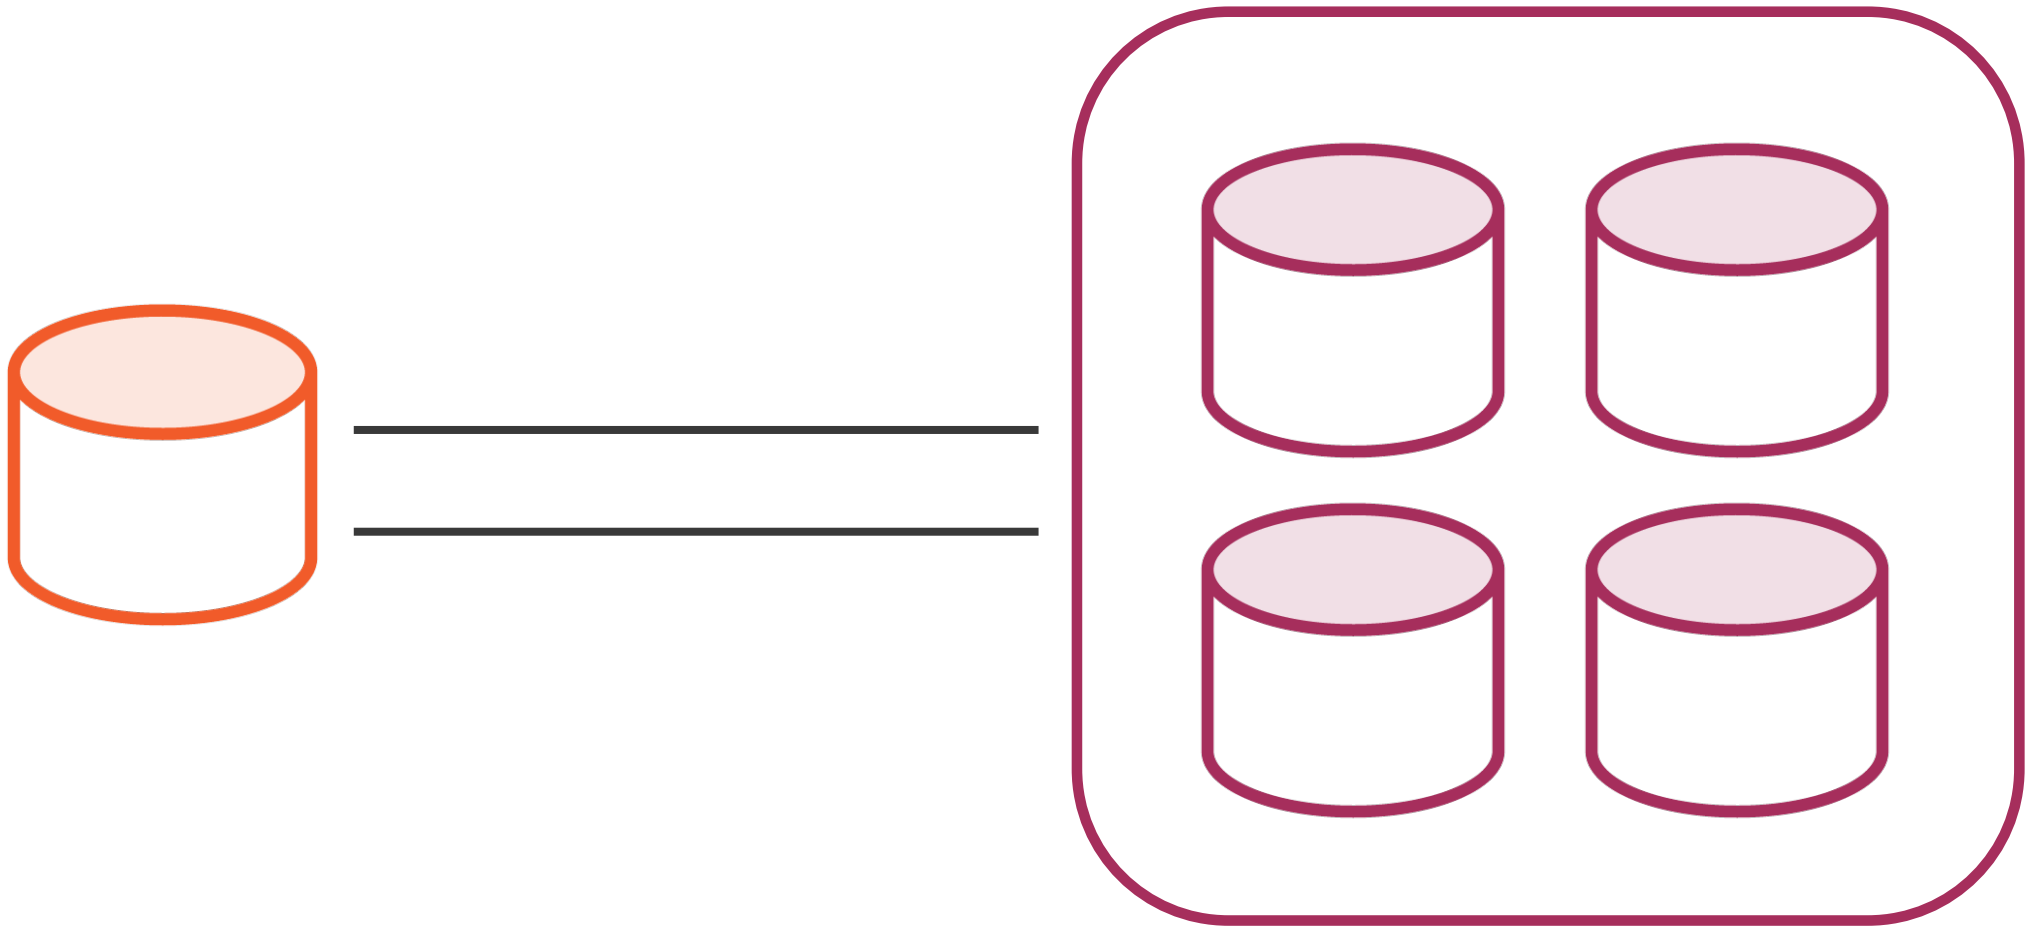

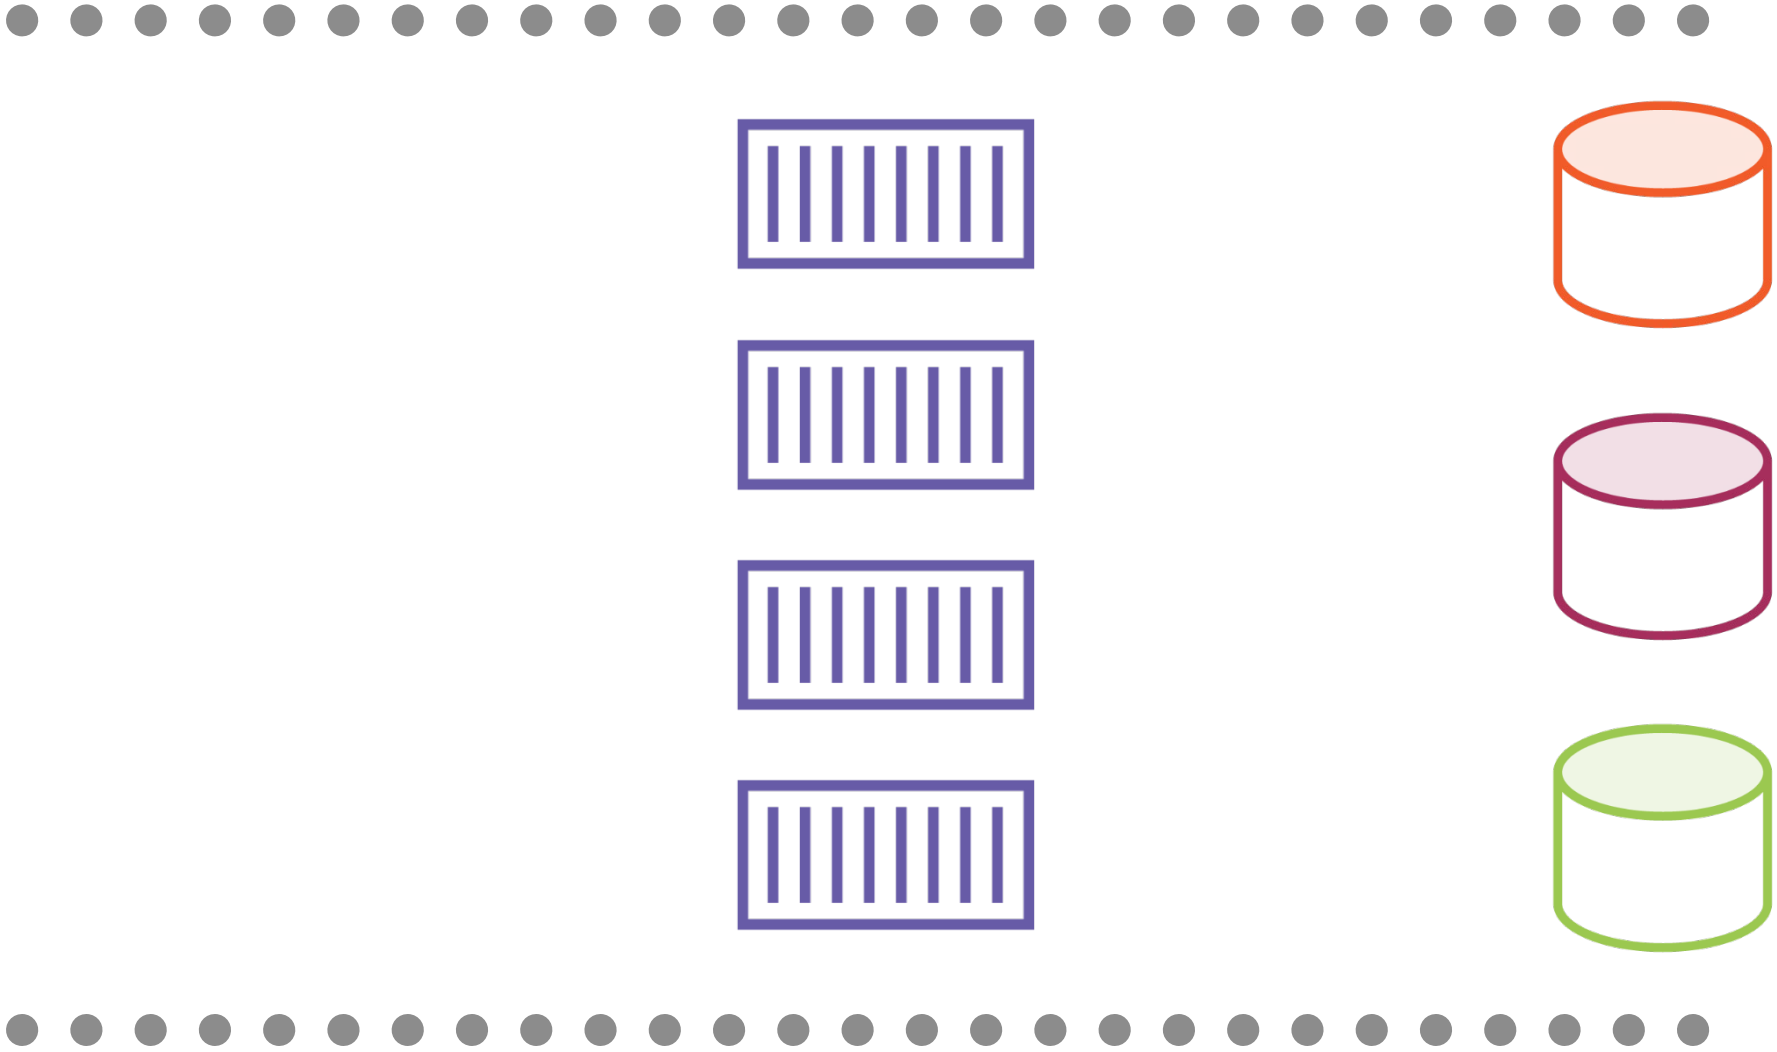

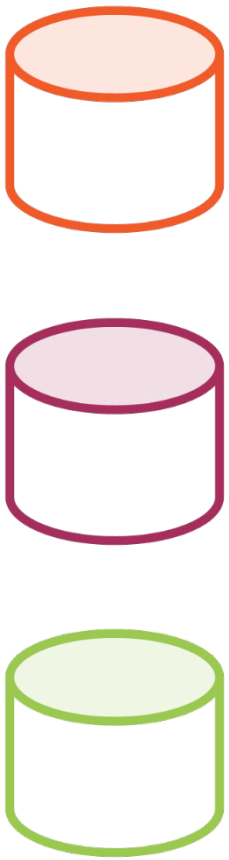

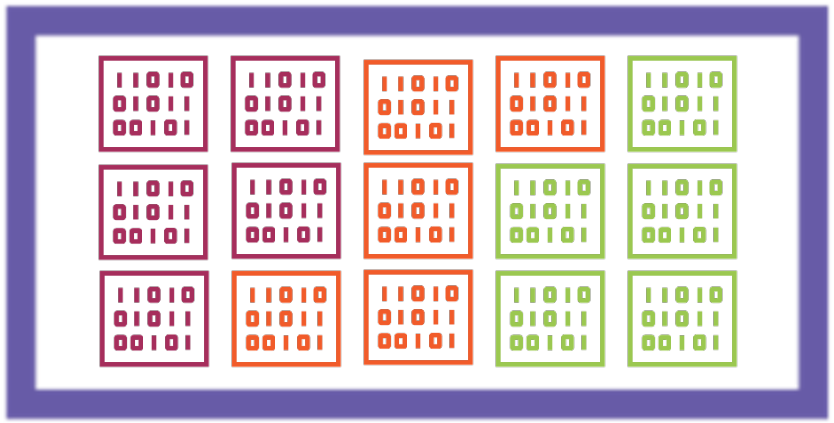

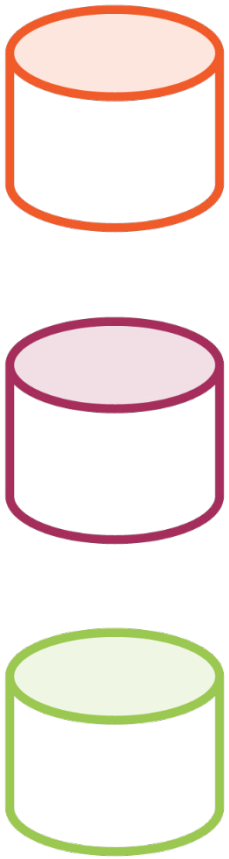

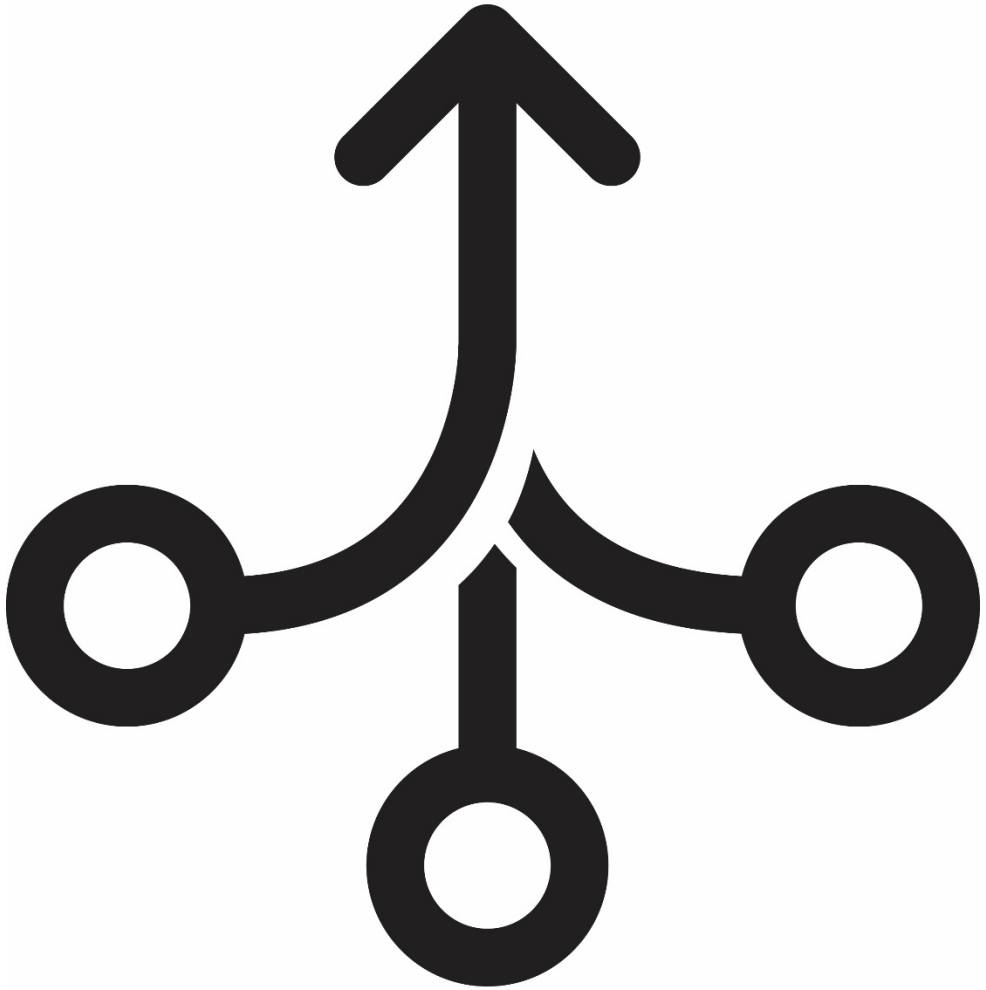

Kafka

Camel

JMS

MQTT

Storm

Flink

Twitter

Flume

ZeroMQ

RocketMQ

StreamMultipleTupleExtractor<T,K,V>

T

Incoming message data type

K

Key data type

V

Value data type

Stream Receiver

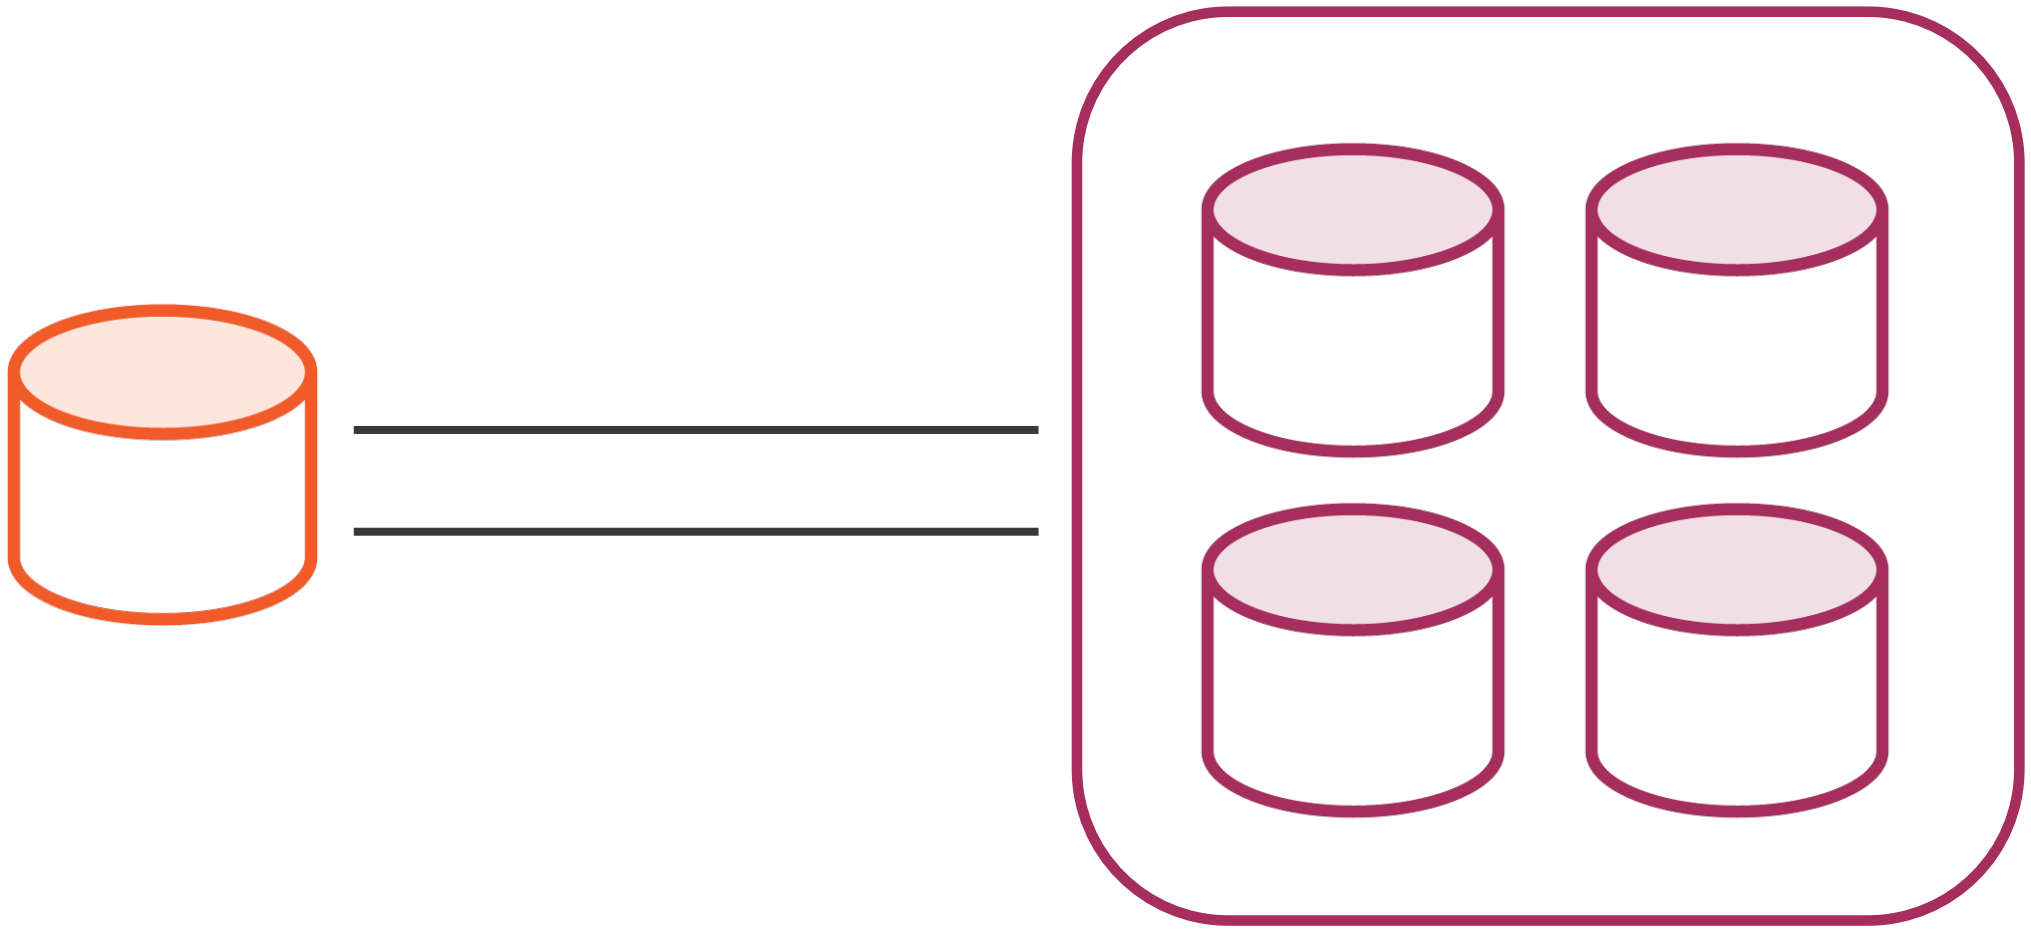

Stream Receiver

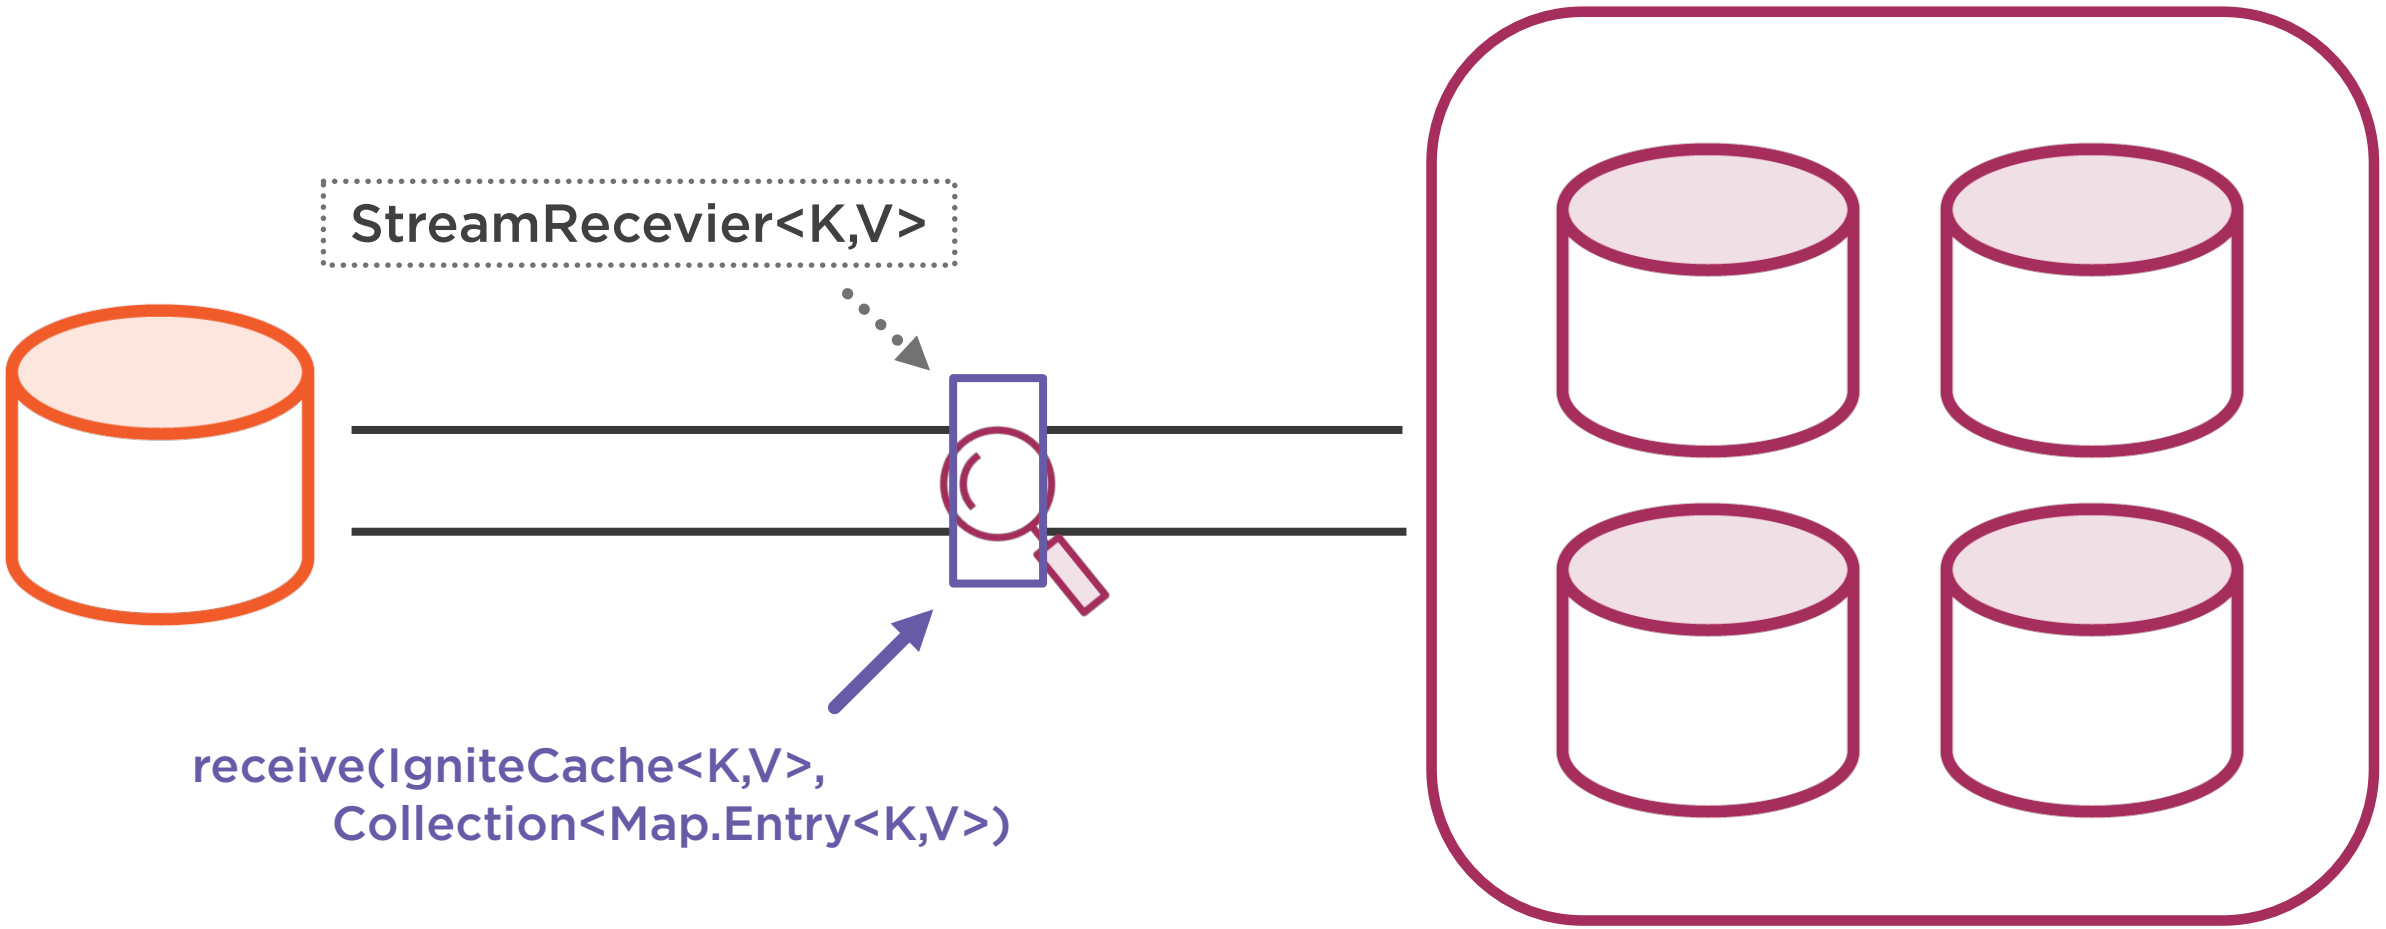

Implementations of the Stream Receiver

Stream Transformer

Stream Visitor

Stream Transformer

Convenience adapter to update existing values in streaming cache based on the previously cached value.

Stream Transformer

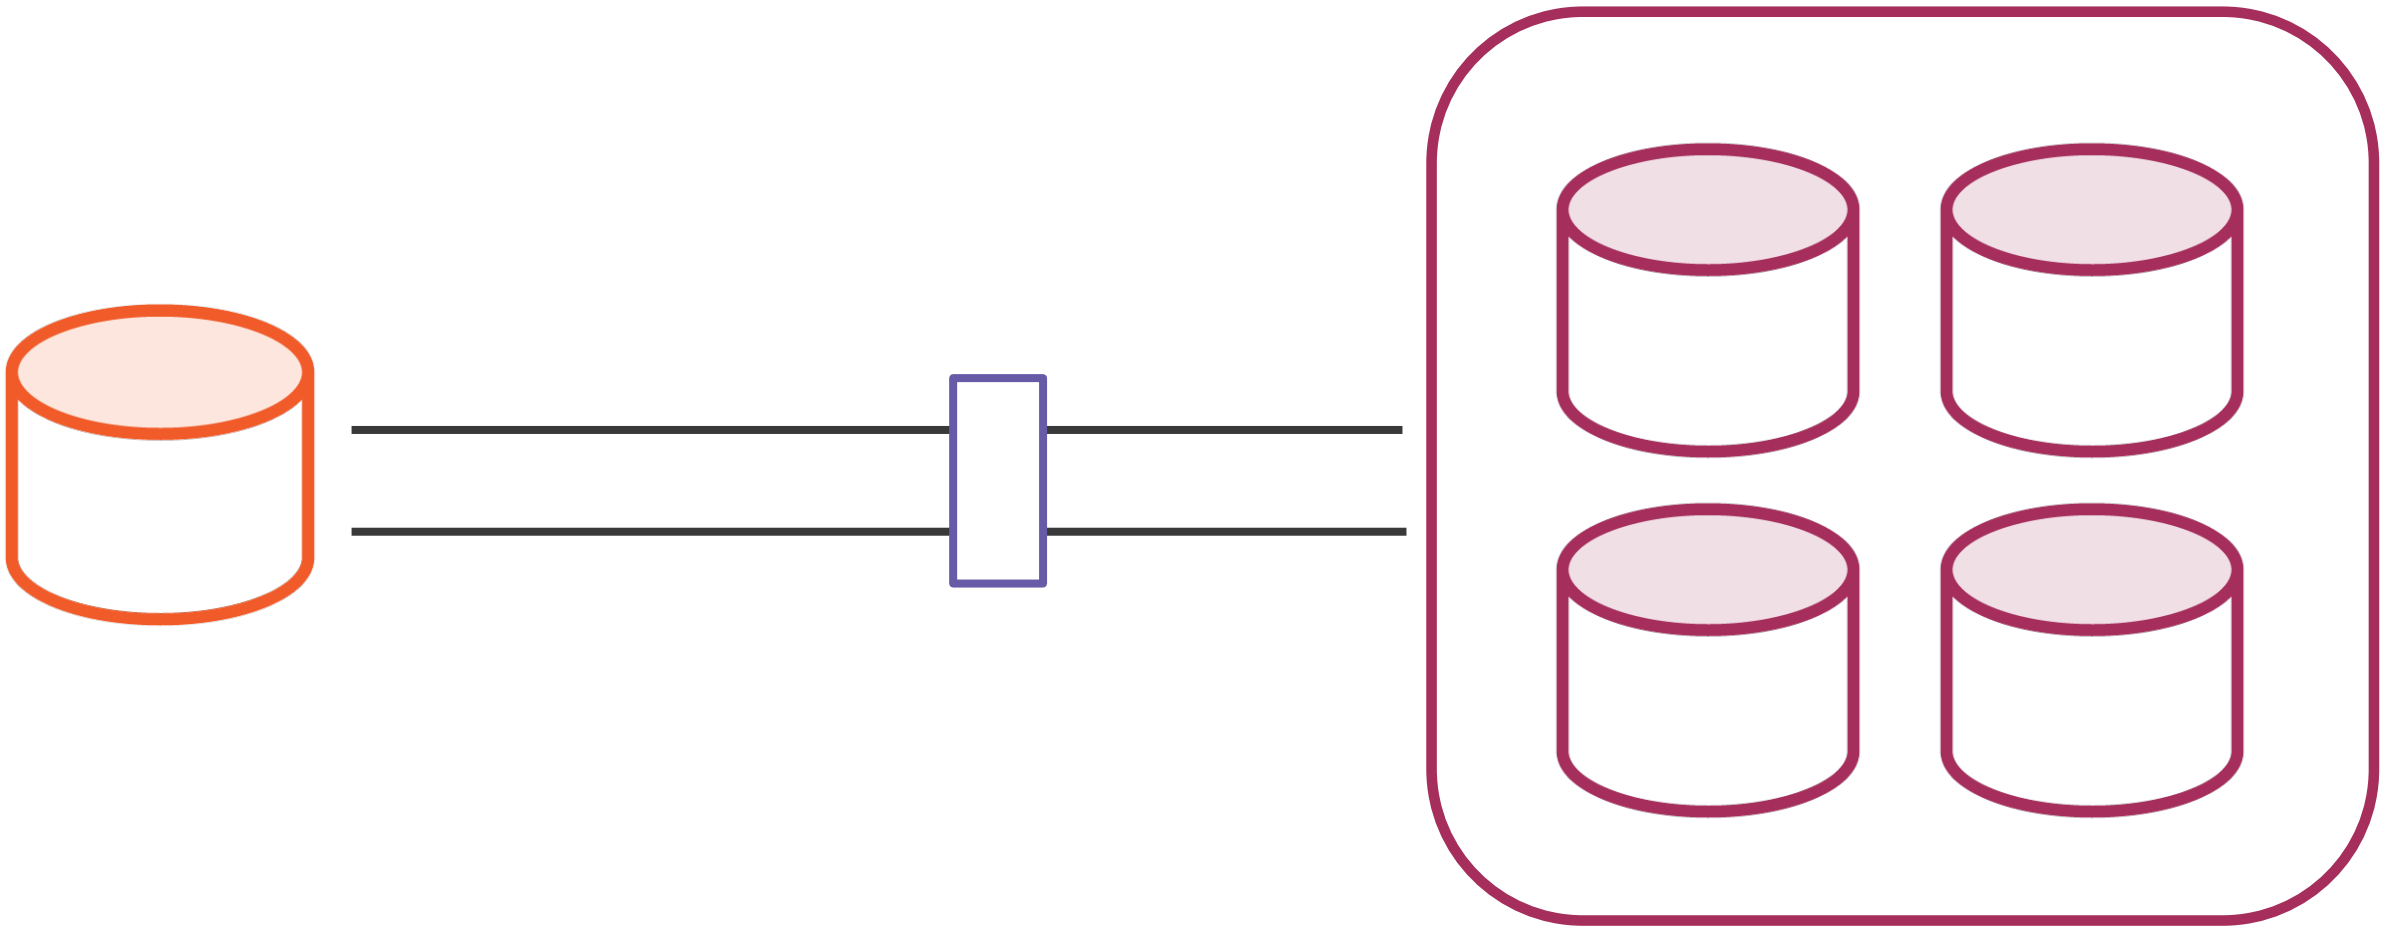

Stream Transformer

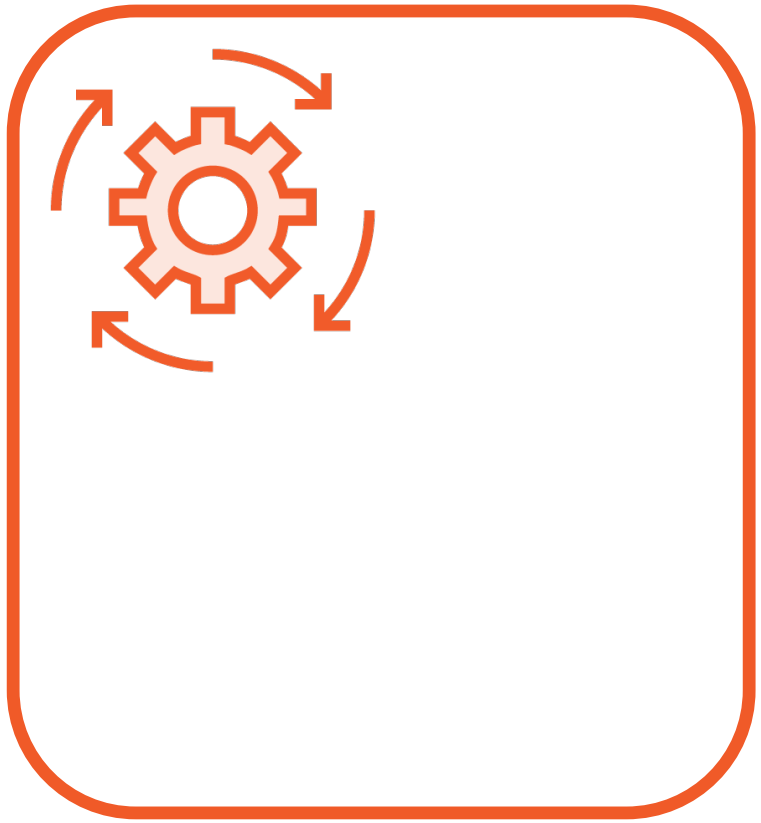

Transformer vs. Visitor

Stream Transformer	Stream Visitor
Entries with matches in the cache	All entries in the stream
Provide an entry processor	Takes “StreamReceiver” cache & entries
Automatically writes to cache	Manually write to cache

Flight number

Name of passenger

Assigned seat

Frequent flyer number

Frequent flyer status

Summary

Affinity

- Function
- Data Collocation

The SQL Grid

Continuous queries

Entry processors

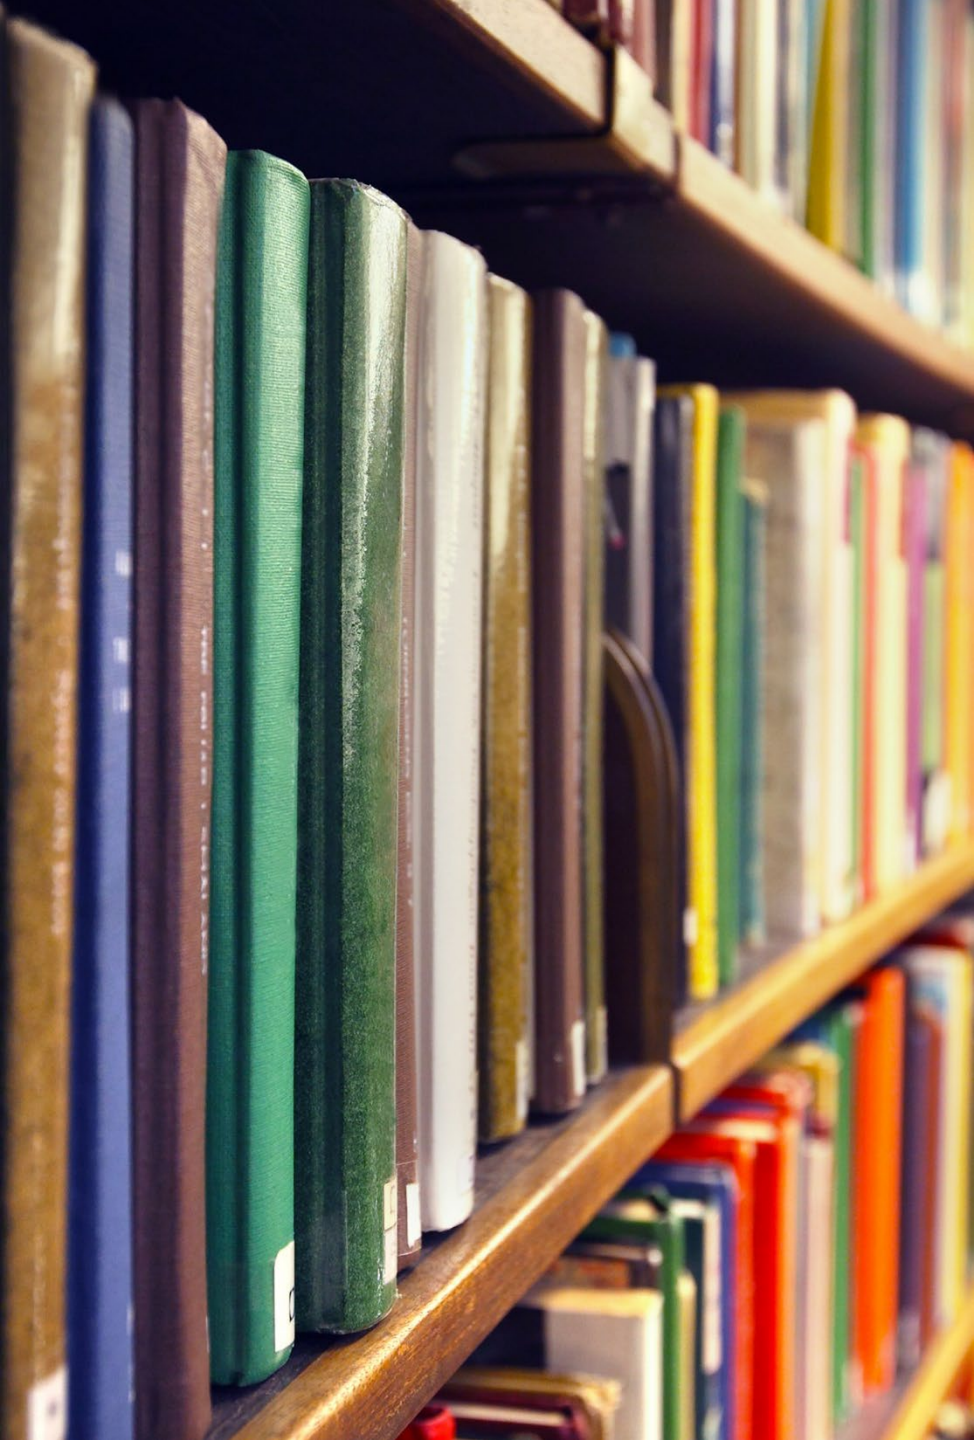

Apache Ignite Websites:

https://ignite.apache.org/

http://apacheignite-sql.readme.io/docs

Support:

https://stackoverflow.com/questions/tagged/ignite