

Exploring Distributed Computing with Ignite





Edward Curren



ENTERPRISE ARCHITECT



@EdwardCurren



<http://www.edwardcurren.com>











Scheduler

Cluster

parallel execution
mechanisms



Overview



Distributed closures

ComputeTask, ComputeJob and
ComputeTaskSession

Job scheduling

Checkpointing


Collocation with data



Distributing a Task's Workload



Distributing a Task's Workload



Task Parallelism

Call & Run




Distributing a Task's Workload



Task Parallelism

Call & Run



Data Parallelism

Apply




Go to Terminal3
Go to Terminal3
Go to Terminal3
129,130
Go to Terminal3
Go to Terminal3
75-82
Baggage drop 75-82
106-109
Go to Terminal3
Go to Terminal3
Go to Terminal3
93,94
95-97
Baggage drop 132-135
122-125

DA781	Enroute
DA294	Ready
DA319	Enroute
DA6804	Enroute
DA984	Arrived
DA681	Ready
DA965	Enroute
DA201	Arrived
DA881	Enroute
DA4901	Arrived
DA198	Ready
DA681	Enroute
DA615	Arrived
DA485	Enroute
DA354	Enroute

Go to Terminal3
Go to Terminal3
Go to Terminal3
129,130
Go to Terminal3
Go to Terminal3
75-82
Baggage drop 75-82
106-109
Go to Terminal3
Go to Terminal3
Go to Terminal3
93,94
95-97
Baggage drop 132-135
122-125

DA781	Enroute
DA294	Ready
DA319	Enroute
DA6804	Enroute
DA681	Ready
DA965	Enroute
DA881	Enroute
DA198	Ready
DA681	Enroute
DA485	Enroute
DA354	Enroute

Distributing a Task's Workload




Task Parallelism

Call & Run



Data Parallelism

Apply




Broadcast

Use ComputeTask and ComputeJob When




You need to control how jobs are mapped to nodes or define more complex reduce logic

You need to implement custom fail-over logic



Apache Ignite's ComputeTask API



`map()`

`reduce()`

Compute Task




```
graph TD; A[Compute Task] --- B[Job]; A --- C[Job]; A --- D[Job]; A --- E[Job];
```

Compute Task

Job


Job

Job

Job



Apache Ignite's ComputeTask API



`map()`

`reduce()`

map(
 <node, data>
 <node, data>
 <node, data>
)



```
graph TD; A[Compute Task] --- B[Job]; A --- C[Job]; A --- D[Job]; A --- E[Job];
```

Compute Task

Job

Job

Job

Job







Job

Job

Job

Job





```
graph TD; A[Compute Task] --- B[Job]; A --- C[Job]; A --- D[Job]; A --- E[Job];
```

Compute Task

Job

Job

Job

Job



Compute Task



Implementations of ComputeTask



ComputeTaskSplitAdapter()

{ Split()
Reduce()




Implementations of ComputeTask



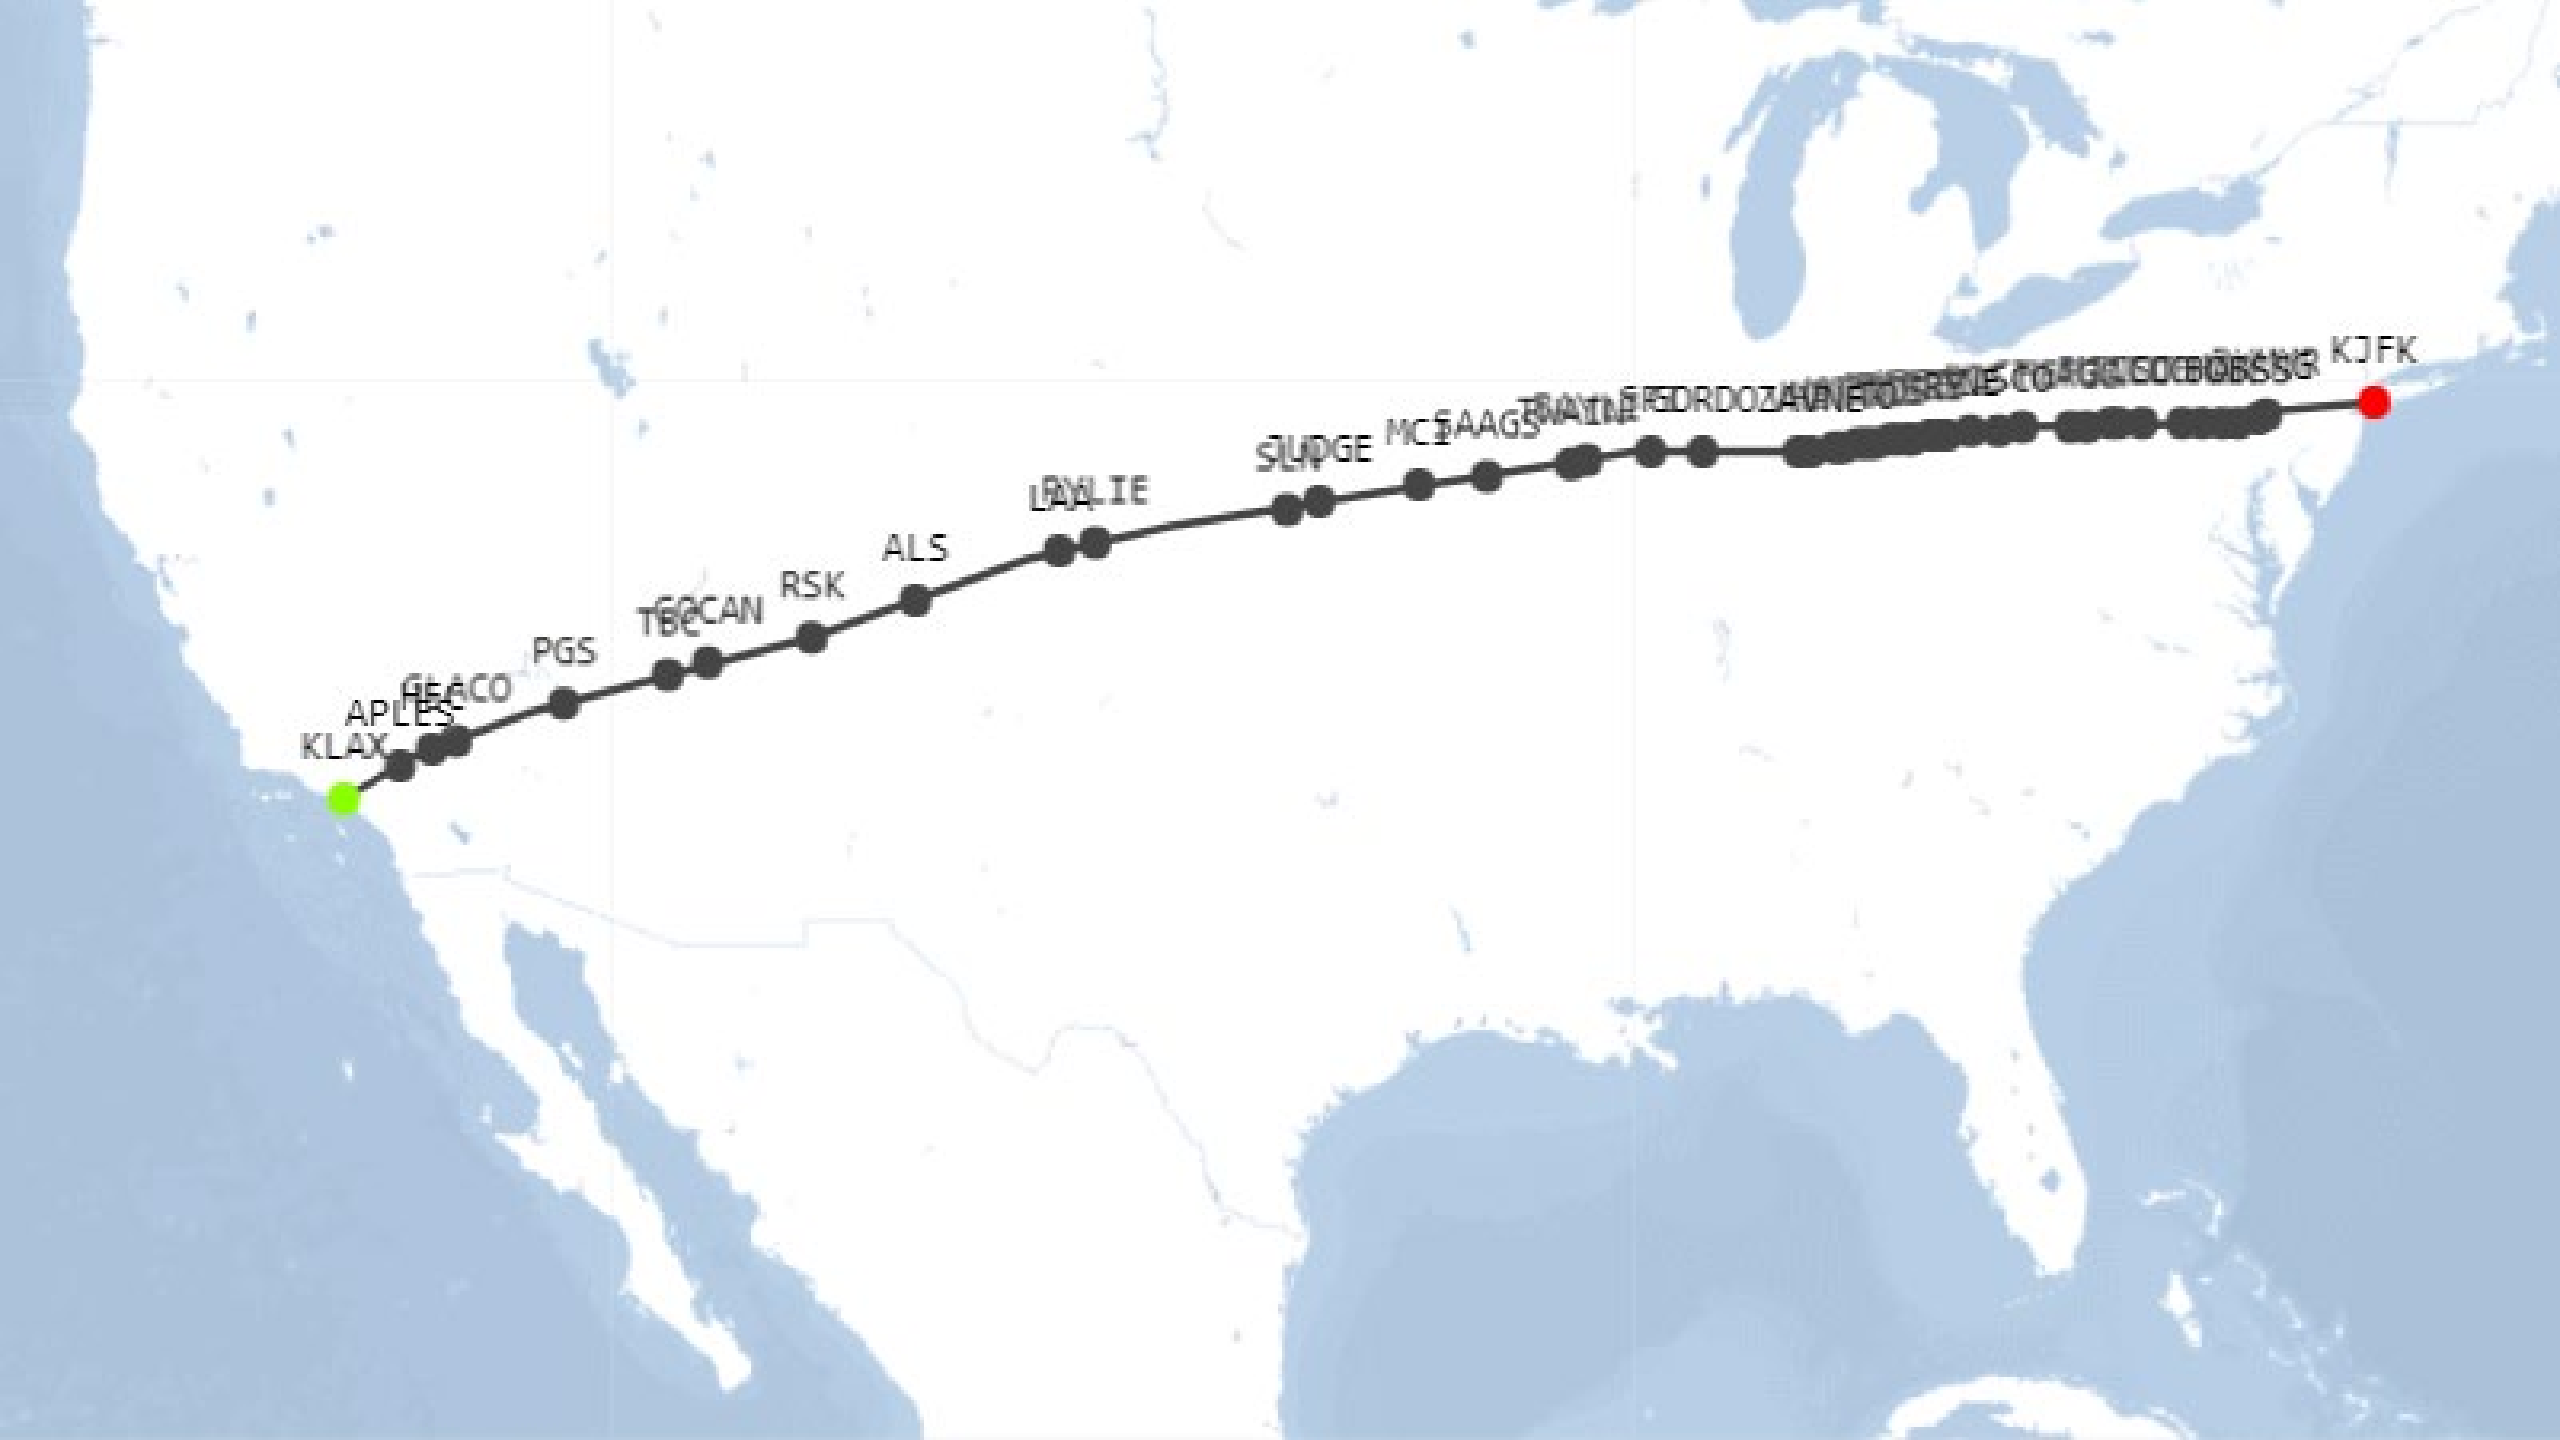
`ComputeTaskSplitAdapter()`

{ `Split()`
`Reduce()`



`ComputeTaskAdapter()`

{ `Map()`
`Reduce()`



Compute Task



```
graph TD; A[Compute Task] --- B[Job]; A --- C[Job]; A --- D[Job]; A --- E[Job];
```

Compute Task

Job


Job

Job

Job



ComputeJobAdapter




Is an abstract class

Exposes `execute()` and `cancel()` methods

Default 'no-op' `cancel()` method provided

Distributed Task Session







Flight
FlightId






Reservation
ReservationId FlightId PassengerId





Passenger
PassengerId FrequentFlyerId

FQFL
FrequentFlyerId

FFFH
FrequentFlyerId FlightId









**Task-Parallelism
("Run" and "Call")**

**Data-Parallelism
("Apply")**

Broadcast





Affinity Run & Affinity Call

Cache Name 1.. ∞

`IgniteRunnable<>()`

Affinity Key

Partition Number





~~SQL Query~~

~~SQL Fields Query~~

Scan Query

~~Continuous Query~~

~~Text Query~~



Summary



Distributed execution of tasks

The “Compute API”

Task parallelism and Data parallelism

ComputeTask and ComputeJob

The “Distributed Task Session”

Compute collocation

